

# SpaceAndTime

## Table of contents

TODO check if I checked all the residuals for all models . . . . .	2
Question 1 Spatial modelling Kingdom of the Netherlands . . . . .	2
1 a) . . . . .	2
1 b) . . . . .	4
1 c) . . . . .	5
1 d) . . . . .	7
1 e) . . . . .	13
1 f) . . . . .	31
1 g) . . . . .	34
1 h) . . . . .	36
Question 2 . . . . .	42
2 a) . . . . .	42
Trend analysis . . . . .	43
2 b) . . . . .	43
ACF . . . . .	43
PACF . . . . .	44
best model selection . . . . .	46
talk about the model being more easily explainability becaues $MA = 0$ . . . . .	51
Best model residual validation . . . . .	51
Forecasting . . . . .	52
2 c) . . . . .	53
Initial assumptions . . . . .	53
model fitting . . . . .	54
Forecast . . . . .	57
2 d) . . . . .	59
2 e) . . . . .	61
Model testing . . . . .	64
Seasonal check . . . . .	64
auto arima check . . . . .	69

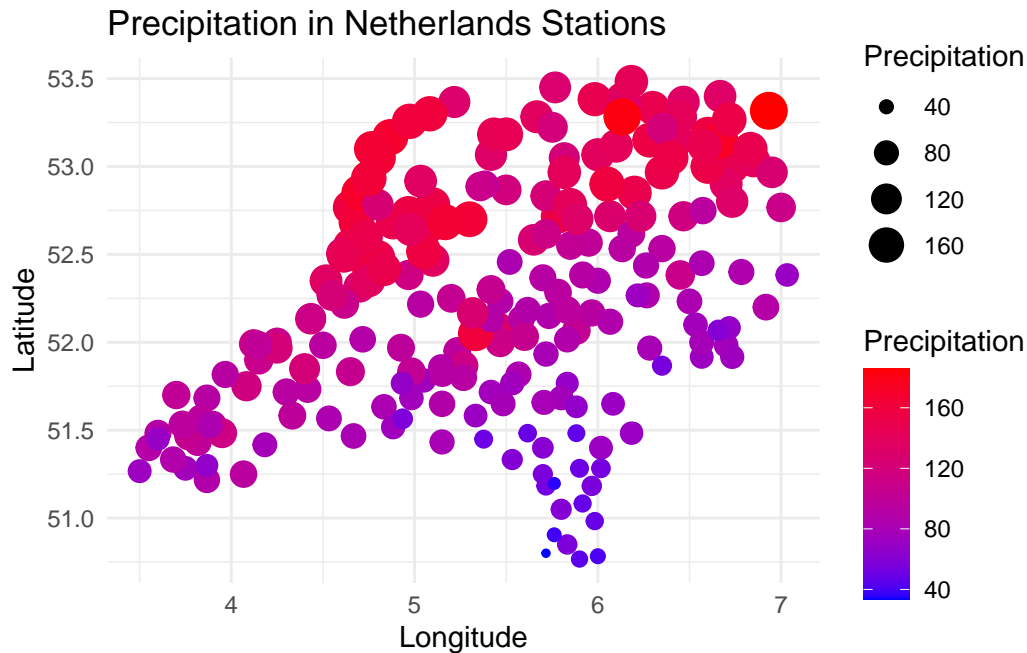
Forecasting . . . . .	70
DLM . . . . .	72
model fitting . . . . .	72
Forecast . . . . .	76
2 f) . . . . .	79
Question 3 . . . . .	83
Question 3 a) . . . . .	83
3 b) . . . . .	85
3 c) . . . . .	86
model checking Nappa . . . . .	91
Forecasting . . . . .	95
Death valley . . . . .	97
Forecast . . . . .	100
3 d) . . . . .	102

**TODO check if I checked all the residuals for all models**

## **Question 1 Spatial modelling Kingdom of the Netherlands**

### **1 a)**

```
ggplot(data = netherlandsDF) + geom_point(aes(x = longitude, y = latitude,
size = precip, color = precip)) + scale_color_continuous(low = "blue",
high = "red") + labs(title = "Precipitation in Netherlands Stations",
x = "Longitude", y = "Latitude", size = "Precipitation", color = "Precipitation") +
theme_minimal()
```



From what we can see from the data it does seem to be spatially correlated as we can see that the Dutch provinces of north Holland, Friesland and Groningen have higher precipitation and as we go south the precipitation decreases as we can see from the Dutch provinces of Zeeland, north Brabant and Limburg where precipitation is significantly lower than their northern counterparts.

From this data, latitude seems to be the biggest factor in the variation of the precipitation as the longitude only suggests some slight variations in the data.

```
# Create geodata object
precipitationNetherlands_geoR = as.geodata(netherlandsDF, coords.col = c("longitude",
  "latitude"), data.col = "precip")

summary(precipitationNetherlands_geoR)
```

Number of data points: 220

```
Coordinates summary
  longitude latitude
min      3.500    50.767
max      7.033    53.483
```

```
Distance summary
      min      max
0.001000 3.998498
```

```
Data summary
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
33.9000  80.8500 100.1500 106.5705 137.3500 185.6000
```

As we can see from the numerical summary of the data the median is different from the mean, which indicates it is not a symmetric distribution of data points and is instead positively skewed since the mean is bigger than the median. As such there are more values on the left side of the distribution.

## 1 b)

```
# set seed for reproducibility
set.seed(26041999)

# Select 3 random rows from the data frame
randomRowsPrecipitation = netherlandsDF %>%
  sample_n(3)

# Add a new column with labels
randomRowsPrecipitation$label = c("A", "B", "C")

# Print the randomly selected rows
randomRowsPrecipitation
```

```
# A tibble: 3 x 5
  station_name longitude latitude precip label
<chr>         <dbl>    <dbl>   <dbl> <chr>
1 NIJKERK          5.47     52.2    89.1 A
2 WOLPHAARTSDIJK    3.73     51.5    95.9 B
3 EEXT             6.73     53     147. C
```

```
# Remove the selected rows from the original dataset
netherlandsDF_filtered = netherlandsDF %>%
  anti_join(randomRowsPrecipitation)
```

```
Joining, by = c("station_name", "longitude", "latitude", "precip")
```

```
# Print the resulting dataframe  
netherlandsDF_filtered
```

```
# A tibble: 217 x 4  
  station_name longitude latitude precip  
  <chr>          <dbl>    <dbl>  <dbl>  
1 WEST TERSCHELLING    5.22    53.4  130.  
2 GRONINGEN-1         6.6     53.2  157.  
3 HOORN              5.07    52.6  146.  
4 HOOFDDORP          4.7     52.3  130.  
5 WINTERSWIJK         6.7     52.0   77.7  
6 KERKWERVE          3.87    51.7   91.8  
7 WESTDORPE-1         3.87    51.2   87.7  
8 OUDENBOSCH          4.53    51.6   84.2  
9 ROERMOND            5.97    51.2   56.4  
10 PETTEN             4.65    52.8  158.  
# ... with 207 more rows
```

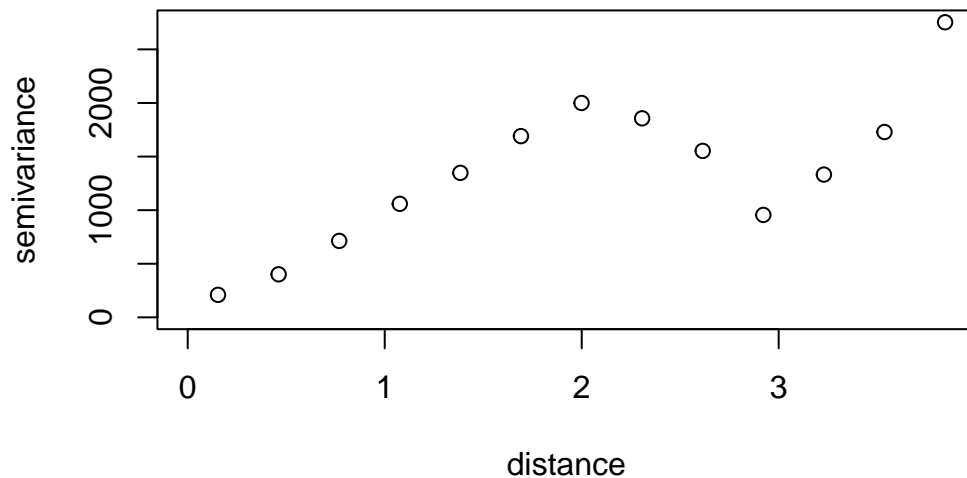
```
## recreate the geoData object with the new filtered dataframe  
precipitationNetherland_geoR = as.geodata(netherlandsDF_filtered, coords.col = c("longitude",  
  "latitude"), data.col = "precip")
```

### 1 c)

```
# # Calculate empirical variogram  
variogramPrecipitationNetherlands = variogram(precipitationNetherland_geoR)
```

variog: computing omnidirectional variogram

```
# Plot empirical variogram  
plot(variogramPrecipitationNetherlands)
```



```
variogramPrecipitationNetherlands$n
```

```
[1] 1069 2482 3384 3834 3647 3221 2534 1541 787 468 318 133 17
```

From the plotted variogram we can see there is a very clear need for a nugget as there is a non-zero value around zero distance, this value seems to be around 75 to 100 at the zero distance from how much it is decreasing.

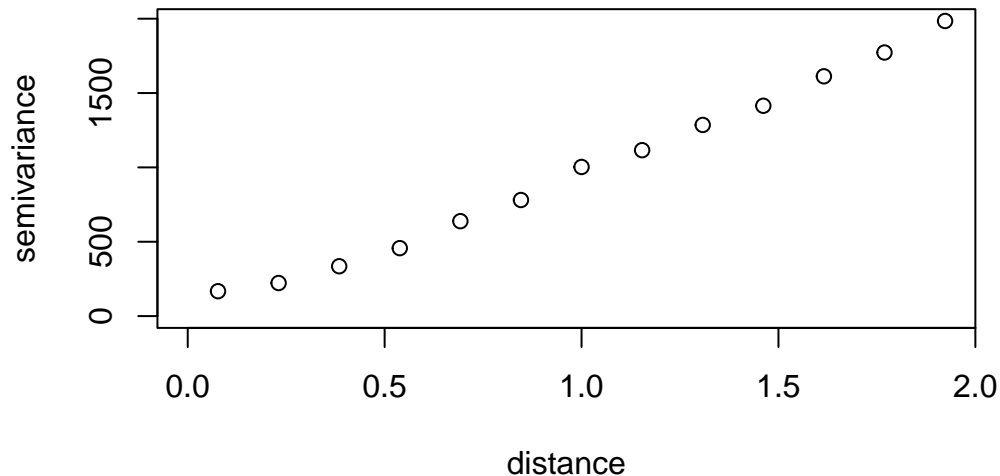
The semi variance continues to increase with distance till around the distance of 2 degrees distance wise, after this there is a decrease in variance that is not representative of the data as we are more and more uncertain the further we are from our known points, as such we will choose the distance of two as the cut off for the maximum distance.

we now change the maximum distance and recut our previous variogram.

```
variogramPrecipitationNetherlands = variog(precipitationNetherlands_geoR,
      option = "bin", max.dist = 2)
```

```
variog: computing omnidirectional variogram
```

```
plot(variogramPrecipitationNetherlands)
```



As we can see from the newly updated variogram the increase is almost linear with a curve near 0 where we can see the need for the nugget.

#### 1 d)

Now that we have the variogram we will start by fitting a model to estimate the covariance via weighted least squares. Fitting this variogram we get the estimated values of  $\sigma^2$ ,  $\phi$  and  $\tau^2$  also known as the nugget

We will first start with the default Matérn = 0,5 which is equivalent to an exponential as the form observed in the previous variogram seems to not fully linear and therefore require the curvature from a function like the exponential function to account for the behaviour at the near 0 distance.

From this we will try different models to search for the model with the best fit.

```
# ?variofit thau = nugget variability sigmasq = if the model can  
# capture more or less of the total variability phi = if the  
# correlation extends over a bigger or smaller distance loss value =
```

```
# goodness of fit (smaller means better fit)
```

```
krigingVariogramFittedDefault = variofit(variogramPrecipitationNetherlands,  
  nugget = 85)
```

```
variofit: covariance model used is matern  
variofit: weights used: npairs  
variofit: minimisation function used: optim
```

Warning in variofit(variogramPrecipitationNetherlands, nugget = 85): initial values not provided - running the default search

```
variofit: searching for best initial value ... selected values:  
          sigmasq  phi   tausq kappa  
initial.value "1984.67" "1.54" "85"  "0.5"  
status        "est"    "est"  "est" "fix"  
loss value: 818327620.928191
```

```
krigingVariogramFittedDefault
```

```
variofit: model parameters estimated by WLS (weighted least squares):  
covariance model is: matern with fixed kappa = 0.5 (exponential)  
parameter estimates:  
      tausq      sigmasq      phi  
      0.00 2561207.09    2590.45  
Practical Range with cor=0.05 for asymptotic range: 7760.294
```

```
variofit: minimised weighted sum of squares = 35360382
```

Now we will increase the kappa of the Matérn to see if the increased flexibility and smoothness leads to a better fit

```
krigingVariogramFittedMatrén1.5 = variofit(variogramPrecipitationNetherlands,  
  kappa = 1.5, nugget = 85)
```

```
variofit: covariance model used is matern  
variofit: weights used: npairs  
variofit: minimisation function used: optim
```



```
Warning in variofit(variogramPrecipitationNetherlands, kappa = 1.5, nugget =
85): initial values not provided - running the default search
```

```
variofit: searching for best initial value ... selected values:
      sigmasq  phi   tausq kappa
initial.value "1984.67" "0.62" "85"  "1.5"
status        "est"    "est"  "est" "fix"
loss value: 190482114.063677
```

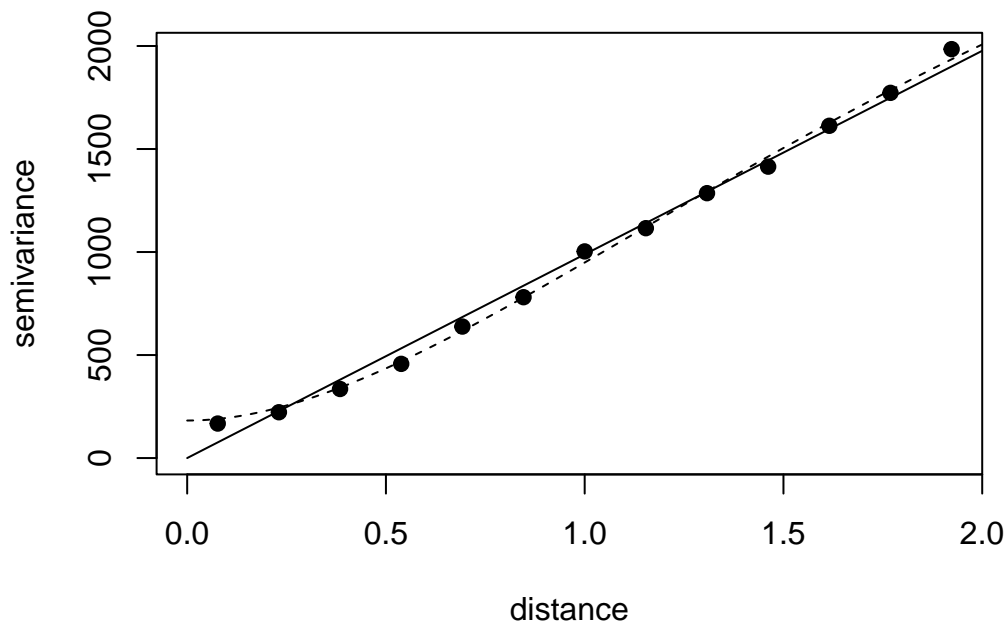
```
krigingVariogramFittedMatrén1.5
```

```
variofit: model parameters estimated by WLS (weighted least squares):
covariance model is: matern with fixed kappa = 1.5
parameter estimates:
      tausq  sigmasq      phi
182.0638 3461.1492    1.1316
Practical Range with cor=0.05 for asymptotic range: 5.368367
```

```
variofit: minimised weighted sum of squares = 15435206
```

We will first visually compare these 2 models to see which one has a better

```
par(mar = c(4, 4, 2, 2))
plot(variogramPrecipitationNetherlands, pch = 19)
lines(krigingVariogramFittedDefault)
lines(krigingVariogramFittedMatrén1.5, lty = 2)
```



```
# lines(krigingVariogramFittedMatrén2.5, lty = 3)
```

Immediately we can see that the extra flexibility of the Matérn 1,5 not only better follows the actual data, it actually accounts correctly for the initial variance from the nugget which the Matérn 0,5 does not as it simply decreases to 0.

We now will test if any additional flexibility changes can improve the model fit

We will first try to again increase the kappa to see if the model again benefits from the extra flexibility

```
krigingVariogramFittedMatrén2.0 = variofit(variogramPrecipitationNetherlands,
      kappa = 2, nugget = 85)
```

```
variofit: covariance model used is matern
variofit: weights used: npairs
variofit: minimisation function used: optim
```

```
Warning in variofit(variogramPrecipitationNetherlands, kappa = 2, nugget = 85):
initial values not provided - running the default search
```

```

variofit: searching for best initial value ... selected values:
           sigmasq  phi   tausq   kappa
initial.value "1984.67" "0.62" "198.47" "2"
status        "est"    "est"  "est"   "fix"
loss value: 227233080.389154

```

```
krigingVariogramFittedMatrén2.0
```

```

variofit: model parameters estimated by WLS (weighted least squares):
covariance model is: matern with fixed kappa = 2
parameter estimates:
      tausq  sigmasq      phi
197.6127 3037.3813   0.8386
Practical Range with cor=0.05 for asymptotic range: 4.502076

```

```
variofit: minimised weighted sum of squares = 18197964
```

Here we will instead see if the model will benefit instead from a cut of flexibility to make it less smooth

```
krigingVariogramFittedMatrén1.0 = variofit(variogramPrecipitationNetherlands,
      kappa = 1, nugget = 85)
```

```

variofit: covariance model used is matern
variofit: weights used: npairs
variofit: minimisation function used: optim

```

```
Warning in variofit(variogramPrecipitationNetherlands, kappa = 1, nugget = 85):
initial values not provided - running the default search
```

```

variofit: searching for best initial value ... selected values:
           sigmasq  phi   tausq   kappa
initial.value "1984.67" "0.92" "198.47" "1"
status        "est"    "est"  "est"   "fix"
loss value: 352001017.756763

```

```
krigingVariogramFittedMatrén1.0
```

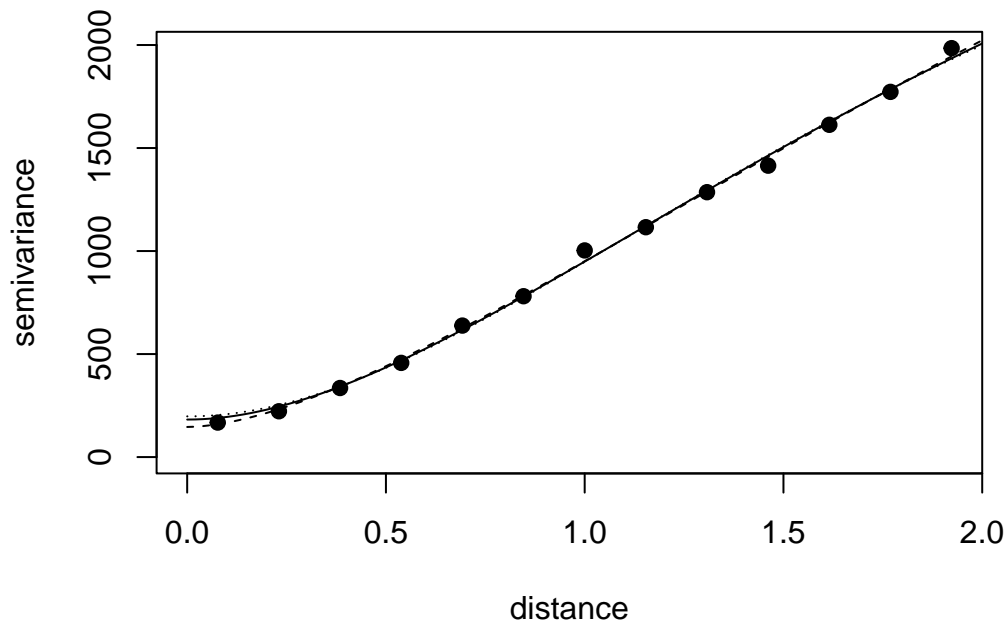
variofit: model parameters estimated by WLS (weighted least squares):  
covariance model is: matern with fixed kappa = 1  
parameter estimates:

tausq	sigmasq	phi
146.1354	4832.1862	2.0470

Practical Range with cor=0.05 for asymptotic range: 8.185098

variofit: minimised weighted sum of squares = 12019899

```
par(mar = c(4, 4, 2, 2))  
plot(variogramPrecipitationNetherlands, pch = 19)  
lines(krigingVariogramFittedMatrén1.5)  
lines(krigingVariogramFittedMatrén1.0, lty = 2)  
lines(krigingVariogramFittedMatrén2.0, lty = 3)
```



As we can see from the new graph it does seem that actually a lower flexibility Matrén has a better fit since the extra flexibility near the start and end of the data points made the models deviate too much from the points.

## 1 e)

To fit a model using the maximum likelihood we will have to try multiple initial values to make sure this is indeed the maximum likelihood and not just a local maximum.

```
# ?likest
```

```
maximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(10, 1))
```

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
       arguments for the maximisation function.  
       For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
       times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

```
maximumLikelihoodNetherlandsInitial1.10 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(1, 10))
```

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
       arguments for the maximisation function.  
       For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
       times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

```
maximumLikelihoodNetherlandsInitial100.10 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(100, 10))
```

```
-----
```

```
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
maximumLikelihoodNetherlandsInitial10.100 = likfit(precipitationNetherlands_geoR,
  ini.cov.pars = c(10, 100))
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

WARNING: estimated range is more than 10 times bigger than the biggest distance between two points

- 1) excluding spatial dependence if estimated sill is too low and/or
- 2) taking trends (covariates) into account

```
maximumLikelihoodNetherlandsInitial1.1 = likfit(precipitationNetherlands_geoR,
  ini.cov.pars = c(1, 1))
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
maximumLikelihoodNetherlandsInitial1000.1000 = likfit(precipitationNetherland_geoR,
  ini.cov.pars = c(1000, 1000))
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

WARNING: estimated range is more than 10 times bigger than the biggest distance between two points

- 1) excluding spatial dependence if estimated sill is too low and/or
- 2) taking trends (covariates) into account

```
maximumLikelihoodNetherlandsInitial500.500 = likfit(precipitationNetherland_geoR,
  ini.cov.pars = c(500, 500))
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

WARNING: estimated range is more than 10 times bigger than the biggest distance between two points

- 1) excluding spatial dependence if estimated sill is too low and/or
- 2) taking trends (covariates) into account

```
maximumLikelihoodNetherlandsInitial10.1
```

```
likfit: estimated model parameters:
```

```

      beta      tausq      sigmasq      phi
" 102.376" " 114.249" "3036.627" "    7.132"
Practical Range with cor=0.05 for asymptotic range: 21.36583

```

```
likfit: maximised log-likelihood = -896.9
```

```
maximumLikelihoodNetherlandsInitial1.10
```

```
likfit: estimated model parameters:
      beta      tausq      sigmasq      phi
" 102.449" " 114.921" "4184.402" "    9.991"
Practical Range with cor=0.05 for asymptotic range: 29.9294

```

```
likfit: maximised log-likelihood = -896.9
```

```
maximumLikelihoodNetherlandsInitial100.10
```

```
likfit: estimated model parameters:
      beta      tausq      sigmasq      phi
" 102.449" " 114.921" "4184.402" "    9.991"
Practical Range with cor=0.05 for asymptotic range: 29.9294

```

```
likfit: maximised log-likelihood = -896.9
```

```
maximumLikelihoodNetherlandsInitial10.100
```

```
likfit: estimated model parameters:
      beta      tausq      sigmasq      phi
" 102.7" " 117.7" "39625.5" " 100.0"
Practical Range with cor=0.05 for asymptotic range: 299.5729

```

```
likfit: maximised log-likelihood = -897.8
```

```
maximumLikelihoodNetherlandsInitial1.1
```

```
likfit: estimated model parameters:
      beta      tausq      sigmasq      phi

```



```
" 102.376" " 114.249" "3036.627" " 7.132"
```

```
Practical Range with cor=0.05 for asymptotic range: 21.36583
```

```
likfit: maximised log-likelihood = -896.9
```

```
maximumLikelihoodNetherlandsInitial1000.1000
```

```
likfit: estimated model parameters:
```

```
      beta      tausq      sigmasq      phi  
" 103.2" " 148.3" "246755.1" " 1000.0"
```

```
Practical Range with cor=0.05 for asymptotic range: 2995.732
```

```
likfit: maximised log-likelihood = -900.9
```

```
maximumLikelihoodNetherlandsInitial500.500
```

```
likfit: estimated model parameters:
```

```
      beta      tausq      sigmasq      phi  
" 103.1" " 145.0" "130046.9" " 500.0"
```

```
Practical Range with cor=0.05 for asymptotic range: 1497.866
```

```
likfit: maximised log-likelihood = -900.1
```

As we can see from the maximised log-likelihoods it does seem that have indeed reached the maximum log-likelihood as none of the values are too fastly different and the likelihood worsened as we started to increase much more our starting values.

Next we will try the REML that takes into account the fact that some of the parameters of the model are related to the variance of the residuals and not the mean.

```
REMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(10, 1), lik.method = "REML")
```

```
-----  
likfit: likelihood maximisation using the function optim.
```

```
likfit: Use control() to pass additional  
arguments for the maximisation function.
```

```
For further details see documentation for optim.
```

```
likfit: It is highly advisable to run this function several
```

times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!

-----  
likfit: end of numerical maximisation.

WARNING: estimated range is more than 10 times bigger than the biggest distance between two  
1) excluding spatial dependence if estimated sill is too low and/or  
2) taking trends (covariates) into account

```
REMLmaximumLikelihoodNetherlandsInitial1.10 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(1, 10), lik.method = "REML")
```

-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
arguments for the maximisation function.  
For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!

-----  
likfit: end of numerical maximisation.

WARNING: estimated range is more than 10 times bigger than the biggest distance between two  
1) excluding spatial dependence if estimated sill is too low and/or  
2) taking trends (covariates) into account

```
REMLmaximumLikelihoodNetherlandsInitial100.1 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(100, 1), lik.method = "REML")
```

-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
arguments for the maximisation function.  
For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!

-----  
likfit: end of numerical maximisation.

WARNING: estimated range is more than 10 times bigger than the biggest distance between two  
 1) excluding spatial dependence if estimated sill is too low and/or  
 2) taking trends (covariates) into account

```
REMLmaximumLikelihoodNetherlandsInitial1.100 = likfit(precipitationNetherlands_geoR,  
  ini.cov.pars = c(1, 100), lik.method = "REML")
```

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
  arguments for the maximisation function.  
  For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
  times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

WARNING: estimated range is more than 10 times bigger than the biggest distance between two  
 1) excluding spatial dependence if estimated sill is too low and/or  
 2) taking trends (covariates) into account

```
REMLmaximumLikelihoodNetherlandsInitial10.1
```

```
likfit: estimated model parameters:  
      beta      tausq      sigmasq      phi  
" 102.61" " 114.60" "18253.12" " 43.35"  
Practical Range with cor=0.05 for asymptotic range: 129.8698  
  
likfit: maximised log-likelihood = -888.9
```

```
REMLmaximumLikelihoodNetherlandsInitial1.10
```

```
likfit: estimated model parameters:  
      beta      tausq      sigmasq      phi  
" 102.62" " 114.83" "18124.93" " 43.19"  
Practical Range with cor=0.05 for asymptotic range: 129.3933  
  
likfit: maximised log-likelihood = -888.9
```

```
REMLmaximumLikelihoodNetherlandsInitial100.1
```

```
likfit: estimated model parameters:
```

```
      beta      tausq      sigmasq      phi  
" 102.61" " 114.60" "18253.12" " 43.35"
```

```
Practical Range with cor=0.05 for asymptotic range: 129.8698
```

```
likfit: maximised log-likelihood = -888.9
```

```
REMLmaximumLikelihoodNetherlandsInitial1.100
```

```
likfit: estimated model parameters:
```

```
      beta      tausq      sigmasq      phi  
" 102.7" " 116.5" "40885.8" " 100.0"
```

```
Practical Range with cor=0.05 for asymptotic range: 299.5729
```

```
likfit: maximised log-likelihood = -888.9
```

As we can see from these new models, the new likelihood method actually did improve our model as we have a lower log-likelihood.

Since the data did not seem to be perfectly stationary as seen in the previous questions, we will now check if adding a linear trend improves our model

```
linearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherland_geoR,  
  trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML")
```

```
-----  
likfit: likelihood maximisation using the function optim.
```

```
likfit: Use control() to pass additional
```

```
  arguments for the maximisation function.
```

```
  For further details see documentation for optim.
```

```
likfit: It is highly advisable to run this function several
```

```
  times with different initial values for the parameters.
```

```
likfit: WARNING: This step can be time demanding!
```

```
-----  
likfit: end of numerical maximisation.
```

```
linearREMLmaximumLikelihoodNetherlandsInitial1.10 = likfit(precipitationNetherland_geoR,
  trend = "1st", ini.cov.pars = c(1, 10), lik.method = "REML")
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
linearREMLmaximumLikelihoodNetherlandsInitial100.10 = likfit(precipitationNetherland_geoR,
  trend = "1st", ini.cov.pars = c(100, 10), lik.method = "REML")
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
linearREMLmaximumLikelihoodNetherlandsInitial10.100 = likfit(precipitationNetherland_geoR,
  trend = "1st", ini.cov.pars = c(10, 100), lik.method = "REML")
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
```

-----  
likfit: end of numerical maximisation.

WARNING: estimated range is more than 10 times bigger than the biggest distance between two  
1) excluding spatial dependence if estimated sill is too low and/or  
2) taking trends (covariates) into account

`linearREMLmaximumLikelihoodNetherlandsInitial10.1`

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2194.6747"	" -11.4587"	" 45.2352"	" 112.7395"	" 191.5497"	" 0.4033"

Practical Range with cor=0.05 for asymptotic range: 1.208203

likfit: maximised log-likelihood = -872

`linearREMLmaximumLikelihoodNetherlandsInitial1.10`

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2084.141"	" -6.973"	" 42.674"	" 125.053"	" 3238.867"	" 9.905"

Practical Range with cor=0.05 for asymptotic range: 29.67433

likfit: maximised log-likelihood = -873

`linearREMLmaximumLikelihoodNetherlandsInitial100.10`

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2084.141"	" -6.973"	" 42.674"	" 125.053"	" 3238.867"	" 9.905"

Practical Range with cor=0.05 for asymptotic range: 29.67433

likfit: maximised log-likelihood = -873

`linearREMLmaximumLikelihoodNetherlandsInitial10.100`

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2084.09"	"-6.91"	"42.67"	"125.89"	"32191.57"	"100.00"

Practical Range with cor=0.05 for asymptotic range: 299.5701

likfit: maximised log-likelihood = -873

From these models we can see that the linear trend does indeed improve our model, now we will check if there are any other covariance functions that can improve the model further.

```
Matren0.5linearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherlan
  trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML", cov.model = "matern",
  kappa = 0.5)
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
Matren1.0linearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherlan
  trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML", cov.model = "matern",
  kappa = 1)
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
Matren1.5linearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherlan
  trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML", cov.model = "matern",
  kappa = 1.5)
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
Matren2.0linearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherlan
  trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML", cov.model = "matern",
  kappa = 2)
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
Matren2.5linearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherlan
  trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML", cov.model = "matern",
  kappa = 2.5)
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
```



```
likfit: It is highly advisable to run this function several
        times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
```

```
-----
likfit: end of numerical maximisation.
```

```
Matren0.5linearREMLmaximumLikelihoodNetherlandsInitial10.1
```

```
likfit: estimated model parameters:
```

beta0	beta1	beta2	tausq	sigmasq	phi
"-2194.6747"	" -11.4587"	" 45.2352"	" 112.7395"	" 191.5497"	" 0.4033"

Practical Range with cor=0.05 for asymptotic range: 1.208203

```
likfit: maximised log-likelihood = -872
```

```
Matren1.0linearREMLmaximumLikelihoodNetherlandsInitial10.1
```

```
likfit: estimated model parameters:
```

beta0	beta1	beta2	tausq	sigmasq	phi
"-2221.5117"	" -12.1627"	" 45.8184"	" 125.4406"	" 160.1807"	" 0.2088"

Practical Range with cor=0.05 for asymptotic range: 0.8347704

```
likfit: maximised log-likelihood = -871.3
```

```
Matren1.5linearREMLmaximumLikelihoodNetherlandsInitial10.1
```

```
likfit: estimated model parameters:
```

beta0	beta1	beta2	tausq	sigmasq	phi
"-2233.7541"	" -12.4219"	" 46.0784"	" 129.6185"	" 149.9124"	" 0.1529"

Practical Range with cor=0.05 for asymptotic range: 0.7254113

```
likfit: maximised log-likelihood = -871.1
```

```
Matren2.0linearREMLmaximumLikelihoodNetherlandsInitial10.1
```

```
likfit: estimated model parameters:
```

beta0	beta1	beta2	tausq	sigmasq	phi
"-2240.7056"	" -12.5595"	" 46.2252"	" 131.4685"	" 144.9394"	" 0.1248"

Practical Range with cor=0.05 for asymptotic range: 0.669996

likfit: maximised log-likelihood = -871.1

```
Matren2.5linearREMLmaximumLikelihoodNetherlandsInitial10.1
```

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2245.1703"	" -12.6453"	" 46.3192"	" 132.4280"	" 142.0693"	" 0.1074"

Practical Range with cor=0.05 for asymptotic range: 0.6358818

likfit: maximised log-likelihood = -871

It does seem that the matrén covariance function did indeed slightly improved the model so we will compare it to a model using a spherical covariance function. The spherical covariance function is appropriate for this scenario has the spatial correlation between data points decreases rapidly as the distance between the points increases and we are limited with the range of correlation has after 2 degrees of distance we loose sensible correlation, hence the cut in the variogram.

```
SphericallinearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherlan  
trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML", cov.model = "spherical")
```

kappa not used for the spherical correlation function

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
arguments for the maximisation function.  
For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

```
SphericallinearREMLmaximumLikelihoodNetherlandsInitial1.10 = likfit(precipitationNetherlan  
trend = "1st", ini.cov.pars = c(1, 10), lik.method = "REML", cov.model = "spherical")
```

kappa not used for the spherical correlation function

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
        arguments for the maximisation function.
        For further details see documentation for optim.
likfit: It is highly advisable to run this function several
        times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
SphericallinearREMLmaximumLikelihoodNetherlandsInitial100.10 = likfit(precipitationNetherl
trend = "1st", ini.cov.pars = c(100, 10), lik.method = "REML", cov.model = "spherical"
```

kappa not used for the spherical correlation function

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
        arguments for the maximisation function.
        For further details see documentation for optim.
likfit: It is highly advisable to run this function several
        times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
SphericallinearREMLmaximumLikelihoodNetherlandsInitial10.100 = likfit(precipitationNetherl
trend = "1st", ini.cov.pars = c(10, 100), lik.method = "REML", cov.model = "spherical"
```

kappa not used for the spherical correlation function

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
        arguments for the maximisation function.
        For further details see documentation for optim.
likfit: It is highly advisable to run this function several
        times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
```

likfit: end of numerical maximisation.

WARNING: estimated range is more than 10 times bigger than the biggest distance between two

- 1) excluding spatial dependence if estimated sill is too low and/or
- 2) taking trends (covariates) into account

SphericallinearREMLmaximumLikelihoodNetherlandsInitial10.1

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2182.422" "	-11.539" "	45.002" "	127.066" "	200.406" "	1.003"

Practical Range with cor=0.05 for asymptotic range: 1.00273

likfit: maximised log-likelihood = -872.2

SphericallinearREMLmaximumLikelihoodNetherlandsInitial1.10

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2082.850" "	-6.686" "	42.622" "	125.174" "	2151.559" "	9.905"

Practical Range with cor=0.05 for asymptotic range: 9.905214

likfit: maximised log-likelihood = -873

SphericallinearREMLmaximumLikelihoodNetherlandsInitial100.10

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2082.850" "	-6.686" "	42.622" "	125.174" "	2151.559" "	9.905"

Practical Range with cor=0.05 for asymptotic range: 9.905214

likfit: maximised log-likelihood = -873

SphericallinearREMLmaximumLikelihoodNetherlandsInitial10.100

likfit: estimated model parameters:

```
      beta0      beta1      beta2      tausq      sigmasq      phi
"-2083.890" "   -6.902" "   42.667" "  125.695" "21538.008" "   99.999"
Practical Range with cor=0.05 for asymptotic range: 99.99893
```

likfit: maximised log-likelihood = -873

As we can see the spherical covariance function does not provide as good of a fit as the Matérn.

We will now validate the model by doing cross-validation on the model.

```
xv.ml = xvalid(precipitationNetherland_geoR, model = Matren2.5linearREMLmaximumLikelihoodM)
```

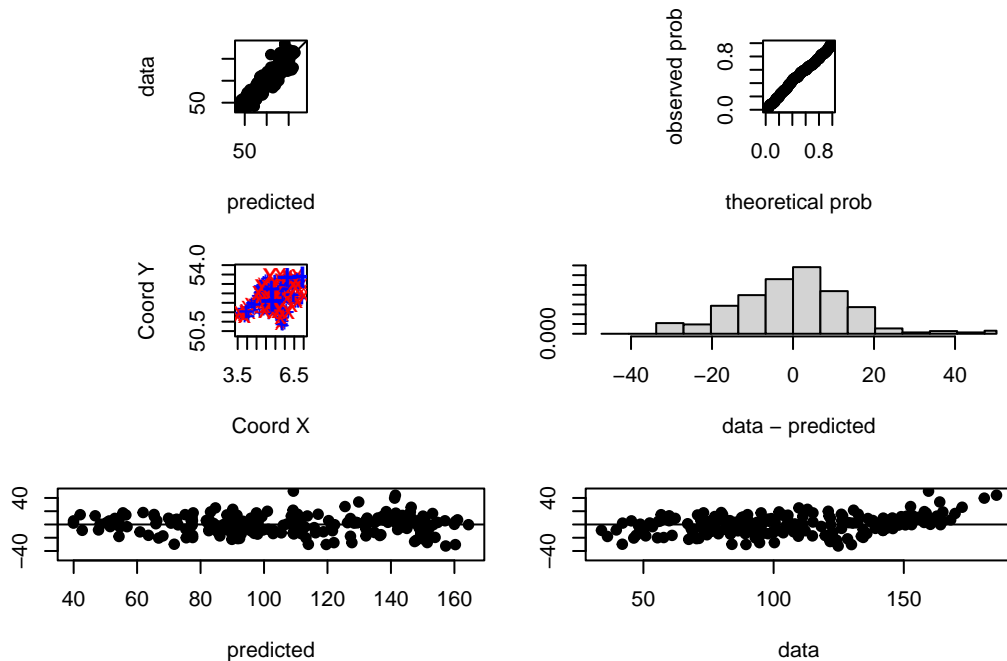
xvalid: number of data locations = 217

xvalid: number of validation locations = 217

xvalid: performing cross-validation at location ... 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13

xvalid: end of cross-validation

```
par(mfrow = c(3, 2), mar = c(4, 2, 2, 2))
plot(xv.ml, error = TRUE, std.error = FALSE, pch = 19)
```



From these plots we can see that the residuals seem mostly normal without any quickly identifiable patterns or bias.

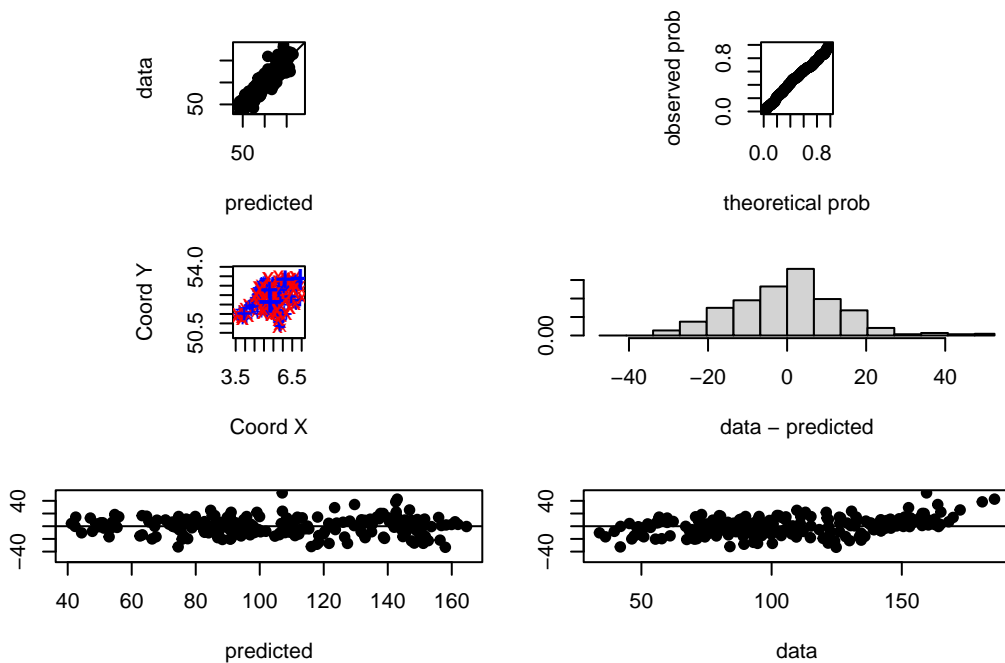
From the first top left graph we can see however that we seem to slightly underestimate more data points.

To account for bias we will also perform cross-validation to the next best performing model using the spherical function instead

```
xv.ml = xvalid(precipitationNetherland_geoR, model = SphericallylinearREMLmaximumLikelihoodN
```

```
xvalid: number of data locations      = 217
xvalid: number of validation locations = 217
xvalid: performing cross-validation at location ... 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13
xvalid: end of cross-validation
```

```
par(mfrow = c(3, 2), mar = c(4, 2, 2, 2))
plot(xv.ml, error = TRUE, std.error = FALSE, pch = 19)
```



As we can see the data at the start seems to be systematically underestimated and at the end it seems to be overestimated. Furthermore the theoretical data plot seems to be less linear.

This confirms that the spherical covariance function is indeed a worse model than the Matérn model.

## 1 f)

First we will start by making the predictions using the variogram

```
spatialPointsABC = randomRowsPrecipitation[, c("longitude", "latitude")]

# Set the kriging control parameters
krigeControl = kriging.control(type.krige = "OK", cov.model = krigingVariogramFittedMatrén1.0,
  cov.pars = krigingVariogramFittedMatrén1.0$cov.pars)

# Kriging with the fitted variogram model
krigeResults = kriging.conv(precipitationNetherlands_geoR, locations = spatialPointsABC,
  krige = krigeControl)
```

krige.conv: model with constant mean

krige.conv: Kriging performed using global neighbourhood

```
# Extract predictions from the kriging results
predictions = krigeResults$predict

# Compare the predicted values with the actual precipitation values

actualPrecipitationValues = randomRowsPrecipitation[, 4]

comparisonvariogram = data.frame(actualPrecipitationValues, predictions)
```

Now for the maximum likelihood function

```
# done above spatialPointsABC = randomRowsPrecipitation[,
# c('longitude', 'latitude')]

# Set the kriging control parameters
krigeControl = kriging.control(type.krige = "OK", cov.model = Matren2.5linearREMLmaximumLikelihoodNetherlandsInitial10.1,
  cov.pars = Matren2.5linearREMLmaximumLikelihoodNetherlandsInitial10.1$cov.pars)
```

```
# Kriging with the fitted variogram model
krigeResults = krige.conv(precipitationNetherlands_geoR, locations = spatialPointsABC,
  krige = krigeControl)
```

krige.conv: model with constant mean  
krige.conv: Kriging performed using global neighbourhood

```
# Extract predictions from the kriging results
predictions = krigeResults$predict

# Compare the predicted values with the actual precipitation values

# done above actualPrecipitationValues = randomRowsPrecipitation[,4]
comparisonMaximumLikelihood = data.frame(actualPrecipitationValues, predictions)
```

Now that we have made the predictions for our 2 models we will check the predicted values compared to the real values for each of the models.

```
comparisonvariogram
```

```
precip predictions
1  89.1    95.41904
2  95.9    98.90558
3 147.2   147.82038
```

```
comparisonMaximumLikelihood
```

```
precip predictions
1  89.1    99.00780
2  95.9    99.36194
3 147.2   140.07734
```

```
# Calculate Mean Absolute Error (MAE)
MAEVariogram = mean(abs(comparisonvariogram$precip - comparisonvariogram$predictions))
MAEMaximumLikelihood = mean(abs(comparisonMaximumLikelihood$precip - comparisonMaximumLike
```



```

# Calculate Mean Squared Error (MSE)
mseVariogram = mean((comparisonvariogram$precip - comparisonvariogram$predictions)^2)
mseMaximumLikelihood = mean((comparisonMaximumLikelihood$precip - comparisonMaximumLikelihood$predictions)^2)

# Calculate Root Mean Squared Error (RMSE)
rmseVariogram = sqrt(mseVariogram)
rmseMaximumLikelihood = sqrt(mseMaximumLikelihood)

# Display the calculated metrics
cat("Mean Absolute Error (MAE) of the Variogram:", MAEVariogram, "\n")

```

Mean Absolute Error (MAE) of the Variogram: 3.315002

```

cat("Mean Squared Error (MSE) of the Variogram:", mseVariogram, "\n")

```

Mean Squared Error (MSE) of the Variogram: 16.44957

```

cat("Root Mean Squared Error (RMSE) of the Variogram:", rmseVariogram, "\n")

```

Root Mean Squared Error (RMSE) of the Variogram: 4.055807

```

cat("\n\n")

```

```

cat("Mean Absolute Error (MAE) of the maximum likelihood:", MAEMaximumLikelihood, "\n")

```

Mean Absolute Error (MAE) of the maximum likelihood: 6.830801

```

cat("Mean Squared Error (MSE) of the maximum likelihood:", mseMaximumLikelihood, "\n")

```

Mean Squared Error (MSE) of the maximum likelihood: 53.62729

```
cat("Root Mean Squared Error (RMSE) of the maximum likelihood:", rmseMaximumLikelihood,
    "\n")
```

Root Mean Squared Error (RMSE) of the maximum likelihood: 7.323066

As we can see from both the real values and the MAE, MSE and RMSE the variogram has as much better performance predicting those 3 points than our maximum likelihood model

1 g)

```
# Determine the range of the coordinates
xRange = range(precipitationNetherlands_geoR$coords[, 1])
yRange = range(precipitationNetherlands_geoR$coords[, 2])

# Create a grid with 0.05-degree spacing
gridPoints = expand.grid(x = seq(xRange[1], xRange[2], by = 0.05), y = seq(yRange[1],
    yRange[2], by = 0.05))

# Kriging with the fitted variogram model
krigeResults = krige.conv(precipitationNetherlands_geoR, locations = gridPoints,
    krige = krigeControl)
```

krige.conv: model with constant mean

krige.conv: Kriging performed using global neighbourhood

```
# Create a data frame for the grid points with the predicted mean and
# variance
gridData = data.frame(gridPoints, mean = krigeResults$predict, variance = krigeResults$kri

# Mean plot
meanPlot = ggplot(gridData, aes(x = x, y = y, fill = mean)) + geom_tile() +
    scale_fill_gradientn(colors = c("blue", "green", "yellow", "red")) +
    theme_minimal() + ggtitle("Mean Plot") + labs(x = "Longitude", y = "Latitude",
    fill = "Mean")

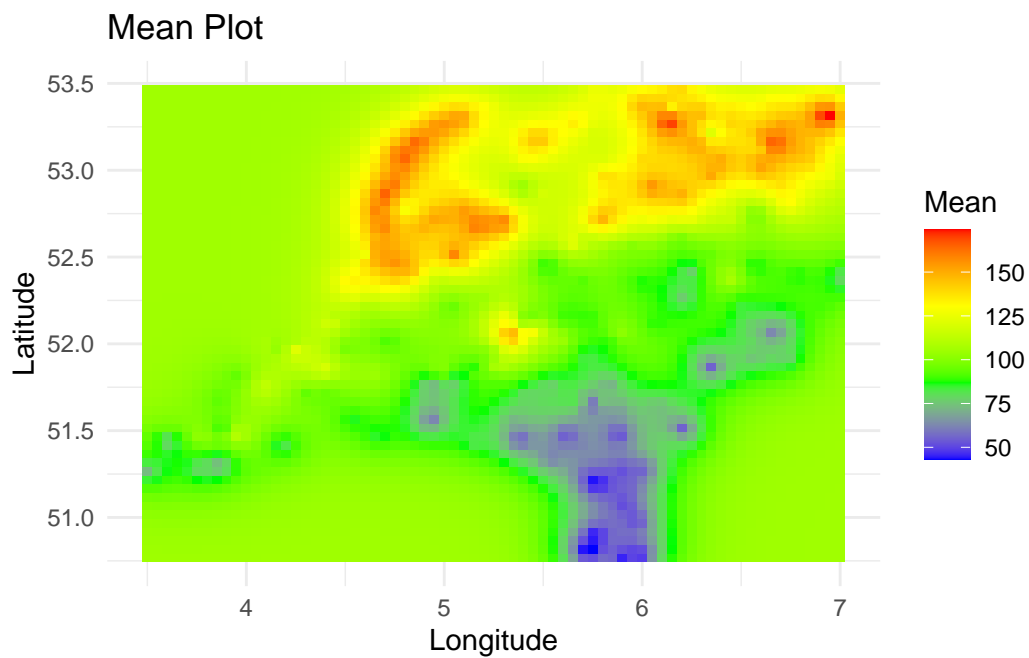
# Variance plot
variancePlot = ggplot(gridData, aes(x = x, y = y, fill = variance)) + geom_tile() +
```

```

scale_fill_gradientn(colors = c("white", "blue", "green", "yellow", "red")) +
theme_minimal() + ggtitle("Variance Plot") + labs(x = "Longitude", y = "Latitude",
fill = "Variance")

# Display the plots
print(meanPlot)

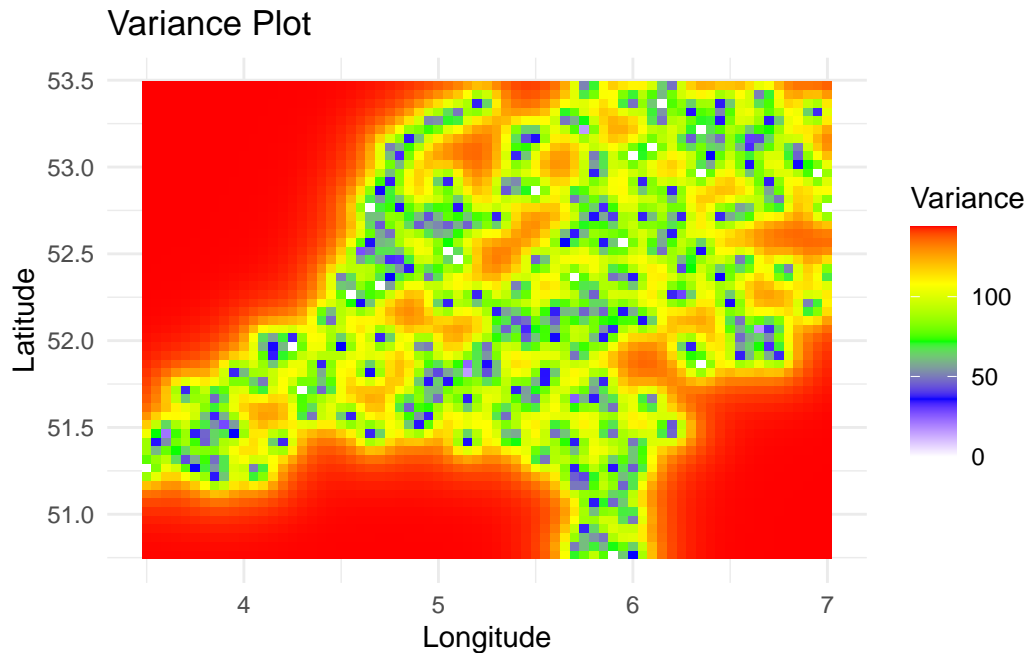
```



```

print(variancePlot)

```



## 1 h)

For the priors I will be using the estimated values from our latest maximum likelihood model.

```
# Extract the estimated parameters
```

```
priorPhiVariogram = krigingVariogramFittedMatrén1.0$cov.pars[1]
priorTauSQVariogram = krigingVariogramFittedMatrén1.0$cov.pars[2]
```

```
priorPhiMaximumLikelihood = Matren1.0linearREMLmaximumLikelihoodNetherlandsInitial10.1$cov
priorTauSQMaximumLikelihood = Matren1.0linearREMLmaximumLikelihoodNetherlandsInitial10.1$cov
```

The function does not support continuous priors directly so we will fit them as discrete priors.

```
# Creating discrete priors for phi and tau^2_rel AKA this is for the
# model with a nugget
phiDiscrete <- seq(min(priorPhiVariogram, priorPhiMaximumLikelihood) * 0.5,
  max(priorPhiVariogram, priorPhiMaximumLikelihood) * 1.5, length.out = 50)

tauSqDiscrete <- seq(min(priorTauSQVariogram, priorTauSQMaximumLikelihood) *
```



```
krigeBayesModelWithoutNugget <- krige.bayes(geodata = precipitationNetherlands_geoR,
      loc = ex.grid, prior = prior.control(phi.prior = phiProbability, phi.discrete = phiDis
```

krige.bayes: model with constant mean

krige.bayes: computing the discrete posterior of phi/tausq.rel

krige.bayes: computing the posterior probabilities.

Number of parameter sets: 50

1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26

krige.bayes: sampling from posterior distribution

krige.bayes: sample from the (joint) posterior of phi and tausq.rel

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	
phi	80.09036	226.3799	372.6695	518.9591	665.2486	811.5382	957.8278	
tausq.rel	0.00000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
frequency	16.00000	20.0000	14.0000	31.0000	23.0000	30.0000	23.0000	
	[,8]	[,9]	[,10]	[,11]	[,12]	[,13]	[,14]	
phi	1104.117	1250.407	1396.696	1542.986	1689.276	1835.565	1981.855	
tausq.rel	0.000	0.000	0.000	0.000	0.000	0.000	0.000	
frequency	24.000	21.000	28.000	21.000	27.000	28.000	31.000	
	[,15]	[,16]	[,17]	[,18]	[,19]	[,20]	[,21]	
phi	2128.144	2274.434	2420.723	2567.013	2713.303	2859.592	3005.882	
tausq.rel	0.000	0.000	0.000	0.000	0.000	0.000	0.000	
frequency	46.000	32.000	37.000	28.000	26.000	26.000	28.000	
	[,22]	[,23]	[,24]	[,25]	[,26]	[,27]	[,28]	[,29]
phi	3152.171	3298.461	3444.75	3591.04	3737.33	3883.619	4029.909	4176.198
tausq.rel	0.000	0.000	0.00	0.00	0.00	0.000	0.000	0.000
frequency	30.000	37.000	20.00	18.00	23.00	14.000	24.000	20.000
	[,30]	[,31]	[,32]	[,33]	[,34]	[,35]	[,36]	
phi	4322.488	4468.777	4615.067	4761.357	4907.646	5053.936	5200.225	
tausq.rel	0.000	0.000	0.000	0.000	0.000	0.000	0.000	
frequency	21.000	19.000	16.000	19.000	19.000	20.000	20.000	
	[,37]	[,38]	[,39]	[,40]	[,41]	[,42]	[,43]	
phi	5346.515	5492.804	5639.094	5785.384	5931.673	6077.963	6224.252	
tausq.rel	0.000	0.000	0.000	0.000	0.000	0.000	0.000	
frequency	17.000	14.000	8.000	17.000	5.000	12.000	9.000	
	[,44]	[,45]	[,46]	[,47]	[,48]	[,49]	[,50]	
phi	6370.542	6516.831	6663.121	6809.411	6955.7	7101.99	7248.279	
tausq.rel	0.000	0.000	0.000	0.000	0.0	0.00	0.000	
frequency	9.000	8.000	7.000	5.000	2.0	3.00	4.000	

krige.bayes: starting prediction at the provided locations

krige.bayes: phi/tausq.rel samples for the predictive are same as for the posterior

```
krige.bayes: computing moments of the predictive distribution
krige.bayes: sampling from the predictive
      Number of parameter sets: 50
1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26
krige.bayes: preparing summaries of the predictive distribution
```

```
summary(krigeBayesModelWithoutNugget$posterior$sample)
```

beta	sigmasq	phi	tausq.rel
Min. : -14806.5	Min. : 375543	Min. : 80.09	Min. : 0
1st Qu.: -2113.0	1st Qu.: 8331936	1st Qu.: 1542.99	1st Qu.: 0
Median : 290.8	Median : 14519487	Median : 2713.30	Median : 0
Mean : 299.1	Mean : 15700517	Mean : 2965.95	Mean : 0
3rd Qu.: 2547.6	3rd Qu.: 22223235	3rd Qu.: 4322.49	3rd Qu.: 0
Max. : 16575.4	Max. : 44053886	Max. : 7248.28	Max. : 0

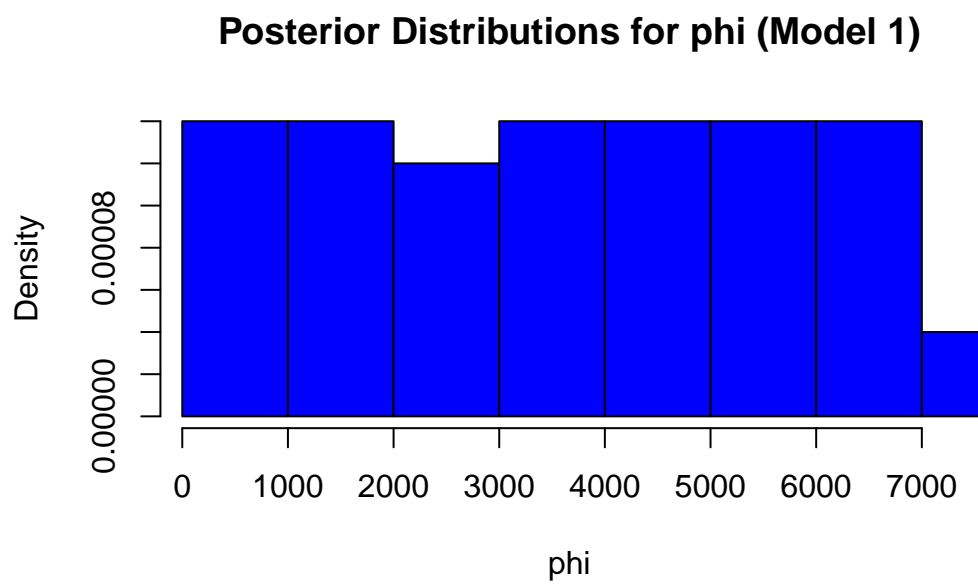
```
summary(krigeBayesModelWithNugget$posterior$sample)
```

beta	sigmasq	phi	tausq.rel
Min. : -49.52	Min. : 2216	Min. : 80.09	Min. : 0.1044
1st Qu.: 66.57	1st Qu.: 3000	1st Qu.: 80.09	1st Qu.: 0.1044
Median : 106.16	Median : 3207	Median : 80.09	Median : 0.1044
Mean : 105.66	Mean : 3238	Mean : 80.09	Mean : 0.1044
3rd Qu.: 145.57	3rd Qu.: 3433	3rd Qu.: 80.09	3rd Qu.: 0.1044
Max. : 282.96	Max. : 4455	Max. : 80.09	Max. : 0.1044

Now we will compare the posterior of both of the models to see the impact of the nugget

```
# Extract posterior samples for phi and tau^2_rel
posterior_samples_model1 <- krigeBayesModelWithNugget$posterior$phi$phi.marginal$phi
posterior_samples_model2 <- krigeBayesModelWithoutNugget$posterior$phi$phi.marginal$phi

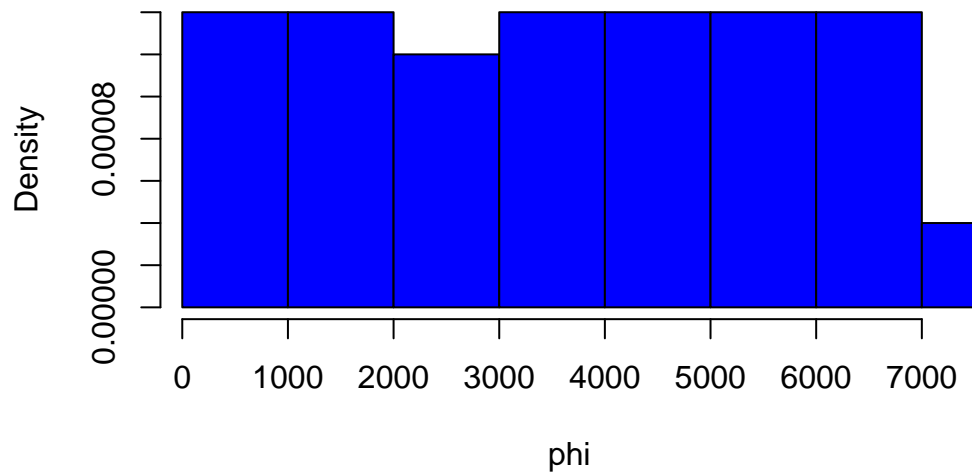
# Plot the posterior distributions
hist(posterior_samples_model1, freq = FALSE, main = "Posterior Distributions for phi (Model 1)",
      xlab = "phi", col = "blue", xlim = range(c(posterior_samples_model1,
          posterior_samples_model2)))
```



```
hist(posterior_samples_model2, freq = FALSE, main = "Posterior Distributions for phi (Model 1)",  
      xlab = "phi", col = "blue", xlim = range(c(posterior_samples_model1,  
          posterior_samples_model2)))
```



## Posterior Distributions for phi (Model 2)



```
# Compare summary statistics
summary_model1 <- summary(posterior_samples_model1)
summary_model2 <- summary(posterior_samples_model2)

cat("model 1 :\n")
```

model 1 :

```
summary_model1
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
80.09	1872.14	3664.18	3664.18	5456.23	7248.28

```
cat("\n\n model 2 :\n")
```

model 2 :

```
summary_model2
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
80.09	1872.14	3664.18	3664.18	5456.23	7248.28

```
# Reset the plot layout  
par(mfrow = c(1, 1))
```

As we can see with a low number of binds we can't see any significant difference in the summaries or histogram between the models with and without a nugget

## Question 2

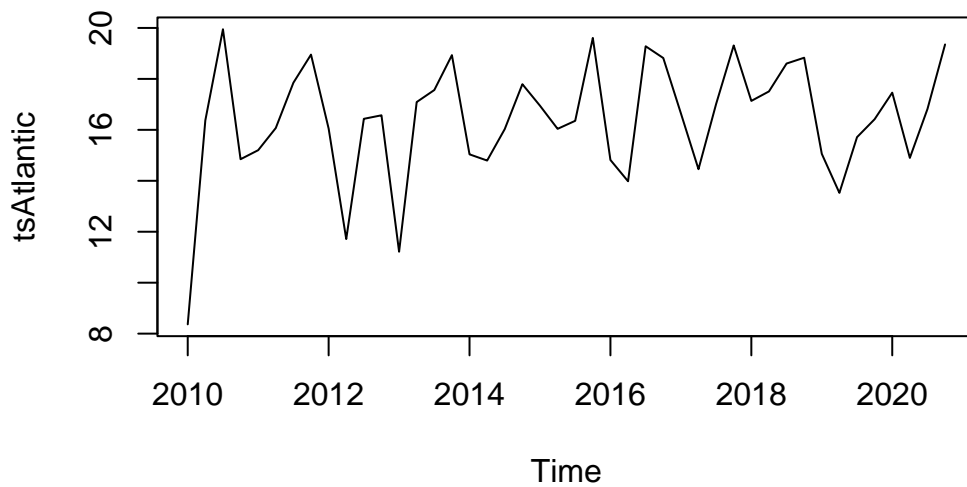
### 2 a)

We first start by making the appropriate changes in the data to average the data to quarterly means

```
AMOCDF$Date = as.Date(AMOCDF$Date, format = "%d/%m/%Y")  
  
## I will now make a column with the quarter and year that I will use  
## to create the averages per quarter  
AMOCDF$YearQuarter = paste(AMOCDF$Year, AMOCDF$Quarter, sep = "-")  
  
YearQuarterAverage = AMOCDF %>%  
  group_by(YearQuarter) %>%  
  summarise(AverageStrength = mean(Strength))
```

Now we will convert the average data to a time series object to be able to plot it

```
tsAtlantic = ts(YearQuarterAverage, start = c(2010, 1), frequency = 4)  
  
tsAtlantic = tsAtlantic[, "AverageStrength"]  
  
plot.ts(tsAtlantic)
```



### Trend analysis

From this graph we can see a yearly oscillation of Sverdrups. We can also identify that the peaks in Sverdrups are usually in the last quarter before the start of a new year and the valleys are on the second quarter of the year.

The data does seem stationary enough that if we were to differentiate we would start losing some of the structure.

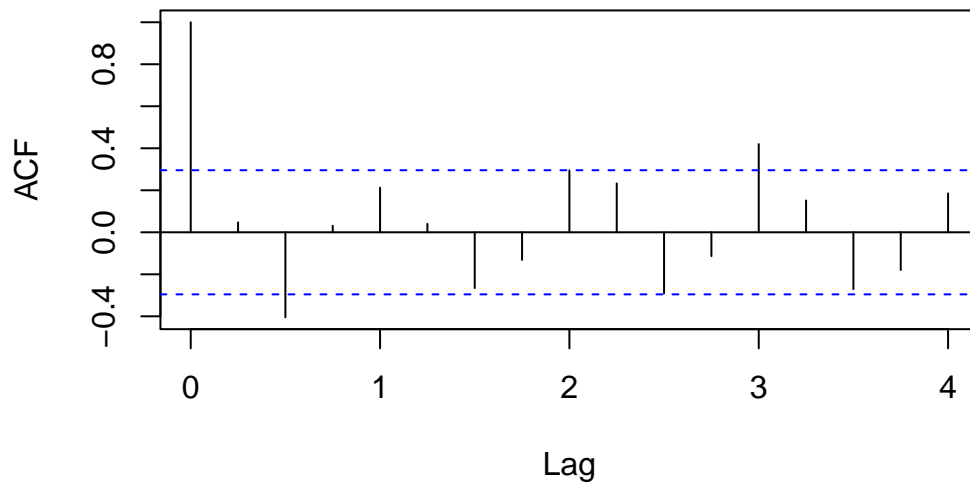
### 2 b)

#### ACF

First we will start by checking the ACF(Autocorrelation Function) and PACF(Partial Auto-correlation Function) to check for if we have stationary data or not to help us decide between an ARMA or an ARIMA model.

```
acf(tsAtlantic)
```

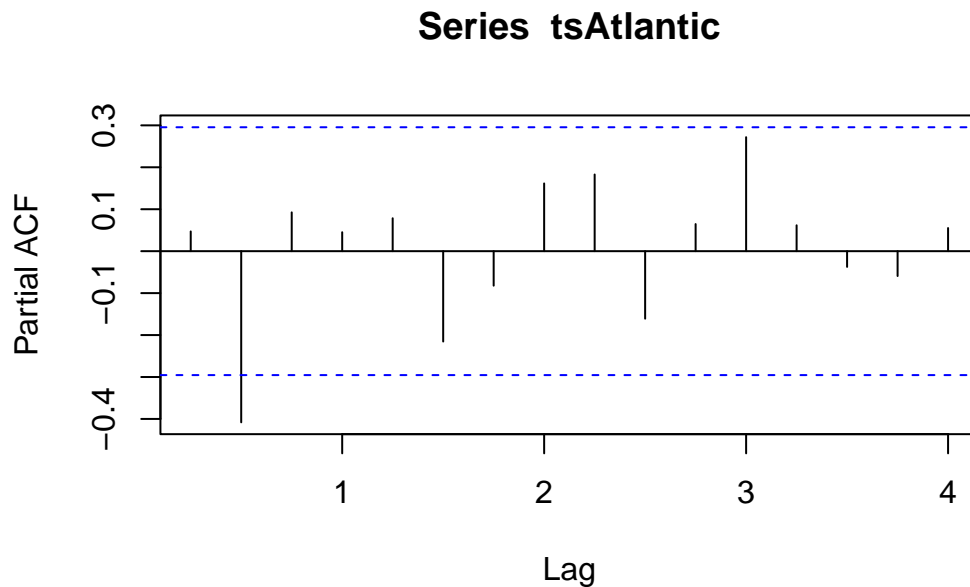
### Series tsAtlantic



We can see that for ACF OF Average strength slowly decreases as lag increases to infinity with lag = 3 still being a significant values, meaning it is not a simple MA model as AR is clearly not quickly cut-off.

### PACF

```
pacf(tsAtlantic)
```



The PACF seems to be cut-off at lag 0,5 indicating an AR model might be a best fit for our data to be a but with some almost significant values after the cut it might be also appropriate to some non-zero q values to confirm our initial assumption

As such we will now proceed to fit multiple model firstly with the initial assumption that, then I will both use models with non-zero q and the model given by the auto.arima function to double check that the assumptions made by the previous analyses is correct.

```
# it is always a good practice to try multiple values of p,d and q to
# see if we can do better we then obviously compare via the AIC of the
# models and their log likelihoods it is never enough to check those we
# also need to check the residuals

## order is p, d ,q

## initial models under our assumptions

model100 = Arima(tsAtlantic, order = c(1, 0, 0))
model200 = Arima(tsAtlantic, order = c(2, 0, 0))
model300 = Arima(tsAtlantic, order = c(3, 0, 0))

## now I will add postive q values
```

```

model101 = Arima(tsAtlantic, order = c(1, 0, 1))
model102 = Arima(tsAtlantic, order = c(1, 0, 2))
model103 = Arima(tsAtlantic, order = c(1, 0, 3))

model201 = Arima(tsAtlantic, order = c(2, 0, 1))
model202 = Arima(tsAtlantic, order = c(2, 0, 2))
model203 = Arima(tsAtlantic, order = c(2, 0, 3))

model301 = Arima(tsAtlantic, order = c(3, 0, 1))
model302 = Arima(tsAtlantic, order = c(3, 0, 2))
model303 = Arima(tsAtlantic, order = c(3, 0, 3))

## lastly we will use auto.arima without seasonality to confirm our
## initial assumptions
modelAuto = auto.arima(tsAtlantic, max.d = 0, max.p = 5, max.q = 5, seasonal = FALSE)

```

## best model selection

```
model100
```

```

Series: tsAtlantic
ARIMA(1,0,0) with non-zero mean

Coefficients:
      ar1      mean
    0.0665  16.3878
s.e.  0.1788   0.3726

sigma^2 = 5.572:  log likelihood = -99.2
AIC=204.41   AICc=205.01   BIC=209.76

```

```
model200
```

```

Series: tsAtlantic
ARIMA(2,0,0) with non-zero mean

Coefficients:

```

```

          ar1      ar2      mean
    0.0990 -0.5565  16.4298
s.e.  0.1576   0.1488   0.2113

sigma^2 = 4.321:  log likelihood = -93.45
AIC=194.9   AICc=195.92   BIC=202.04

```

```
model300
```

```

Series: tsAtlantic
ARIMA(3,0,0) with non-zero mean

```

```

Coefficients:
          ar1      ar2      ar3      mean
    0.1626 -0.5690  0.1464  16.4227
s.e.  0.1729   0.1479  0.1708   0.2409

sigma^2 = 4.35:  log likelihood = -93.09
AIC=196.17   AICc=197.75   BIC=205.1

```

As we can see from these initial models ARIMA(2,0,0) is the model that has the best fit has we can see from its lower AIC score of 194,9.

Now we will check against the other models to check the validity of our assumptions.

```
model101
```

```

Series: tsAtlantic
ARIMA(1,0,1) with non-zero mean

```

```

Coefficients:
          ar1      ma1      mean
    -0.4204  0.7718  16.3721
s.e.  0.2390  0.1466   0.4067

sigma^2 = 5.045:  log likelihood = -96.64
AIC=201.29   AICc=202.31   BIC=208.43

```

```
model102
```

Series: tsAtlantic  
ARIMA(1,0,2) with non-zero mean

Coefficients:

	ar1	ma1	ma2	mean
	0.0230	0.1275	-0.4485	16.4289
s.e.	0.3051	0.2420	0.1348	0.2224

sigma<sup>2</sup> = 4.651: log likelihood = -94.41  
AIC=198.81 AICc=200.39 BIC=207.73

model103

Series: tsAtlantic  
ARIMA(1,0,3) with non-zero mean

Coefficients:

	ar1	ma1	ma2	ma3	mean
	-0.5284	0.7214	-0.3646	-0.3072	16.4299
s.e.	0.9228	0.8649	0.2077	0.3545	0.2195

sigma<sup>2</sup> = 4.733: log likelihood = -94.25  
AIC=200.5 AICc=202.77 BIC=211.21

model201

Series: tsAtlantic  
ARIMA(2,0,1) with non-zero mean

Coefficients:

	ar1	ar2	ma1	mean
	-0.0669	-0.5475	0.2187	16.4255
s.e.	0.2740	0.1555	0.2883	0.2300

sigma<sup>2</sup> = 4.366: log likelihood = -93.15  
AIC=196.31 AICc=197.88 BIC=205.23

model202



Series: tsAtlantic  
ARIMA(2,0,2) with non-zero mean

Coefficients:

	ar1	ar2	ma1	ma2	mean
	0.0787	-0.9982	-0.0255	0.9999	16.4015
s.e.	0.0285	0.0066	0.0899	0.1158	0.2684

sigma<sup>2</sup> = 3.378: log likelihood = -89.46  
AIC=190.91 AICc=193.18 BIC=201.62

model203

Series: tsAtlantic  
ARIMA(2,0,3) with non-zero mean

Coefficients:

	ar1	ar2	ma1	ma2	ma3	mean
	0.0325	-0.9621	0.0499	0.8487	0.4147	16.4028
s.e.	0.0645	0.0442	0.1987	0.2041	0.2511	0.3044

sigma<sup>2</sup> = 3.315: log likelihood = -89.07  
AIC=192.13 AICc=195.25 BIC=204.62

model301

Series: tsAtlantic  
ARIMA(3,0,1) with non-zero mean

Coefficients:

	ar1	ar2	ar3	ma1	mean
	0.4092	-0.5931	0.2864	-0.2467	16.4191
s.e.	0.6330	0.1651	0.3580	0.6291	0.2537

sigma<sup>2</sup> = 4.449: log likelihood = -93.03  
AIC=198.06 AICc=200.34 BIC=208.77

model302

```
Series: tsAtlantic
ARIMA(3,0,2) with non-zero mean
```

Coefficients:

	ar1	ar2	ar3	ma1	ma2	mean
	0.2684	-0.9851	0.2222	-0.3030	1.0000	16.4144
s.e.	0.1999	0.0305	0.1995	0.1453	0.1921	0.2922

```
sigma^2 = 3.392: log likelihood = -89.53
AIC=193.06 AICc=196.17 BIC=205.54
```

```
model303
```

```
Series: tsAtlantic
ARIMA(3,0,3) with non-zero mean
```

Coefficients:

	ar1	ar2	ar3	ma1	ma2	ma3	mean
	-0.3983	-0.9518	-0.4263	0.4381	0.7816	0.7352	16.4197
s.e.	0.3385	0.0412	0.3442	0.2946	0.1690	0.2422	0.2691

```
sigma^2 = 3.352: log likelihood = -88.54
AIC=193.08 AICc=197.19 BIC=207.35
```

In this initial analysis we have found models that do have a lower AIC lower log likelihood than our previous best model, however these model ma's standard error are to close the the ma values indicating that while we are getting a better fit we might be overfitting to our data.

As such this does confirm our initial assumption for the choice of a zero q value.

Now lastly we will check if the auto.arima function does confirm our initial assumptions.

```
modelAuto
```

```
Series: tsAtlantic
ARIMA(2,0,0) with non-zero mean
```

Coefficients:

	ar1	ar2	mean
	0.0990	-0.5565	16.4298
s.e.	0.1576	0.1488	0.2113

```
sigma^2 = 4.321: log likelihood = -93.45
AIC=194.9 AICc=195.92 BIC=202.04
```

The function does confirm our assumption that ARIMA(2,0,0) is indeed the best model.

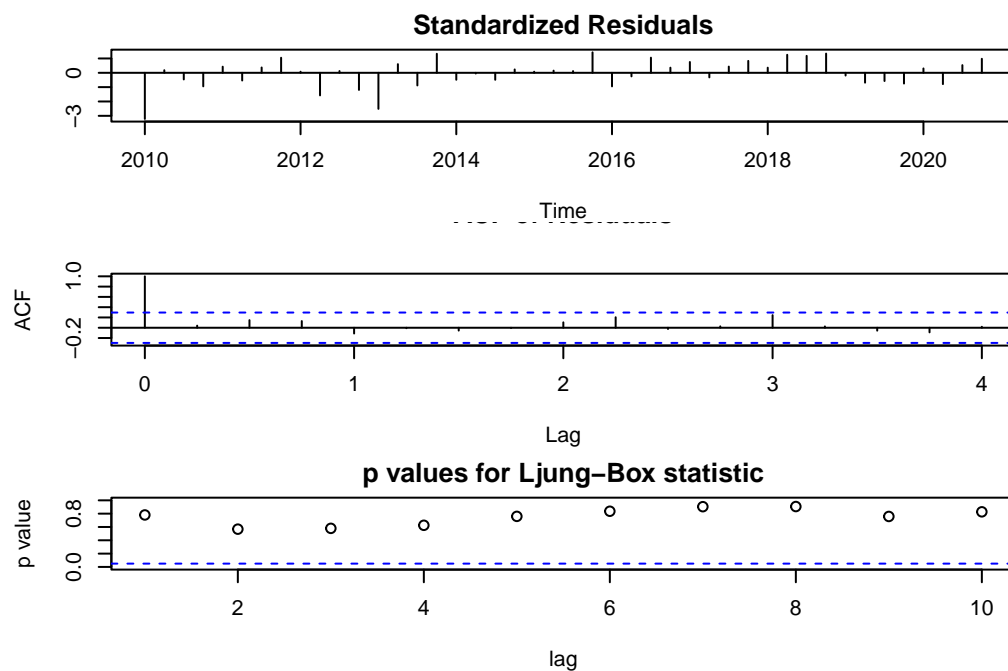
We will now check the residuals to verify if any of our previously selected model validates well or if it is simply the best of bad models.

**talk about the model being more easily explainability because MA = 0**

### Best model residual validation

```
# Set smaller margins
par(mar = c(4, 4, 2, 2))

tsdiag(model200)
```



```
# Reset margins
par(mar = c(5, 4, 4, 2) + 0.1)
```

Initially from the standardised residuals plot we can identify some sort of sinusoidal pattern, this implies that there is a seasonal trend that is not being accounted for in our model and

as such this trends needs to be accounted in future models to better explain and increase the prediction power of a new model.

## Forecasting

Now using the forecast function we will forecast the next 4 quarters of 2021

```
forecast(model200, 4)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2021 Q1	16.50240	13.83841	19.16639	12.42818	20.57662
2021 Q2	14.81104	12.13403	17.48804	10.71691	18.90517
2021 Q3	16.22919	13.18168	19.27669	11.56843	20.88994
2021 Q4	17.31076	14.24941	20.37212	12.62882	21.99271

But this data is better visualized in a graph to better understand if the predictions are sensible compared to our real data.

```
predictedArimaDF = data.frame(forecast(model200, 4))

predictedArimaDF$YearQuarter = c("2021-Q1", "2021-Q2", "2021-Q3", "2021-Q4")

# Combine real_data and pred_data into a single data frame
combinedDataframeAMOC = rbind(data.frame(Date = YearQuarterAverage$YearQuarter,
    Temperature = YearQuarterAverage$AverageStrength, Type = "Real"), data.frame(Date = predictedArimaDF$YearQuarter,
    Temperature = predictedArimaDF$Point.Forecast, Type = "Predicted"))

predictedArimaDF$Temperature = predictedArimaDF$Point.Forecast

predictedArimaDF$Type = "Predicted"

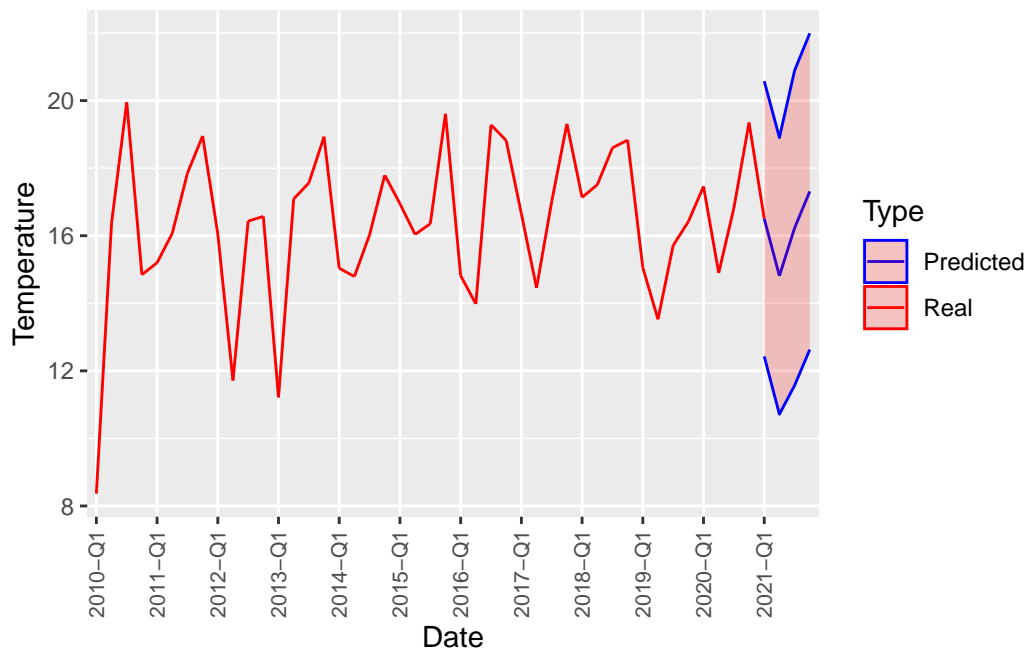
# Create the ggplot
plotARIMA = ggplot(combinedDataframeAMOC, aes(x = Date, y = Temperature,
    color = Type, group = 1)) + geom_line() + scale_color_manual(values = c("blue",
    "red"))

# Add the 95% confidence interval
plotARIMA = plotARIMA + geom_ribbon(data = predictedArimaDF, aes(x = YearQuarter,
    ymin = Lo.95, ymax = Hi.95), fill = "red", alpha = 0.2)
```

```
# Adjust the x-axis labels
plotARIMA = plotARIMA + scale_x_discrete(breaks = combinedDataframeAMOC$Date[c(TRUE,
  rep(FALSE, 3))], labels = combinedDataframeAMOC$Date[c(TRUE, rep(FALSE,
  3))])

plotARIMA = plotARIMA + theme(axis.text.x = element_text(angle = 90, vjust = 0.5,
  size = 8))

# Display the plot
print(plotARIMA)
```



As we can see from the graph the ARIMA (2,0,0) seems to give us a sensible forecast for the 2021 quarter values, however as we can see the interval of the prediction accuracy our model is not too certain on the values most likely due to our model not accounting for the seasonal cycle of our data.

## 2 c)

### Initial assumptions

From the previous exploratory analysis of the data we have established that the data did not

need to be differentiated since it was constant, this translates to polynomial DLM component of order 2 that will use linear model to account for this type of changes in the data.

Furthermore, from the residual analysis we have inferred that there is an underlying seasonal trend present on the data, this seasonal trend will be represented by a seasonal component of frequency 4 to represent the 4 quarters per year.

### model fitting

```
## linear model, order = 2, quadratic order = 3 , etc

## what we want is a linear model with a seasonal component so we add
## the 2 components together in a model

## things to try, another term like quadratic, or a arma component
## stacked on top of this

## Initial model with a linear polynomial and a seasonal component

buildFun = function(x) {
  dlmModPoly(order = 2, dV = exp(x[1]), dW = c(0, exp(x[2]))) + dlmModSeas(frequency = 4,
    dV = 0, dW = c(exp(x[3]), rep(0, 2)))
}

linearDLM = dlmMLE(tsAtlantic, parm = c(0, 0, 0), build = buildFun)

linearDLM$par
```

```
[1] 1.151339 -18.078101 -2.189479
```

```
fittedLinearDLM = buildFun(linearDLM$par)
```

```
V(fittedLinearDLM)
```

```
      [,1]
[1,] 3.162425
```

```
W(fittedLinearDLM)
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	0	0.000000e+00	0.000000	0	0
[2,]	0	1.408576e-08	0.000000	0	0
[3,]	0	0.000000e+00	0.111975	0	0
[4,]	0	0.000000e+00	0.000000	0	0
[5,]	0	0.000000e+00	0.000000	0	0

```
## second model with a quadratic polynomial and a seasonal component
```

```
buildFunQuad = function(x) {
  dlmModPoly(order = 3, dV = exp(x[1]), dW = c(0, exp(x[2]), exp(x[3]))) +
  dlmModSeas(frequency = 4, dV = 0, dW = c(exp(x[4]), rep(0, 2)))
}
```

```
quadraticDLM = dlmMLE(tsAtlantic, parm = c(0, 0, 0, 0), build = buildFunQuad)
```

```
quadraticDLM$par
```

```
[1] 1.161355 -17.807081 -28.603103 -2.352292
```

```
fittedQuadraticDLM = buildFunQuad(quadraticDLM$par)
```

```
V(fittedQuadraticDLM)
```

	[,1]
[1,]	3.194257

```
W(fittedQuadraticDLM)
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	0	0.000000e+00	0.000000e+00	0.00000000	0	0
[2,]	0	1.847069e-08	0.000000e+00	0.00000000	0	0
[3,]	0	0.000000e+00	3.782948e-13	0.00000000	0	0
[4,]	0	0.000000e+00	0.000000e+00	0.09515082	0	0
[5,]	0	0.000000e+00	0.000000e+00	0.00000000	0	0
[6,]	0	0.000000e+00	0.000000e+00	0.00000000	0	0

Now we will compare both models through their log likelihood using the `dlmLL` function and see if the extra flexibility from the extra polynomial function is providing a better fit

```
dmlLL(tsAtlantic, fittedLinearDLM)
```

```
[1] 94.98804
```

```
dmlLL(tsAtlantic, fittedQuadraticDLM)
```

```
[1] 108.043
```

As we can see the dlm model using only a linear polynomial has a lower log likelihood than the model with an extra quadratic term, meaning this extra flexibility does not contribute to a better model fit and as such we will use the linear fitted model to do our forecasting.

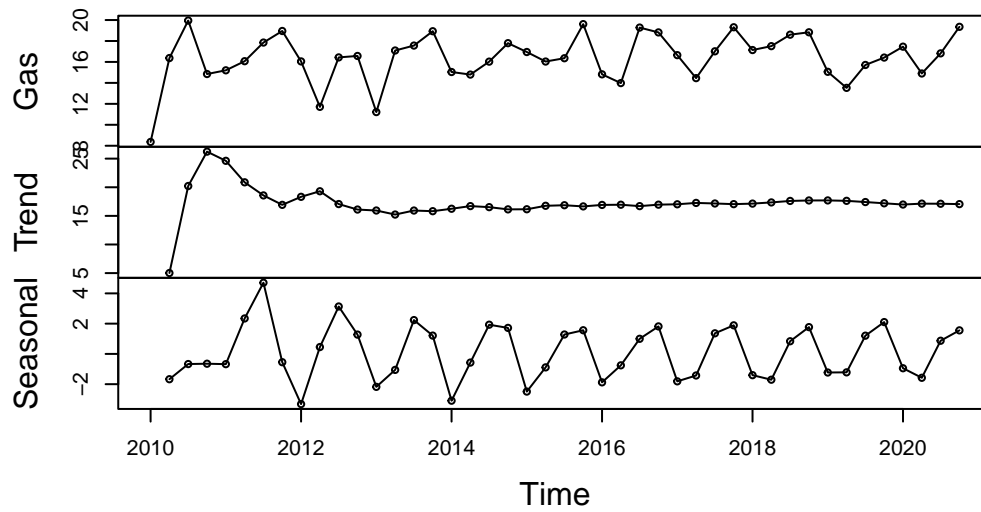
```
amocPredict = dlmFilter(tsAtlantic, mod = fittedLinearDLM)
summary(amocPredict)
```

	Length	Class	Mode
y	44	ts	numeric
mod	10	dlm	list
m	225	mts	numeric
U.C	45	-none-	list
D.C	225	-none-	numeric
a	220	mts	numeric
U.R	44	-none-	list
D.R	220	-none-	numeric
f	44	ts	numeric

```
x = cbind(tsAtlantic, dropFirst(amocPredict$a[, c(1, 3)]))
x = window(x, start = c(2010, 1))
colnames(x) = c("Gas", "Trend", "Seasonal")
plot(x, type = "o", main = "Atlantic AMOC at 26,5N 2010-2020")
```



## Atlantic AMOC at 26,5N 2010–2020



### Forecast

```
amocForecast = dlmForecast(amocPredict, nAhead = 4)
summary(amocForecast)
```

```

Length Class Mode
a 20      mts   numeric
R  4      -none- list
f  4      ts    numeric
Q  4      -none- list
```

```
dim(amocForecast$a)
```

```
[1] 4 5
```

```
dim(amocForecast$f)
```

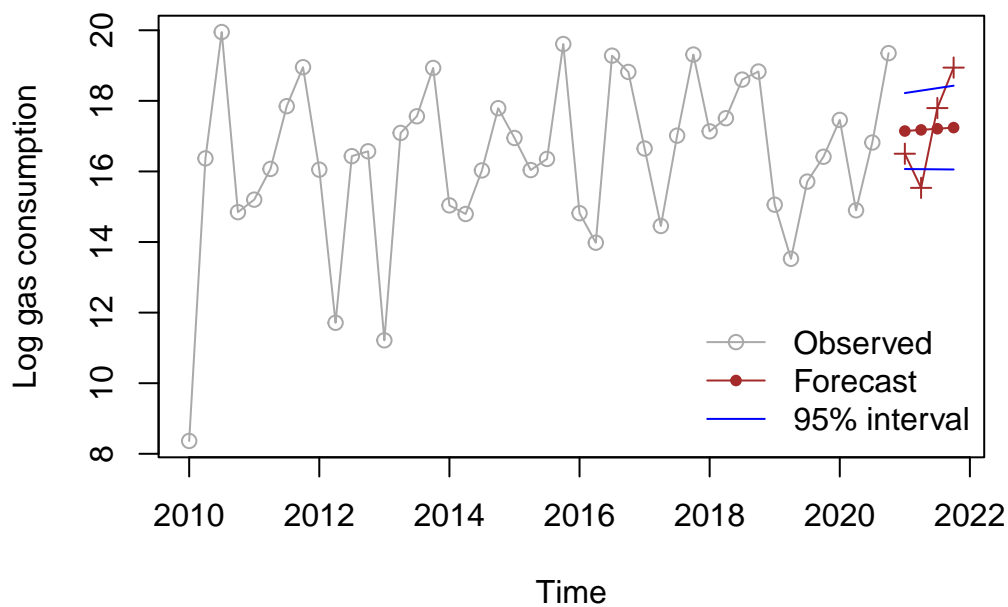
```
[1] 4 1
```

```

sqrtR = sapply(amocForecast$R, function(x) sqrt(x[1, 1]))
pl = amocForecast$a[, 1] + qnorm(0.025, sd = sqrtR)
pu = amocForecast$a[, 1] + qnorm(0.975, sd = sqrtR)
x = ts.union(window(tsAtlantic, start = c(2010, 1)), amocForecast$a[, 1],
              amocForecast$f, pl, pu)
par(mar = c(4, 4, 2, 2))
plot(x, plot.type = "single", type = "o", pch = c(1, 20, 3, NA, NA), col = c("darkgrey",
  "brown", "brown", "blue", "blue"), ylab = "Log gas consumption")

legend("bottomright", legend = c("Observed", "Forecast", "95% interval"),
      bty = "n", pch = c(1, 20, NA), lty = 1, col = c("darkgrey", "brown",
  "blue"))

```

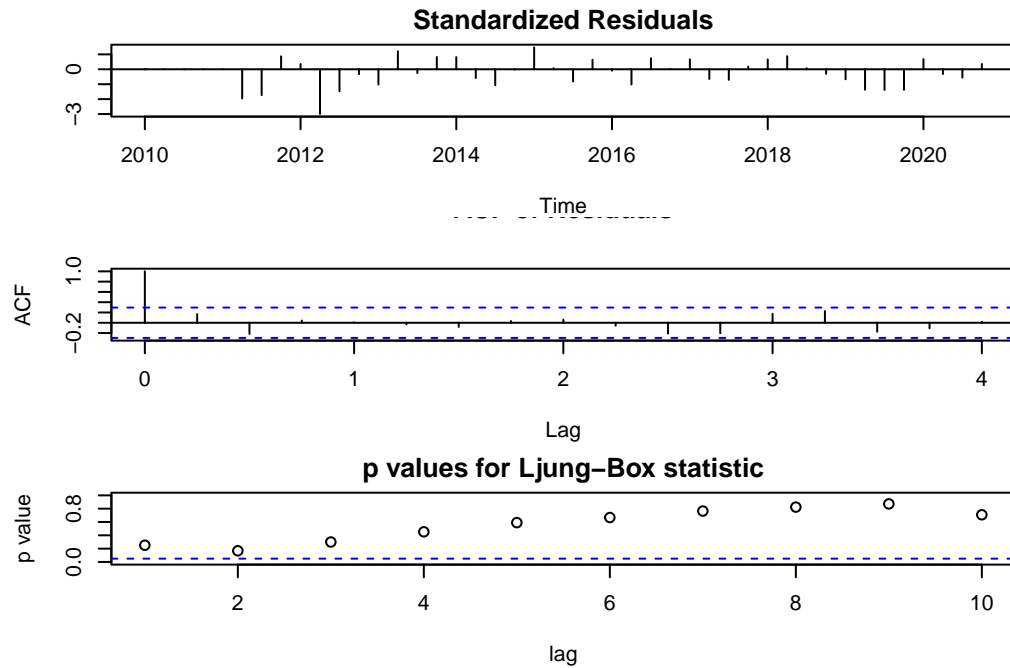


```

# Set smaller margins
par(mar = c(4, 4, 2, 2))

tsdiag(amocPredict)

```

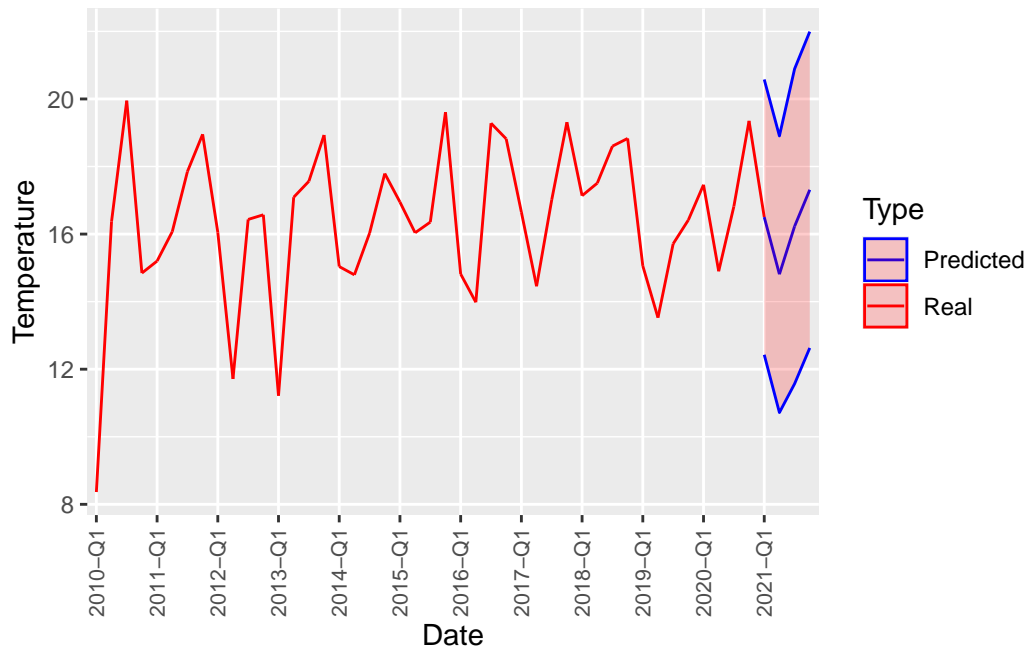


```
# Reset margins
par(mar = c(5, 4, 4, 2) + 0.1)
```

## 2 d)

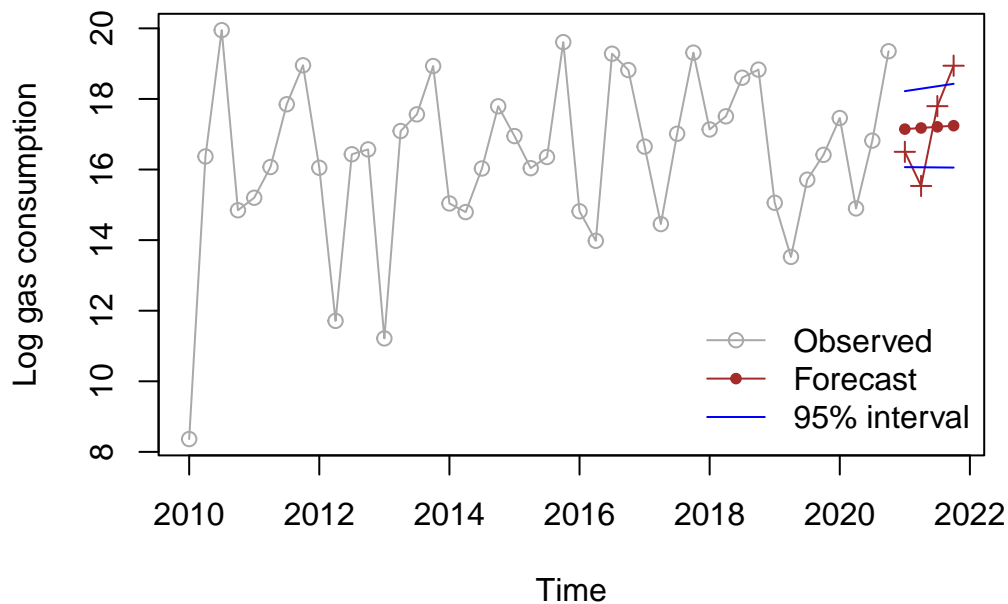
Again comparing the forecast values and their respective prediction intervals as we can see from the graphs below the dlm model has smaller prediction intervals, most likely due to being able to explain the underlying seasonal trend reducing therefore the uncertainty in comparison the ARIMA model.

```
print(plotARIMA)
```



```
sqaretRoot = sapply(amocForecast$R, function(x) sqrt(x[1, 1]))
predictionLow = amocForecast$a[, 1] + qnorm(0.025, sd = sqaretRoot) ## Low
predictionUpper = amocForecast$a[, 1] + qnorm(0.975, sd = sqaretRoot) ## Upper
x = ts.union(window(tsAtlantic, start = c(2010, 1)), amocForecast$a[, 1],
              amocForecast$f, predictionLow, predictionUpper)
par(mar = c(4, 4, 2, 2))
plot(x, plot.type = "single", type = "o", pch = c(1, 20, 3, NA, NA), col = c("darkgrey",
  "brown", "brown", "blue", "blue"), ylab = "Log gas consumption")

legend("bottomright", legend = c("Observed", "Forecast", "95% interval"),
      bty = "n", pch = c(1, 20, NA), lty = 1, col = c("darkgrey", "brown",
  "blue"))
```



2 e)

```
# AMOCDFMonthly =AMOCDF %>% mutate(YearMonth = paste0(year(Date), '-',
# month(Date, label = TRUE, abbr = FALSE)))
```

```
## I will now make a column with the month and year that I will use to
## create the monthly averages
```

```
AMOCDF$YearMonth = paste(AMOCDF$Year, AMOCDF$Month, sep = "-")
```

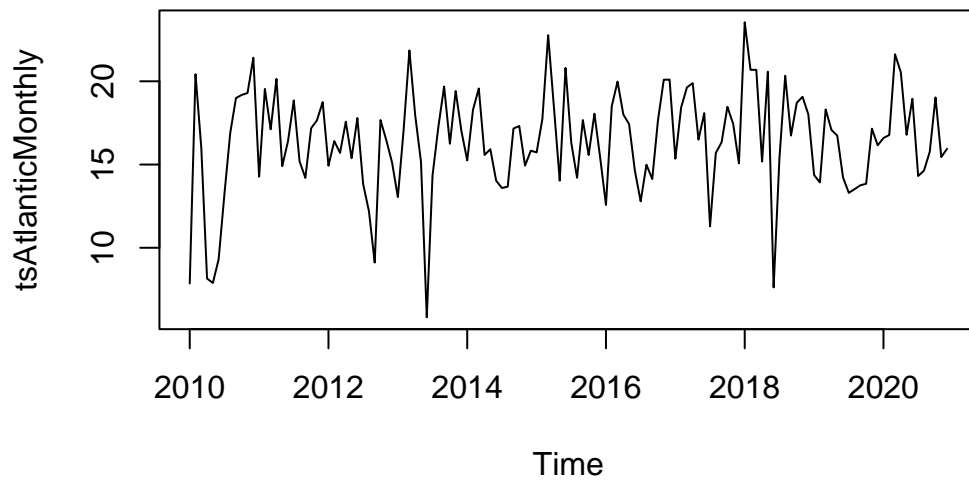
```
YearMonthlyAverage = AMOCDF %>%
  group_by(YearMonth) %>%
  summarise(AverageStrength = mean(Strength))
```

Now we will create a new montly time series object and make it univariate

```
tsAtlanticMonthly = ts(YearMonthlyAverage, start = c(2010, 1), frequency = 12)

tsAtlanticMonthly = tsAtlanticMonthly[, "AverageStrength"]

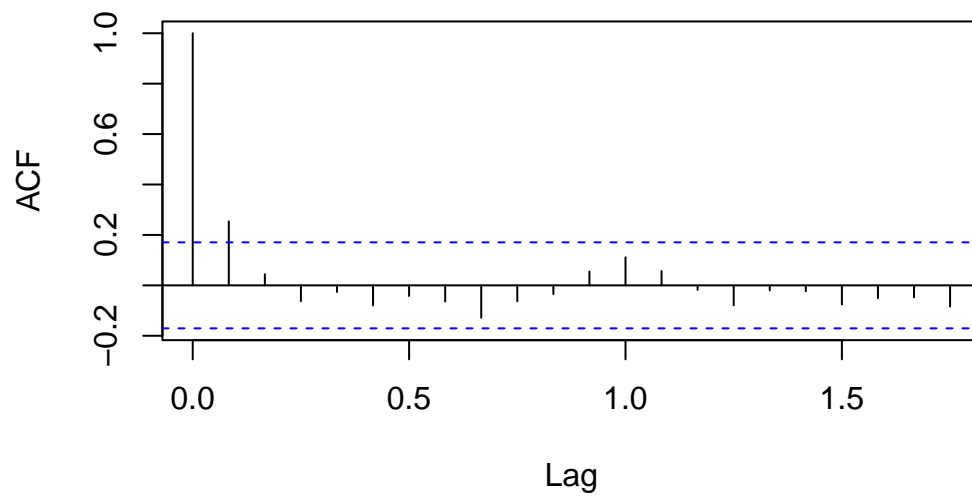
plot.ts(tsAtlanticMonthly)
```



Seeing this graph we can observe that the data continues being stationary for the ARIMA model but a seasonal trend not only is more apparently but it also appear to need to be differentiated as it seems to have a decreasing linear trend

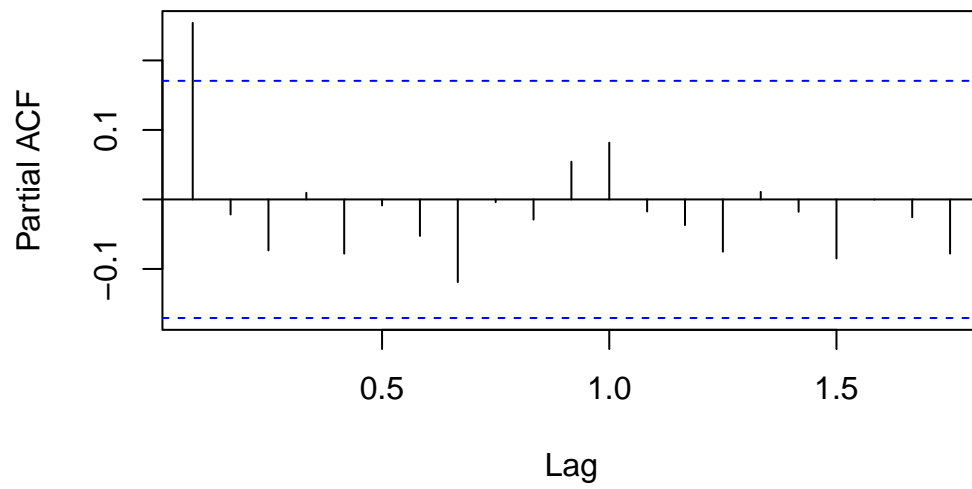
```
acf(tsAtlanticMonthly)
```

**Series tsAtlanticMonthly**



```
pacf(tsAtlanticMonthly)
```

**Series tsAtlanticMonthly**



The acf has a very clear cut-off as only 3 the values are significant which is very similar to what we had observed previously.

The main difference is in the pacf, where we can now say for sure that there is a very clear cut-off after the first value.

## Model testing

These pattern suggests that an ARMA/ARIMA model might be the most appropriate so first we will check them out with the seasonal component of order 1, the so quick cut-off of both the ACF and PACF also might suggest that p and q will be smaller values.

## Seasonal check

```
## initial assumption

modelMonthlySeasonal100.110 = Arima(tsAtlanticMonthly, order = c(1, 0, 0),
  seasonal = list(order = c(1, 1, 0), period = 12))

modelMonthlySeasonal100.011 = Arima(tsAtlanticMonthly, order = c(1, 0, 0),
  seasonal = list(order = c(0, 1, 1), period = 12))

modelMonthlySeasonal200.210 = Arima(tsAtlanticMonthly, order = c(2, 0, 0),
  seasonal = list(order = c(2, 1, 0), period = 12))
modelMonthlySeasonal200.012 = Arima(tsAtlanticMonthly, order = c(2, 0, 0),
  seasonal = list(order = c(0, 1, 2), period = 12))

modelMonthlySeasonal001.110 = Arima(tsAtlanticMonthly, order = c(0, 0, 1),
  seasonal = list(order = c(1, 1, 0), period = 12))

modelMonthlySeasonal001.011 = Arima(tsAtlanticMonthly, order = c(0, 0, 1),
  seasonal = list(order = c(0, 1, 1), period = 12))

modelMonthlySeasonal002.210 = Arima(tsAtlanticMonthly, order = c(0, 0, 2),
  seasonal = list(order = c(2, 1, 0), period = 12))
modelMonthlySeasonal002.012 = Arima(tsAtlanticMonthly, order = c(0, 0, 2),
  seasonal = list(order = c(0, 1, 2), period = 12))

modelMonthlySeasonal100.110
```



Series: tsAtlanticMonthly  
ARIMA(1,0,0)(1,1,0)[12]

Coefficients:

	ar1	sar1
	0.1779	-0.4618
s.e.	0.0909	0.0839

sigma^2 = 12.06: log likelihood = -320.1  
AIC=646.2 AICc=646.4 BIC=654.56

modelMonthlySeasonal100.011

Series: tsAtlanticMonthly  
ARIMA(1,0,0)(0,1,1)[12]

Coefficients:

	ar1	sma1
	0.1844	-0.8500
s.e.	0.0918	0.1123

sigma^2 = 8.74: log likelihood = -306.88  
AIC=619.76 AICc=619.97 BIC=628.12

modelMonthlySeasonal200.210

Series: tsAtlanticMonthly  
ARIMA(2,0,0)(2,1,0)[12]

Coefficients:

	ar1	ar2	sar1	sar2
	0.115	0.0374	-0.7271	-0.4814
s.e.	0.094	0.0930	0.0932	0.0933

sigma^2 = 9.741: log likelihood = -309.65  
AIC=629.29 AICc=629.82 BIC=643.23

modelMonthlySeasonal200.012

```
Series: tsAtlanticMonthly  
ARIMA(2,0,0)(0,1,2)[12]
```

```
Coefficients:
```

	ar1	ar2	sma1	sma2
	0.1566	0.0546	-0.9817	0.1895
s.e.	0.0947	0.0950	0.1384	0.1503

```
sigma^2 = 8.778: log likelihood = -305.88  
AIC=621.76 AICc=622.29 BIC=635.7
```

```
modelMonthlySeasonal001.110
```

```
Series: tsAtlanticMonthly  
ARIMA(0,0,1)(1,1,0)[12]
```

```
Coefficients:
```

	ma1	sar1
	0.1680	-0.4634
s.e.	0.0861	0.0840

```
sigma^2 = 12.07: log likelihood = -320.19  
AIC=646.39 AICc=646.59 BIC=654.75
```

```
modelMonthlySeasonal001.011
```

```
Series: tsAtlanticMonthly  
ARIMA(0,0,1)(0,1,1)[12]
```

```
Coefficients:
```

	ma1	sma1
	0.1606	-0.8446
s.e.	0.0847	0.1093

```
sigma^2 = 8.804: log likelihood = -307.14  
AIC=620.28 AICc=620.49 BIC=628.65
```

```
modelMonthlySeasonal002.210
```

```
Series: tsAtlanticMonthly
ARIMA(0,0,2)(2,1,0)[12]
```

```
Coefficients:
```

	ma1	ma2	sar1	sar2
	0.1145	0.0452	-0.7275	-0.4806
s.e.	0.0943	0.0882	0.0933	0.0936

```
sigma^2 = 9.745: log likelihood = -309.66
AIC=629.33 AICc=629.85 BIC=643.26
```

```
modelMonthlySeasonal002.012
```

```
Series: tsAtlanticMonthly
ARIMA(0,0,2)(0,1,2)[12]
```

```
Coefficients:
```

	ma1	ma2	sma1	sma2
	0.1583	0.0745	-0.9786	0.1868
s.e.	0.0953	0.0893	0.1377	0.1495

```
sigma^2 = 8.784: log likelihood = -305.89
AIC=621.78 AICc=622.3 BIC=635.72
```

So as suspected from both the time series plot and the last exercise analysis the added seasonality does increase our model goodness of fit while also penalising the increased in complexity with so far.

Now lets compare them to bigger p and q values to see if our initial assumptions do hold up

```
modelMonthlySeasonal301 = Arima(tsAtlanticMonthly, order = c(3, 0, 1), seasonal = list(ord
  1, 1), period = 12))
modelMonthlySeasonal302 = Arima(tsAtlanticMonthly, order = c(3, 0, 2), seasonal = list(ord
  1, 0), period = 12))
modelMonthlySeasonal303 = Arima(tsAtlanticMonthly, order = c(3, 1, 3), seasonal = list(ord
  1, 0), period = 12))

modelMonthlySeasonal103 = Arima(tsAtlanticMonthly, order = c(1, 0, 3), seasonal = list(ord
  1, 1), period = 12))
modelMonthlySeasonal203 = Arima(tsAtlanticMonthly, order = c(2, 1, 3), seasonal = list(ord
  1, 0), period = 12))
```

### modelMonthlySeasonal301

Series: tsAtlanticMonthly  
ARIMA(3,0,1)(1,1,1)[12]

Coefficients:

	ar1	ar2	ar3	ma1	sar1	sma1
	-0.5728	0.1878	-0.0048	0.7442	-0.1158	-0.8073
s.e.	0.7595	0.1720	0.1369	0.7564	0.1224	0.1140

$\sigma^2 = 8.939$ : log likelihood = -306.02  
AIC=626.04 AICc=627.04 BIC=645.55

### modelMonthlySeasonal302

Series: tsAtlanticMonthly  
ARIMA(3,0,2)(1,1,0)[12]

Coefficients:

	ar1	ar2	ar3	ma1	ma2	sar1
	-1.4367	-0.5902	0.1291	1.6983	1.0000	-0.4382
s.e.	0.0944	0.1532	0.0945	0.0436	0.0489	0.0871

$\sigma^2 = 11.34$ : log likelihood = -316.59  
AIC=647.19 AICc=648.19 BIC=666.7

### modelMonthlySeasonal303

Series: tsAtlanticMonthly  
ARIMA(3,1,3)(1,1,0)[12]

Coefficients:

	ar1	ar2	ar3	ma1	ma2	ma3	sar1
	-1.4279	-0.5751	0.1373	0.6984	-0.6984	-1.0000	-0.4313
s.e.	0.0952	0.1547	0.0954	0.0533	0.0552	0.0576	0.0877

$\sigma^2 = 11.54$ : log likelihood = -316.86  
AIC=649.73 AICc=651.04 BIC=671.96

```
modelMonthlySeasonal103
```

```
Series: tsAtlanticMonthly  
ARIMA(1,0,3)(1,1,1)[12]
```

```
Coefficients:
```

	ar1	ma1	ma2	ma3	sar1	sma1
	-0.6951	0.8708	0.1931	-0.0013	-0.1134	-0.8029
s.e.	0.4854	0.4812	0.1503	0.1321	0.1194	0.1136

```
sigma^2 = 8.954: log likelihood = -306.02  
AIC=626.03 AICc=627.03 BIC=645.55
```

```
modelMonthlySeasonal203
```

```
Series: tsAtlanticMonthly  
ARIMA(2,1,3)(1,1,0)[12]
```

```
Coefficients:
```

	ar1	ar2	ma1	ma2	ma3	sar1
	-1.3087	-0.6017	0.5001	-0.6702	-0.8299	-0.4545
s.e.	0.2576	0.1557	0.2045	0.1332	0.1180	0.0979

```
sigma^2 = 12.02: log likelihood = -318.1  
AIC=650.19 AICc=651.2 BIC=669.65
```

As we can see here the initial assumption that a smaller p and q value would better fit the model.

Now we will use auto.arima to verify if our assumptions were indeed correct

### auto arima check

```
`?`(auto.arima)
```

```
starting httpd help server ... done
```

```
monthlyModelAuto = auto.arima(tsAtlanticMonthly, max.d = 0, max.p = 5, max.q = 5,
                               D = 1)
monthlyModelAuto
```

Series: tsAtlanticMonthly  
ARIMA(1,0,0)(0,1,1)[12]

Coefficients:

```
          ar1      sma1
      0.1844  -0.8500
s.e.  0.0918   0.1123
```

```
sigma^2 = 8.74:  log likelihood = -306.88
AIC=619.76   AICc=619.97   BIC=628.12
```

From what we can see the auto arima has indeed confirmed our initial assumption by picking a model that we already had seen as the best performer ARIMA(1,0,0)(0,1,1)[12]

## Forecasting

Now using the forecast function we will forecast the next 4 quarters of 2021

```
forecast(modelMonthlySeasonal100.011, 12)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Jan 2021	15.29474	11.49064	19.09884	9.476872	21.11261
Feb 2021	17.68442	13.81636	21.55247	11.768741	23.60009
Mar 2021	19.72032	15.85011	23.59053	13.801347	25.63929
Apr 2021	17.53922	13.66894	21.40951	11.620140	23.45831
May 2021	16.17383	12.30355	20.04411	10.254741	22.09292
Jun 2021	14.65775	10.78746	18.52803	8.738657	20.57683
Jul 2021	14.12800	10.25772	17.99829	8.208914	20.04709
Aug 2021	15.37476	11.50447	19.24504	9.455670	21.29385
Sep 2021	15.71081	11.84052	19.58109	9.791718	21.62989
Oct 2021	17.33660	13.46632	21.20689	11.417514	23.25569
Nov 2021	17.56114	13.69086	21.43141	11.642055	23.48022
Dec 2021	16.87677	13.00663	20.74691	10.957906	22.79563

But this data is better visualized in a graph to better understand if the predictions are sensible compared to our real data.

```

predictedArimaSeasonalDF = data.frame(forecast(modelMonthlySeasonal100.011,
12))

predictedArimaSeasonalDF$YearMonth = c("2021-1", "2021-2", "2021-3", "2021-4",
"2021-5", "2021-6", "2021-7", "2021-8", "2021-9", "2021-10", "2021-11",
"2021-12")

predictedArimaSeasonalDF$Type = "Predicted"

predictedArimaSeasonalDF$Temperature = predictedArimaSeasonalDF$Point.Forecast

YearMonthlyAverage$Type = "Real"
YearMonthlyAverage$Temperature = YearMonthlyAverage$AverageStrength

# Combine real_data and pred_data into a single data frame
combinedDataframeAMOC = rbind(data.frame(Date = YearMonthlyAverage$YearMonth,
Temperature = YearMonthlyAverage$Temperature, Type = "Real"), data.frame(Date = predictedArimaSeasonalDF$YearMonth,
Temperature = predictedArimaSeasonalDF$Temperature, Type = "Predicted"))

# Create the ggplot
plotARIMA2 = ggplot(combinedDataframeAMOC, aes(x = Date, y = Temperature,
color = Type, group = 1)) + geom_line() + scale_color_manual(values = c("blue",
"red"))

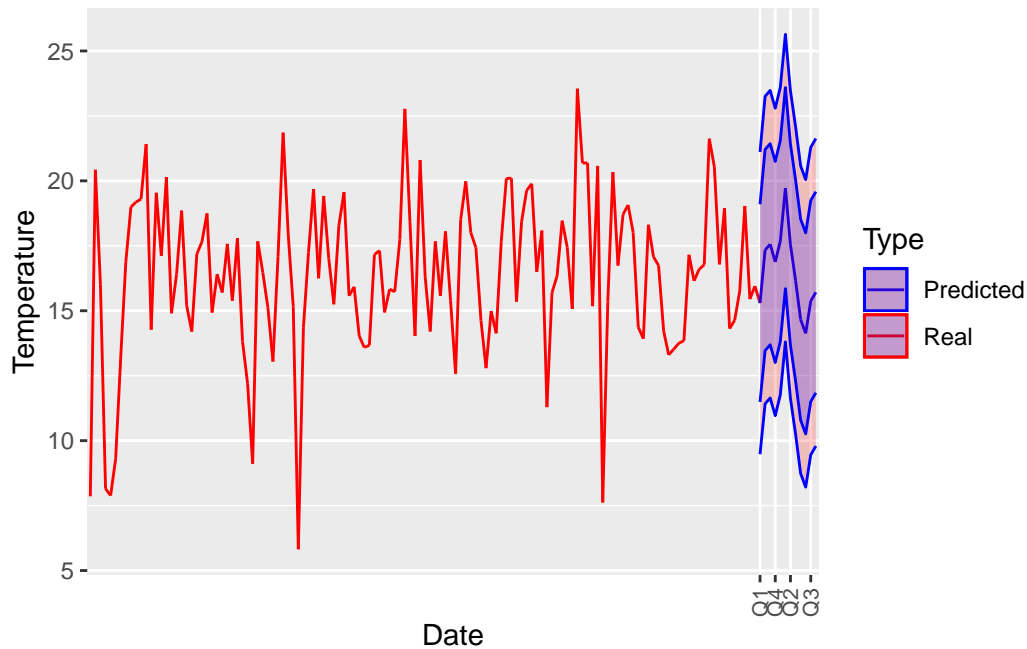
# Add the 80% and 95% confidence intervals
plotARIMA2 = plotARIMA2 + geom_ribbon(data = predictedArimaSeasonalDF, aes(x = YearMonth,
ymin = Lo.95, ymax = Hi.95), fill = "red", alpha = 0.2) + geom_ribbon(data = predictedArimaSeasonalDF,
aes(x = YearMonth, ymin = Lo.80, ymax = Hi.80), fill = "blue", alpha = 0.2)

# Adjust the x-axis labels
plotARIMA2 = plotARIMA2 + scale_x_discrete(breaks = c("2021-1", "2021-4",
"2021-8", "2021-12"), labels = c("Q1", "Q2", "Q3", "Q4"))

plotARIMA2 = plotARIMA2 + theme(axis.text.x = element_text(angle = 90, vjust = 0.5,
size = 8))

# Display the plot
print(plotARIMA2)

```



## DLM

### model fitting

```
## linear model, order = 2, quadratic order = 3 , etc

## what we want is a linear model with a seasonal component so we add
## the 2 components together in a model

## things to try, another term like quadratic, or a arma component
## stacked on top of this

## Initial model with a linear polynomial and a seasonal component

buildFun = function(x) {
  dlmModPoly(order = 2, dV = exp(x[1]), dW = c(0, exp(x[2]))) + dlmModSeas(frequency = 12)
  dV = 0, dW = c(exp(x[3]), rep(0, 10))
}

linearDLM = dlmMLE(tsAtlanticMonthly, parm = c(0, 0, 0), build = buildFun)
```



```
linearDLM$par
```

```
[1] 2.044026 -11.586366 -4.421181
```

```
fittedLinearDLM = buildFun(linearDLM$par)
```

```
V(fittedLinearDLM)
```

```
[,1]
```

```
[1,] 7.721632
```

```
W(fittedLinearDLM)
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]
[1,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[2,]	0	9.291914e-06	0.00000000	0	0	0	0	0	0	0	0
[3,]	0	0.000000e+00	0.01202003	0	0	0	0	0	0	0	0
[4,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[5,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[6,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[7,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[8,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[9,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[10,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[11,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[12,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[13,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
	[,12]	[,13]									
[1,]	0	0									
[2,]	0	0									
[3,]	0	0									
[4,]	0	0									
[5,]	0	0									
[6,]	0	0									
[7,]	0	0									
[8,]	0	0									
[9,]	0	0									
[10,]	0	0									
[11,]	0	0									

```
[12,]    0    0
[13,]    0    0
```

```
## second model with a quadratic polynomial and a seasonal component
```

```
buildFunQuad = function(x) {
  dlmModPoly(order = 3, dV = exp(x[1]), dW = c(0, exp(x[2]), exp(x[3]))) +
  dlmModSeas(frequency = 12, dV = 0, dW = c(exp(x[4]), rep(0, 10)))
}

quadraticDLM = dlmMLE(tsAtlanticMonthly, parm = c(0, 0, 0, 0), build = buildFunQuad)

quadraticDLM$par
```

```
[1] 2.047135 -21.060474 -56.104784 -12.300531
```

```
fittedQuadraticDLM = buildFunQuad(quadraticDLM$par)

V(fittedQuadraticDLM)
```

```
[,1]
[1,] 7.745678
```

```
W(fittedQuadraticDLM)
```

```
      [,1]      [,2]      [,3]      [,4] [,5] [,6] [,7] [,8] [,9]
[1,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[2,] 0 7.137603e-10 0.000000e+00 0.000000e+00 0 0 0 0 0
[3,] 0 0.000000e+00 4.305286e-25 0.000000e+00 0 0 0 0 0
[4,] 0 0.000000e+00 0.000000e+00 4.549326e-06 0 0 0 0 0
[5,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[6,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[7,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[8,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[9,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[10,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[11,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[12,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
```

```

[13,]    0 0.000000e+00 0.000000e+00 0.000000e+00    0    0    0    0    0
[14,]    0 0.000000e+00 0.000000e+00 0.000000e+00    0    0    0    0    0
      [,10] [,11] [,12] [,13] [,14]
[1,]      0      0      0      0      0
[2,]      0      0      0      0      0
[3,]      0      0      0      0      0
[4,]      0      0      0      0      0
[5,]      0      0      0      0      0
[6,]      0      0      0      0      0
[7,]      0      0      0      0      0
[8,]      0      0      0      0      0
[9,]      0      0      0      0      0
[10,]     0      0      0      0      0
[11,]     0      0      0      0      0
[12,]     0      0      0      0      0
[13,]     0      0      0      0      0
[14,]     0      0      0      0      0

```

Now we will compare both models through their log likelihood using the `dmlL` function and see if the extra flexibility from the extra polynomial function is providing a better fit

```
dmlL(tsAtlanticMonthly, fittedLinearDLM)
```

```
[1] 309.5446
```

```
dmlL(tsAtlanticMonthly, fittedQuadraticDLM)
```

```
[1] 324.4748
```

As we can see the dlm model using only a linear polynomial has a lower log likelihood than the model with an extra quadratic term, meaning this extra flexibility does not contribute to a better model fit and as such we will use the linear fitted model to do our forecasting.

```
amocPredict = dlmFilter(tsAtlanticMonthly, mod = fittedLinearDLM)
summary(amocPredict)
```

```

      Length Class  Mode
y       132   ts     numeric
mod      10   dlm     list

```

```

m 1729 mts numeric
U.C 133 -none- list
D.C 1729 -none- numeric
a 1716 mts numeric
U.R 132 -none- list
D.R 1716 -none- numeric
f 132 ts numeric

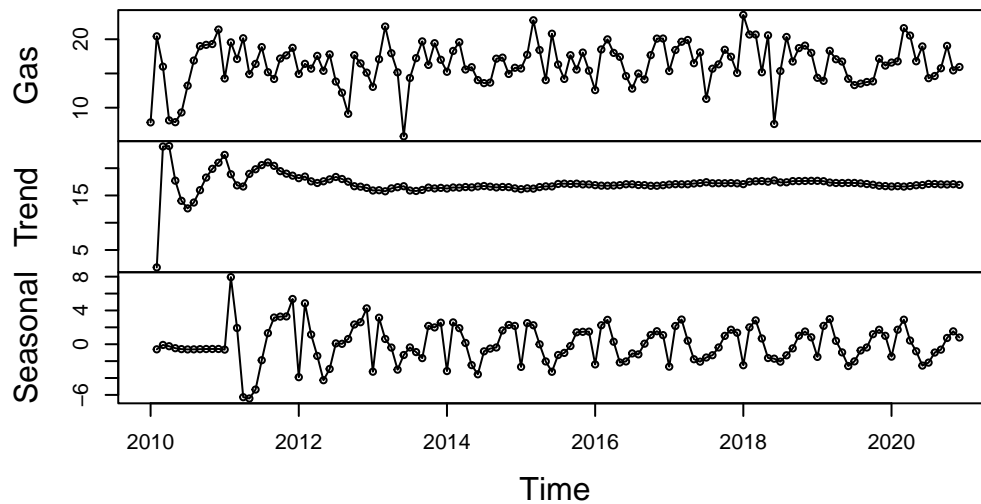
```

```

x = cbind(tsAtlanticMonthly, dropFirst(amocPredict$a[, c(1, 3)]))
x = window(x, start = c(2010, 1))
colnames(x) = c("Gas", "Trend", "Seasonal")
plot(x, type = "o", main = "Atlantic AMOC at 26,5N 2010-2020")

```

### Atlantic AMOC at 26,5N 2010–2020



### Forecast

```

amocForecastMonthly = dlmForecast(amocPredict, nAhead = 12)
summary(amocForecastMonthly)

```

```

Length Class Mode
a 156 mts numeric

```

```
R 12 -none- list
f 12 ts      numeric
Q 12 -none- list
```

```
dim(amocForecastMonthly$a)
```

```
[1] 12 13
```

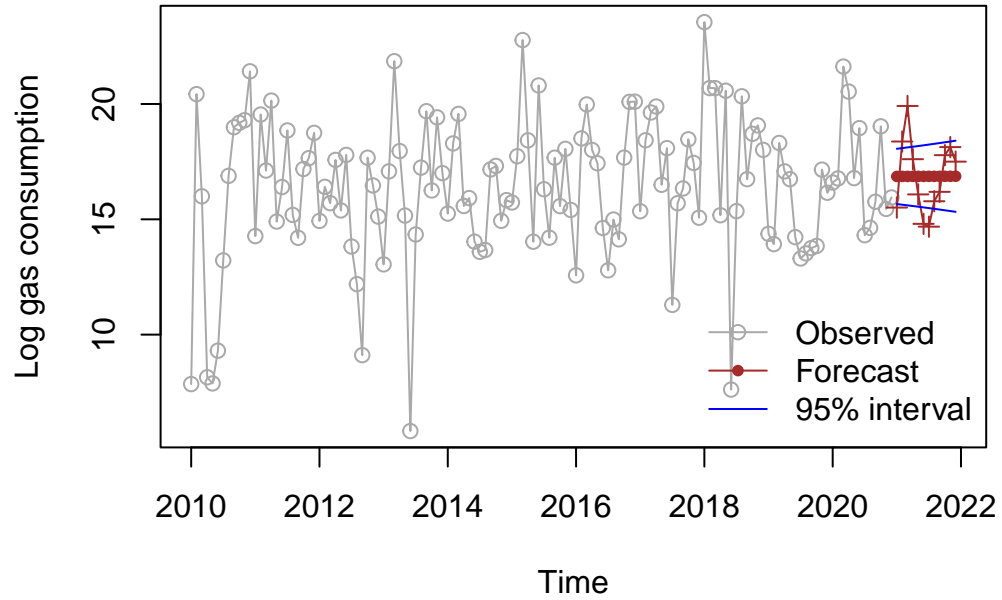
```
dim(amocForecastMonthly$f)
```

```
[1] 12 1
```

```
sqrtR = sapply(amocForecastMonthly$R, function(x) sqrt(x[1, 1]))
pl = amocForecastMonthly$a[, 1] + qnorm(0.025, sd = sqrtR)
pu = amocForecastMonthly$a[, 1] + qnorm(0.975, sd = sqrtR)

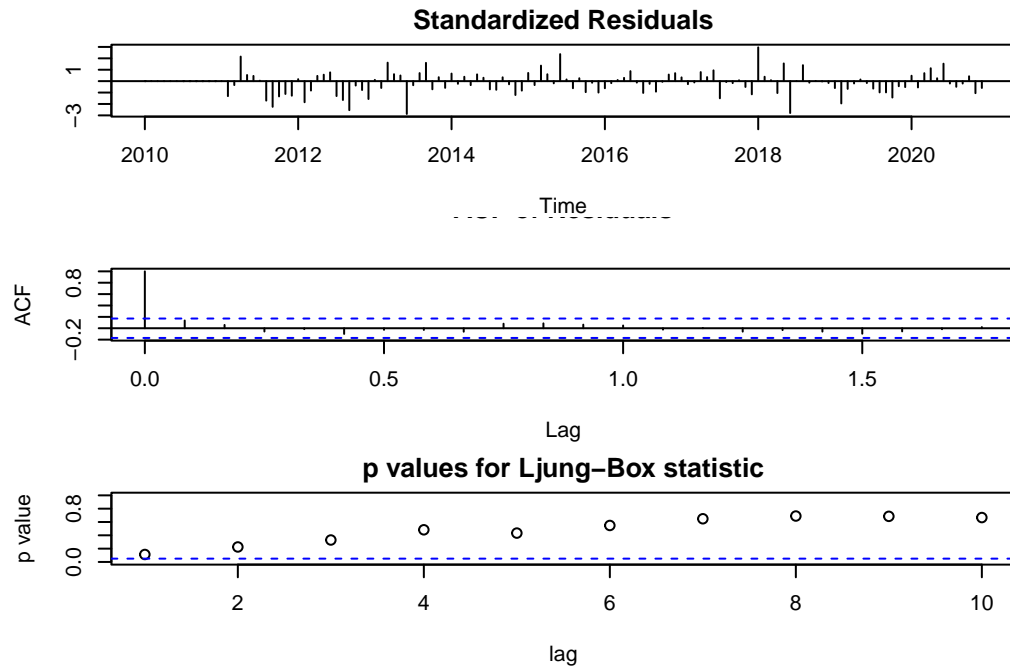
x = ts.union(window(tsAtlanticMonthly, start = c(2010, 1)), amocForecastMonthly$a[,
  1], amocForecastMonthly$f, pl, pu)
par(mar = c(4, 4, 2, 2))
plot(x, plot.type = "single", type = "o", pch = c(1, 20, 3, NA, NA), col = c("darkgrey",
  "brown", "brown", "blue", "blue"), ylab = "Log gas consumption")

legend("bottomright", legend = c("Observed", "Forecast", "95% interval"),
  bty = "n", pch = c(1, 20, NA), lty = 1, col = c("darkgrey", "brown",
  "blue"))
```



```
# Set smaller margins
par(mar = c(4, 4, 2, 2))

tsdiag(amocPredict)
```



```
# Reset margins
par(mar = c(5, 4, 4, 2) + 0.1)
```

Lastly checking the residuals, they seem to be mostly normally distributed with a good mixture of over and under estimations, especially in the middle with some slight seasonality on both ends being present

## 2 f)

Now starting with the ARIMA models

```
predictedArimaDF = data.frame(forecast(model200, 4))

predictedArimaDF$YearQuarter = c("2021-Q1", "2021-Q2", "2021-Q3", "2021-Q4")

# Combine real_data and pred_data into a single data frame
combinedDataframeAMOC = rbind(data.frame(Date = YearQuarterAverage$YearQuarter,
  Temperature = YearQuarterAverage$AverageStrength, Type = "Real"), data.frame(Date = pr
  Temperature = predictedArimaDF$Point.Forecast, Type = "Predicted"))

predictedArimaDF$Temperature = predictedArimaDF$Point.Forecast
```

```

predictedArimaDF$Type = "Predicted"

# Create the ggplot
plotARIMA = ggplot(combinedDataframeAMOC, aes(x = Date, y = Temperature,
  color = Type, group = 1)) + geom_line() + scale_color_manual(values = c("blue",
  "red"))

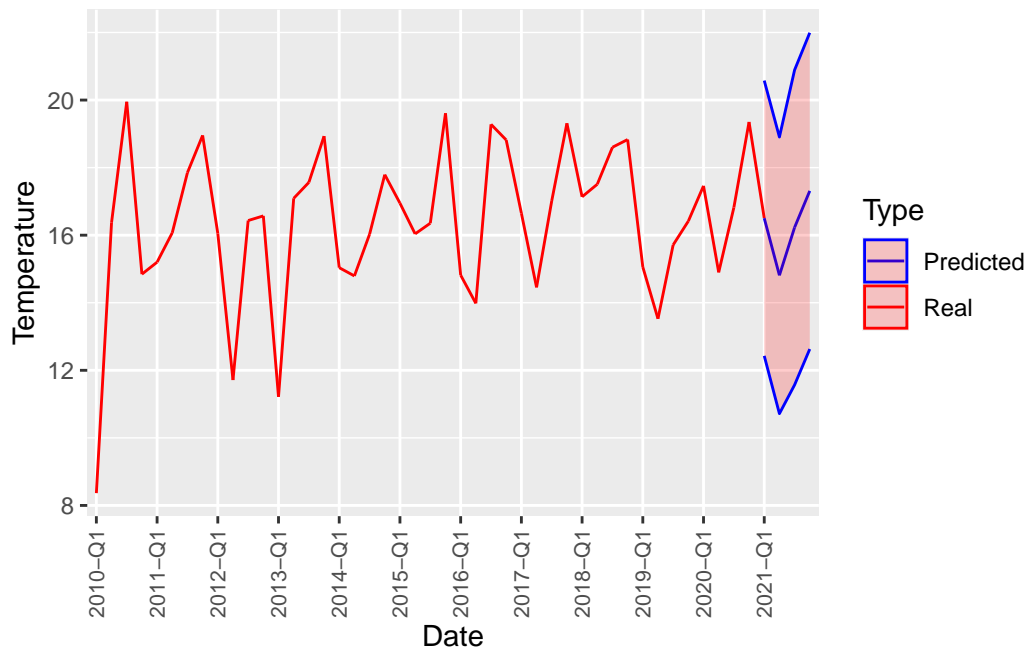
# Add the 95% confidence interval
plotARIMA = plotARIMA + geom_ribbon(data = predictedArimaDF, aes(x = YearQuarter,
  ymin = Lo.95, ymax = Hi.95), fill = "red", alpha = 0.2)

# Adjust the x-axis labels
plotARIMA = plotARIMA + scale_x_discrete(breaks = combinedDataframeAMOC$Date[c(TRUE,
  rep(FALSE, 3))], labels = combinedDataframeAMOC$Date[c(TRUE, rep(FALSE,
  3))])

plotARIMA = plotARIMA + theme(axis.text.x = element_text(angle = 90, vjust = 0.5,
  size = 8))

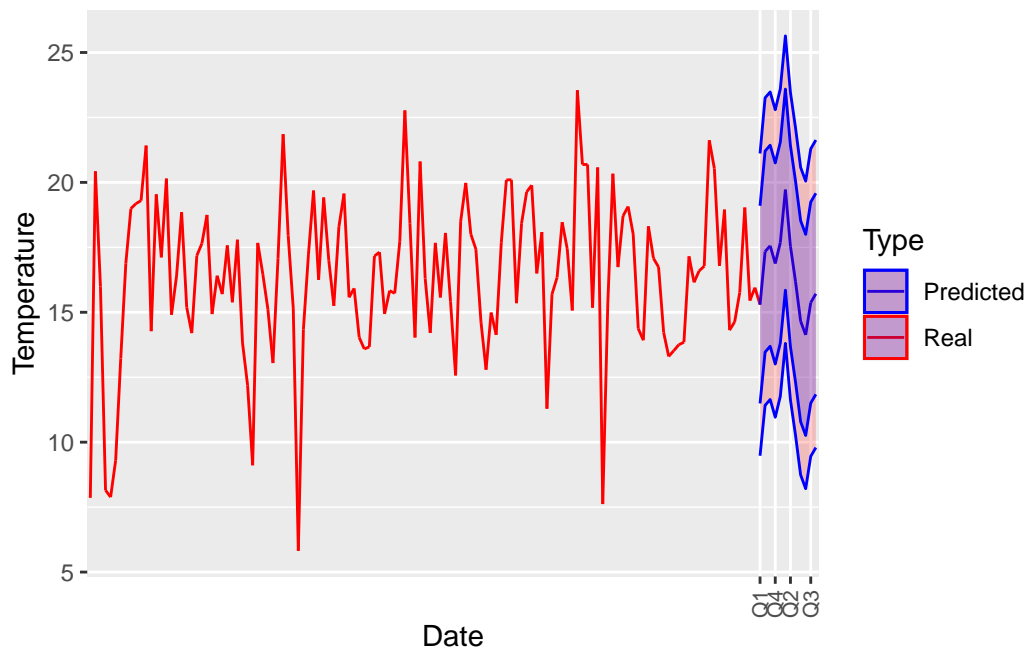
print(plotARIMA)

```





```
print(plotARIMA2)
```



The most obvious difference is the level of detail on the seasonality, with the monthly averages capturing an almost opposite effect than the quarterly averages in both the predicted and also in some of the real data, the forecast also has the opposite trend, with an actual expected decrease in Sverdrups around between the 2nd quarter and the mid 3rd quarter which again seems to follow the opposite trend on the quarterly data.

```
sqrtr = sapply(amocForecastMonthly$R, function(x) sqrt(x[1, 1]))
pl = amocForecastMonthly$a[, 1] + qnorm(0.025, sd = sqrtr)
pu = amocForecastMonthly$a[, 1] + qnorm(0.975, sd = sqrtr)

x = ts.union(window(tsAtlanticMonthly, start = c(2010, 1)), amocForecastMonthly$a[,
  1], amocForecastMonthly$f, pl, pu)
par(mar = c(4, 4, 2, 2))
plot(x, plot.type = "single", type = "o", pch = c(1, 20, 3, NA, NA), col = c("darkgrey",
  "brown", "brown", "blue", "blue"), ylab = "Log gas consumption")

legend("bottomright", legend = c("Observed", "Forecast", "95% interval"),
  bty = "n", pch = c(1, 20, NA), lty = 1, col = c("darkgrey", "brown",
  "blue"))
```

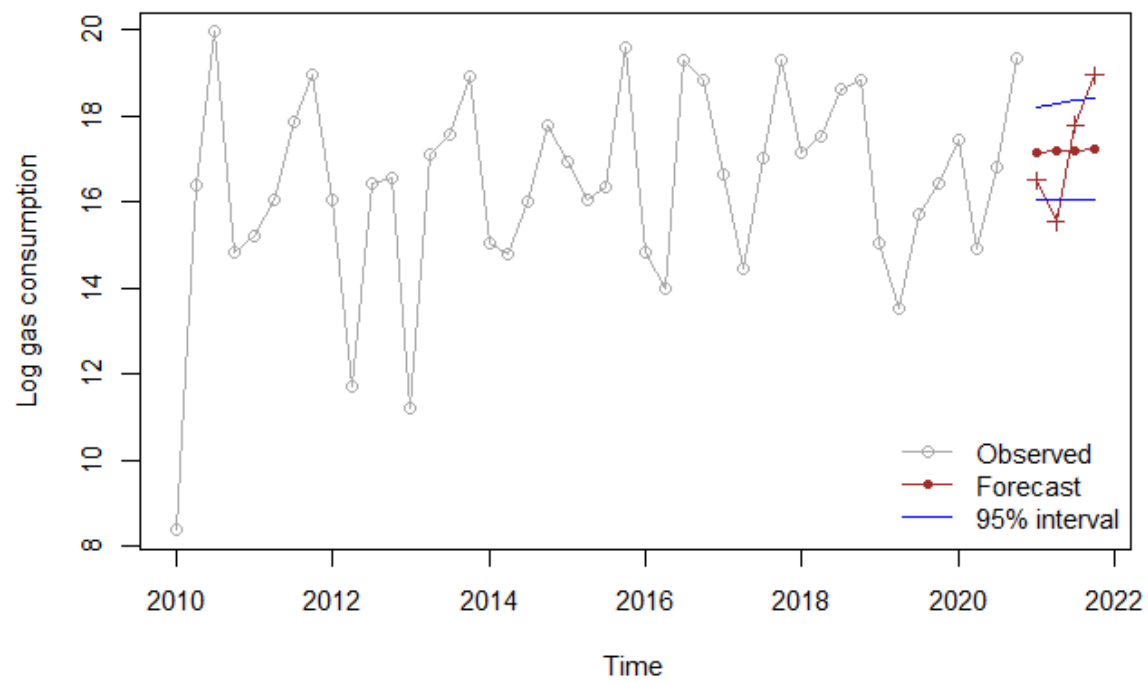
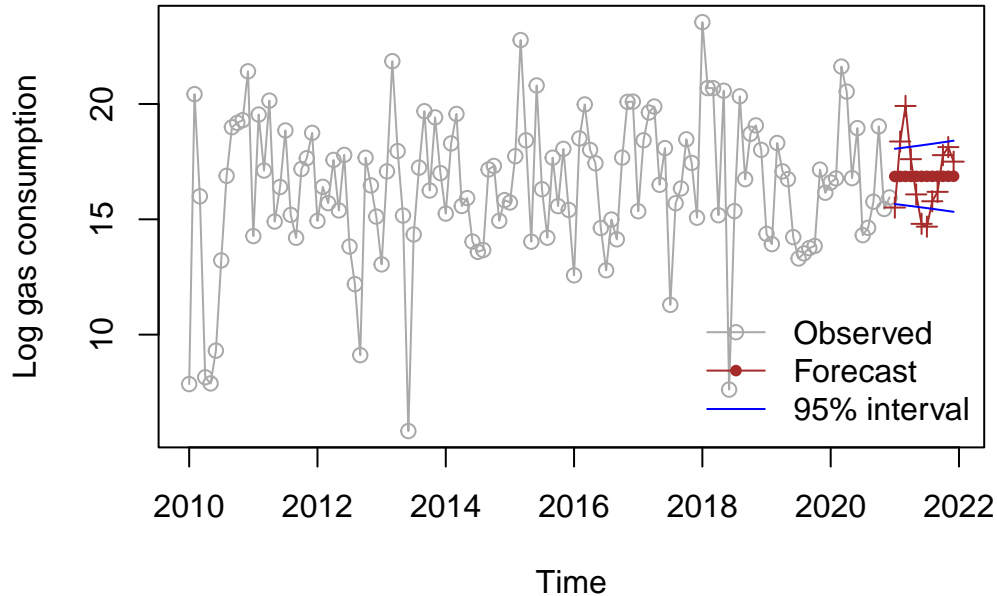


Figure 1: 1st dlm



Observing the 2 graphs of the dlm models we can see that although the quarterly predictions do not completely predict the spike during the first half of the year compared to the monthly data the second half of the year seems to be relative similarly forecasted with the exception of December where the monthly data show again a decrease but the quarterly data is not capable of capturing.

Despite these changes the predicted overall trend is quite similar with the quarterly trend very slightly increasing the monthly data seeming to remain constant.

### Question 3

#### Question 3 a)

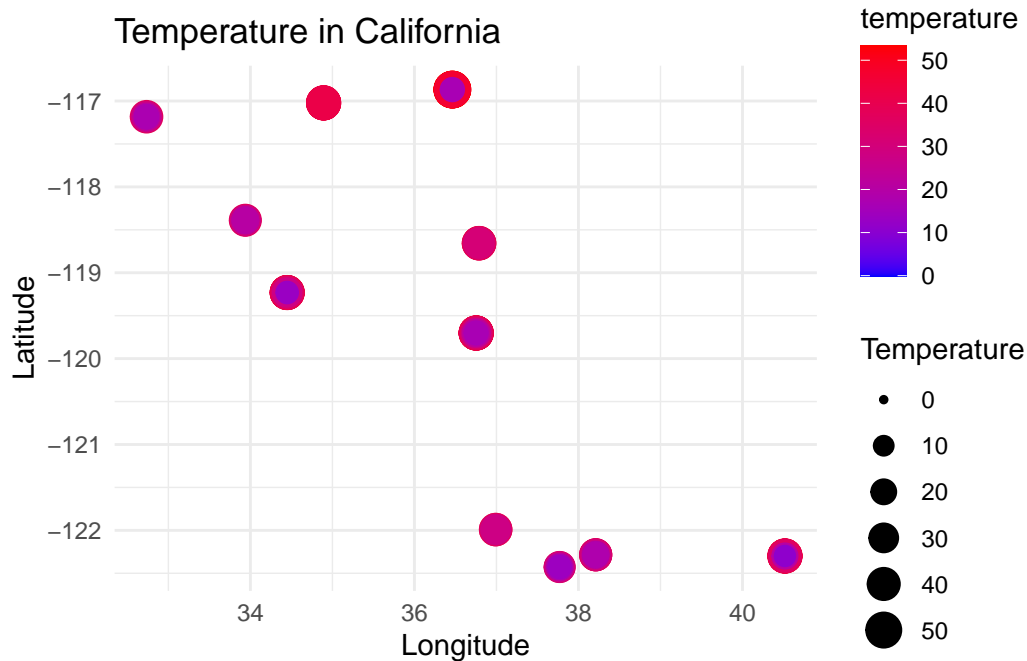
my approach is to see the max temp in the entire state with 8 cities

```
californiaLongTempDF = pivot_longer(californiaTempDF, cols = -Date, names_to = "Location",
                                     values_to = "Temperature")

spatialTemperatureCaliforniaDF = merge(californiaLongTempDF, californiaSpatialDataDF)

ggplot(data = spatialTemperatureCaliforniaDF) + geom_point(aes(x = Lat, y = Long,
```

```
color = Temperature, size = Temperature)) + scale_color_continuous(low = "blue",
high = "red") + labs(title = "Temperature in California", x = "Longitude",
y = "Latitude", color = "temperature") + theme_minimal()
```



```
californiaTempDF$Date = as.Date(as.character(californiaTempDF$Date), format = "%Y%m%d",
origin = "1970-01-01")

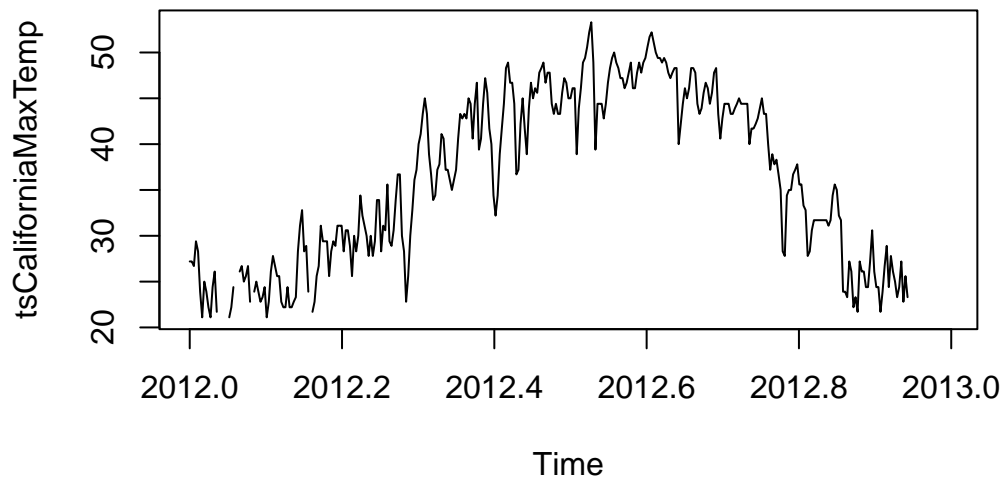
californiaTempDF$max <- apply(californiaTempDF, 1, max, na.rm = TRUE)

tsCaliforniaMaxTemp = ts(californiaTempDF$max, start = c(2012, 1), frequency = 366)

plot.ts(tsCaliforniaMaxTemp)
```

Warning in xy.coords(x, NULL, log = log, setLab = FALSE): NAs introduced by coercion

Warning in xy.coords(x, y): NAs introduced by coercion



As we can see from the time series trend, august to September seems to be the hottest months while january to february seems to be the coldest months in the californian state.

### 3 b)

```
geoDataCalifornia = as.geodata(spatialTemperatureCaliforniaDF, coords.col = 4:5,
                               data.col = "Temperature", covar.col = "Elev")
```

as.geodata: 4004 replicated data locations found.

Consider using jitterDupCoords() for jittering replicated locations.

WARNING: there are data at coincident or very closed locations, some of the geoR's functions

Use function dup.coords() to locate duplicated coordinates.

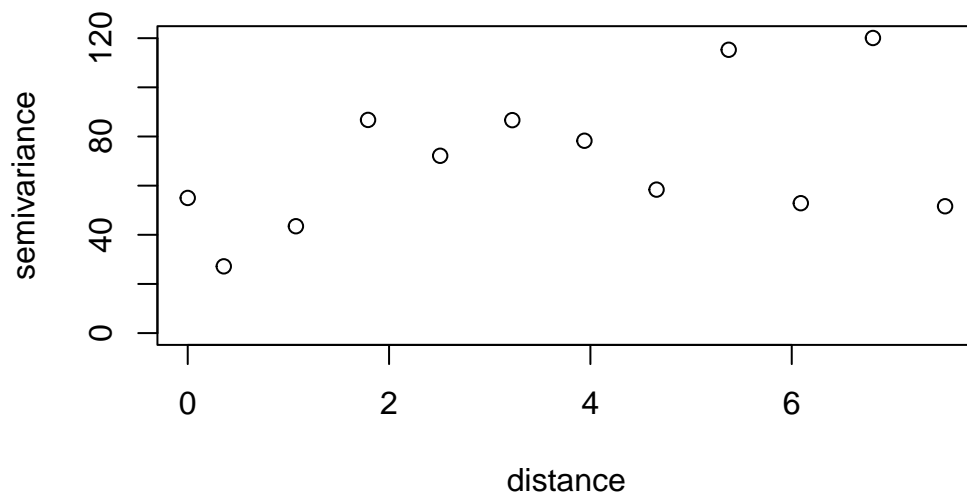
Consider using jitterDupCoords() for jittering replicated locations

```
variogramCalifornia = variog(geoDataCalifornia)
```

variog: computing omnidirectional variogram

variog: co-located data found, adding one bin at the origin

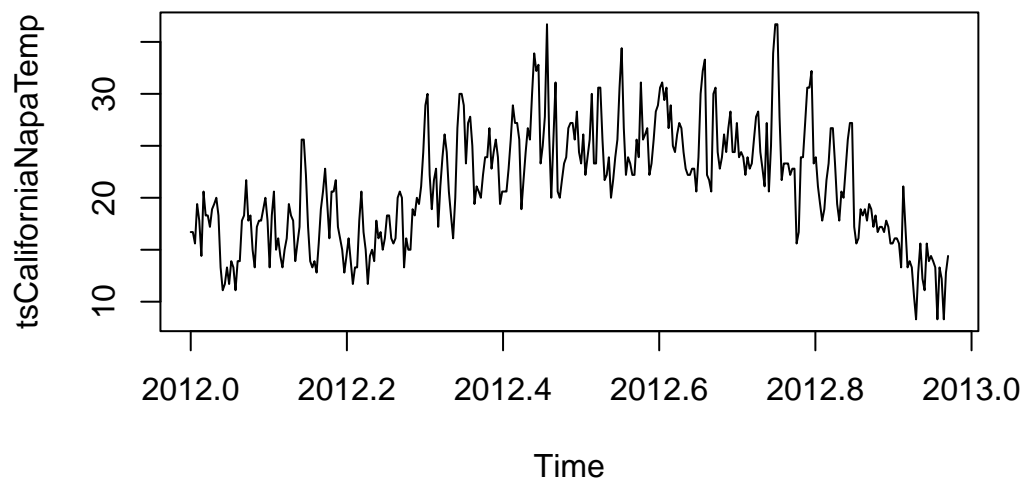
```
plot(variogramCalifornia)
```



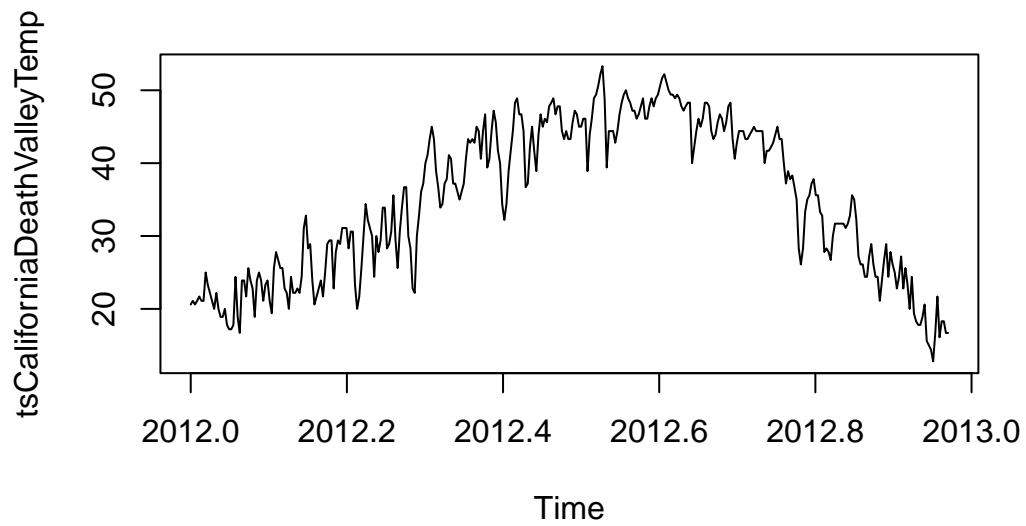
### 3 c)

Here we create a ts object with frequency 366 because 2012 does indeed have the 29th of February, removing the dates from the 9th to the 17th of november

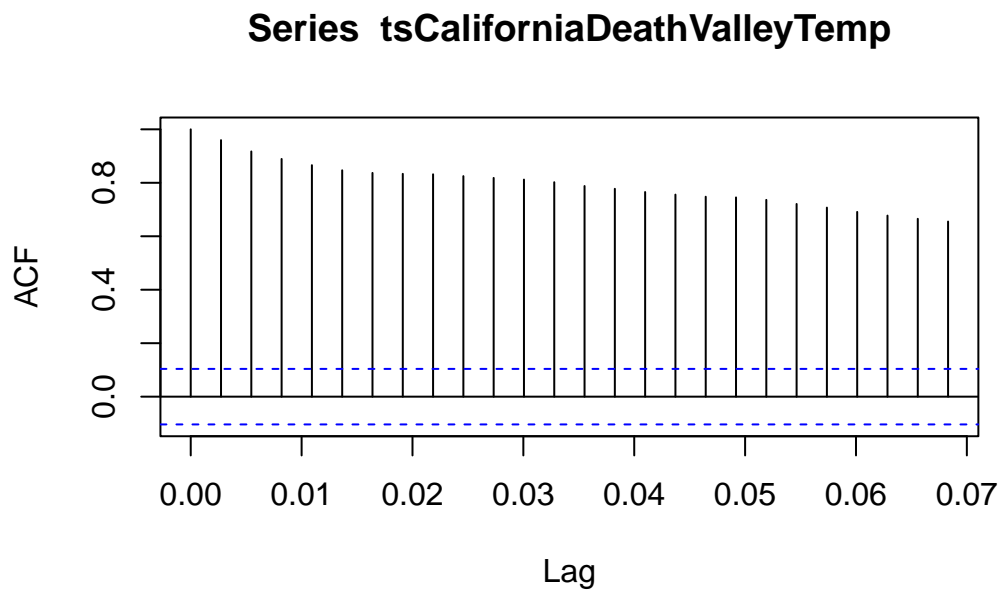
```
toPredictValues = californiaTempDF %>%  
  filter((Date >= as.Date("2012-11-09") & Date <= as.Date("2012-11-17")))  
  
californiaTempDF = californiaTempDF %>%  
  filter(!(Date >= as.Date("2012-11-09") & Date <= as.Date("2012-11-17")))  
  
tsCaliforniaNapaTemp = ts(californiaTempDF$Napa, start = c(2012, 1), frequency = 366)  
  
plot.ts(tsCaliforniaNapaTemp)
```



```
tsCaliforniaDeathValleyTemp = ts(californiaTempDF$`Death Valley`, start = c(2012,  
1), frequency = 366)  
  
plot.ts(tsCaliforniaDeathValleyTemp)
```

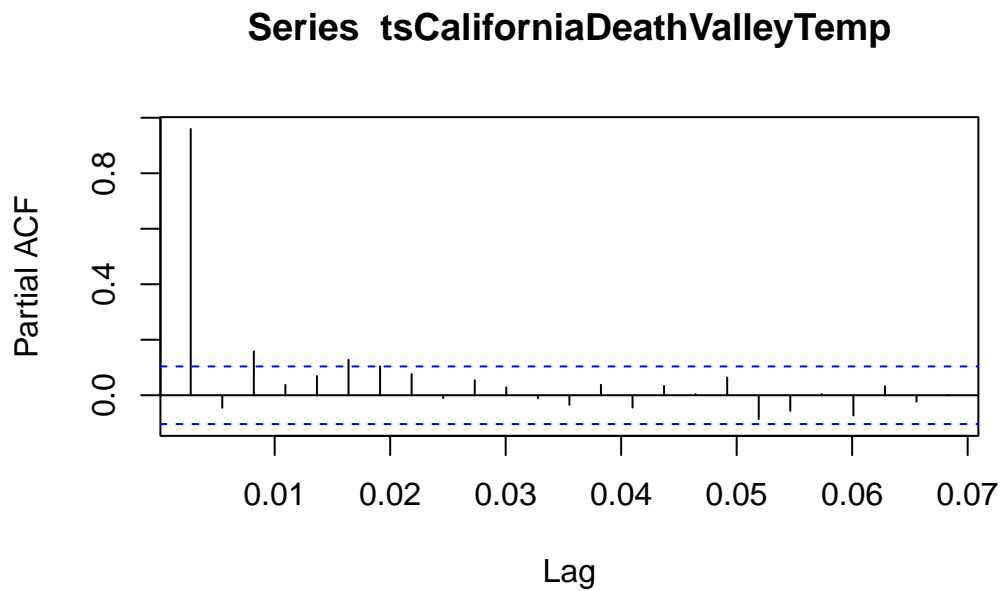


```
acf(tsCaliforniaDeathValleyTemp)
```





```
pacf(tsCaliforniaDeathValleyTemp)
```

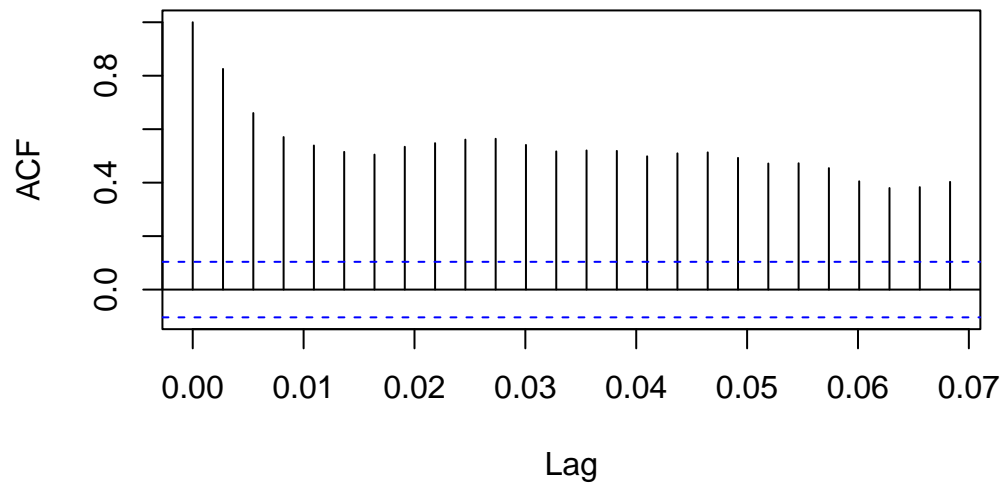


In death valley we can clearly see that ACF is very slowly descending without any cut-off and PAC seems to have a very quick cut off, this means that the best model will most likely will be AR with a possibly larger p value

Now checking for Napa

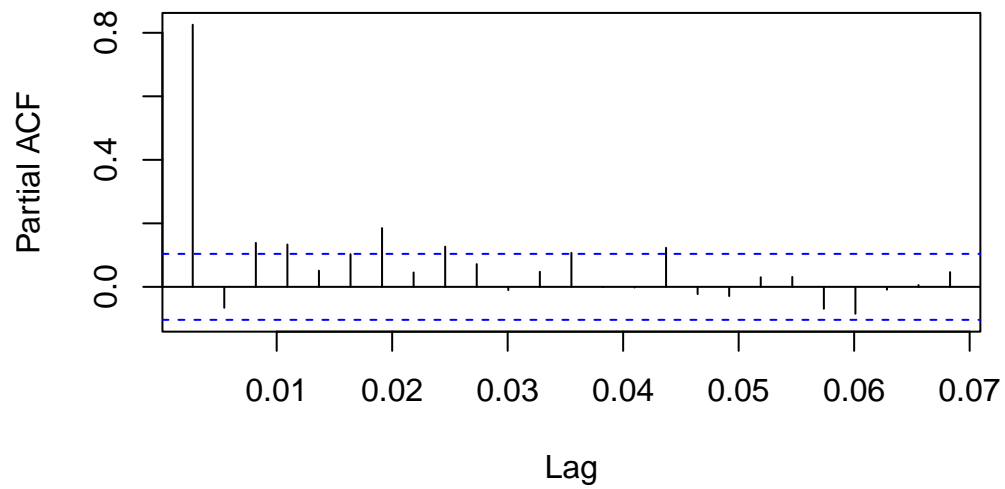
```
acf(tsCaliforniaNapaTemp)
```

**Series tsCaliforniaNapaTemp**



```
pacf(tsCaliforniaNapaTemp)
```

**Series tsCaliforniaNapaTemp**



We can once again see the same pattern of a very slowly decreasing ACF but this time PACF does seem to at least not has such a clear cut off meanig there might a smaller q non-zero q value requiring an ARIMA for this city

The model for both cities does seem stationary with a linear trend on seasonality however it is not clear enough since we only have 1 year worth of data

### model checking Nappa

```
## initial assumption
```

```
modelMonthlySeasonal100.110Napa = Arima(tsCaliforniaNapaTemp, order = c(1,  
0, 0), seasonal = list(order = c(1, 1, 0), period = 12))
```

```
modelMonthlySeasonal100.011Napa = Arima(tsCaliforniaNapaTemp, order = c(1,  
0, 0), seasonal = list(order = c(0, 1, 1), period = 12))
```

```
modelMonthlySeasonal200.210Napa = Arima(tsCaliforniaNapaTemp, order = c(2,  
0, 0), seasonal = list(order = c(2, 1, 0), period = 12))
```

```
modelMonthlySeasonal200.012Napa = Arima(tsCaliforniaNapaTemp, order = c(2,  
0, 0), seasonal = list(order = c(0, 1, 2), period = 12))
```

```
modelMonthlySeasonal001.110Napa = Arima(tsCaliforniaNapaTemp, order = c(0,  
0, 1), seasonal = list(order = c(1, 1, 0), period = 12))
```

```
modelMonthlySeasonal001.011Napa = Arima(tsCaliforniaNapaTemp, order = c(0,  
0, 1), seasonal = list(order = c(0, 1, 1), period = 12))
```

```
modelMonthlySeasonal002.210Napa = Arima(tsCaliforniaNapaTemp, order = c(0,  
0, 2), seasonal = list(order = c(2, 1, 0), period = 12))
```

```
modelMonthlySeasonal002.012Napa = Arima(tsCaliforniaNapaTemp, order = c(0,  
0, 2), seasonal = list(order = c(0, 1, 2), period = 12))
```

```
modelMonthlySeasonal100.110Napa
```

```
Series: tsCaliforniaNapaTemp  
ARIMA(1,0,0)(1,1,0)[12]
```

```
Coefficients:
```

```
ar1      sar1
```

```
      0.6323  -0.4915
s.e.  0.0418   0.0474
```

```
sigma^2 = 14.03:  log likelihood = -943.36
AIC=1892.73  AICc=1892.8   BIC=1904.25
```

```
modelMonthlySeasonal100.011Napa
```

```
Series: tsCaliforniaNapaTemp
ARIMA(1,0,0)(0,1,1)[12]
```

```
Coefficients:
```

```
      ar1      sma1
      0.8307  -0.9637
s.e.  0.0329   0.0780
```

```
sigma^2 = 10.12:  log likelihood = -900.8
AIC=1807.61  AICc=1807.68  BIC=1819.13
```

```
modelMonthlySeasonal200.210Napa
```

```
Series: tsCaliforniaNapaTemp
ARIMA(2,0,0)(2,1,0)[12]
```

```
Coefficients:
```

```
      ar1      ar2      sar1      sar2
      0.7592  -0.1464  -0.6312  -0.2890
s.e.  0.0535   0.0540   0.0539   0.0534
```

```
sigma^2 = 12.47:  log likelihood = -923.09
AIC=1856.19  AICc=1856.36  BIC=1875.39
```

```
modelMonthlySeasonal200.012Napa
```

```
Series: tsCaliforniaNapaTemp
ARIMA(2,0,0)(0,1,2)[12]
```

```
Coefficients:
```

```
      ar1      ar2      sma1      sma2
```

```

          0.8807  -0.0576  -0.9749  0.0349
s.e.    0.0546   0.0550   0.0730  0.0634

sigma^2 = 10.26:  log likelihood = -900.05
AIC=1810.11   AICc=1810.29   BIC=1829.31

```

```
modelMonthlySeasonal001.110Napa
```

```

Series: tsCaliforniaNapaTemp
ARIMA(0,0,1)(1,1,0)[12]

```

Coefficients:

```

          ma1      sar1
          0.5787  -0.4244
s.e.    0.0365   0.0493

```

```

sigma^2 = 15.15:  log likelihood = -955.99
AIC=1917.98   AICc=1918.05   BIC=1929.5

```

```
modelMonthlySeasonal001.011Napa
```

```

Series: tsCaliforniaNapaTemp
ARIMA(0,0,1)(0,1,1)[12]

```

Coefficients:

```

          ma1      sma1
          0.6150  -0.5909
s.e.    0.0358   0.0459

```

```

sigma^2 = 13.72:  log likelihood = -940.43
AIC=1886.85   AICc=1886.92   BIC=1898.37

```

```
modelMonthlySeasonal002.210Napa
```

```

Series: tsCaliforniaNapaTemp
ARIMA(0,0,2)(2,1,0)[12]

```

Coefficients:

```

          ma1      ma2      sar1      sar2

```

```

      0.7380  0.3614  -0.6225  -0.2583
s.e.  0.0488  0.0552   0.0561   0.0531

```

```

sigma^2 = 12.66:  log likelihood = -925.48
AIC=1860.95   AICc=1861.13   BIC=1880.16

```

```

modelMonthlySeasonal002.012Napa

```

```

Series: tsCaliforniaNapaTemp
ARIMA(0,0,2)(0,1,2)[12]

```

Coefficients:

```

      ma1      ma2      sma1      sma2
      0.7903  0.3957  -0.7285  0.0584
s.e.  0.0494  0.0537   0.0609  0.0608

```

```

sigma^2 = 11.87:  log likelihood = -915.85
AIC=1841.71   AICc=1841.88   BIC=1860.91

```

So far it seems the second model is the most adequate

```

NapaModelAuto = auto.arima(tsCaliforniaNapaTemp, max.d = 0, max.p = 5, max.q = 5)

NapaModelAuto

```

```

Series: tsCaliforniaNapaTemp
ARIMA(1,0,0) with non-zero mean

```

Coefficients:

```

      ar1      mean
      0.8280  21.0196
s.e.  0.0296   0.9505

```

```

sigma^2 = 9.813:  log likelihood = -911.23
AIC=1828.45   AICc=1828.52   BIC=1840.08

```

it seems there isn't enough data for auto.arima to fully caught the seasonality needed

## Forecasting

Now using the forecast function we will produce an out of time cross-validation to forecast the values in the 2 weeks of November

Here we simply

```
forecast(modelMonthlySeasonal100.011Napa, h = 8)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2012.9727	14.92598	10.83022	19.02173	8.662061	21.18990
2012.9754	16.26633	10.94209	21.59057	8.123605	24.40905
2012.9781	18.45885	12.43152	24.48617	9.240851	27.67684
2012.9809	18.64136	12.17375	25.10897	8.750004	28.53272
2012.9836	18.76622	12.01212	25.52031	8.436723	29.09571
2012.9863	18.98751	12.04219	25.93284	8.365557	29.60947
2012.9891	19.08760	12.01335	26.16185	8.268467	29.90673
2012.9918	19.32427	12.16243	26.48611	8.371186	30.27735

```
toPredictValues$Napa
```

```
[1] 14.4 15.0 14.4 16.7 20.6 22.2 20.0 17.2 17.8
```

```
# Create a new time series object with missing values replaced by zeros
tsCaliforniaNapaTempFilled <- tsCaliforniaNapaTemp
tsCaliforniaNapaTempFilled[is.na(tsCaliforniaNapaTempFilled)] <- 0

# Update the ARIMA model with the filled time series (missing values
# replaced by zeros)
updatedArimaModelNappa <- Arima(tsCaliforniaNapaTempFilled, model = modelMonthlySeasonal100.011Napa)

# Generate predictions for the entire time series (including missing
# values)
predictionsDeath <- fitted(updatedArimaModelNappa)

predictionsDeath
```

Time Series:

Start = c(2012, 1)

End = c(2012, 356)

Frequency = 366

[1]	16.68330	16.68330	15.58440	19.38060	17.78220	14.38560	20.57940	18.28170
[9]	18.28170	17.18280	18.88110	19.38060	18.59940	19.06863	15.85261	15.63821
[17]	11.05962	10.34994	17.35326	11.64472	14.13739	12.37380	13.71735	14.73191
[25]	15.85803	17.22831	16.18963	20.64107	17.05881	16.43539	16.39671	14.24005
[33]	16.78616	15.99548	19.14825	18.69138	20.76960	16.95252	13.74890	17.65155
[41]	19.08643	13.73266	16.42205	15.46028	13.91842	14.30270	17.91641	19.28214
[49]	18.69308	16.17816	14.99305	15.97614	16.53961	23.12920	23.83668	21.20682
[57]	17.00966	13.63755	15.67451	14.46407	14.18706	14.41783	19.32189	20.46238
[65]	22.10434	18.14145	16.10632	19.48240	19.61224	19.86635	18.51169	16.00786
[73]	15.96576	12.67467	15.74682	16.79950	14.27980	11.64648	14.21259	13.80501
[81]	17.70835	18.58276	17.70569	14.83564	13.01406	14.20577	16.16423	14.85718
[89]	17.22505	15.51704	16.86009	15.61713	16.65817	16.59552	18.66010	15.46579
[97]	16.36397	15.86294	20.14824	20.63980	18.72487	13.40721	16.15754	15.63116
[105]	16.00752	17.28146	18.50916	18.73588	19.54886	20.47454	24.01532	27.57129
[113]	26.45504	20.92447	18.50603	21.19723	22.50109	16.12126	21.03860	22.14762
[121]	25.17380	23.40072	21.11120	18.84717	14.77182	19.18808	25.35061	28.24964
[129]	28.11613	26.08474	23.17983	25.23080	26.46381	23.59399	19.95740	20.97688
[137]	18.88297	19.74981	21.97044	23.31088	23.19376	24.03030	23.05936	23.13590
[145]	24.51160	22.34398	20.04087	20.53741	18.98992	20.45608	22.65819	24.85427
[153]	27.55166	24.42946	26.76010	24.15347	18.97385	20.42745	24.27930	25.60679
[161]	23.38174	28.51865	32.05372	30.62186	30.60903	21.43196	24.92374	25.88870
[169]	33.73334	24.22488	20.85281	24.81283	28.17746	20.99571	20.53706	22.04656
[177]	22.32909	22.21424	26.50611	26.02657	25.40991	23.65819	27.94395	24.14387
[185]	21.36465	25.45095	22.45047	23.96747	24.41776	27.44987	23.70999	22.97793
[193]	28.27307	28.01716	25.45375	21.85969	20.74744	23.39587	20.71130	22.24764
[201]	23.33094	23.67682	29.81107	32.49853	25.10717	20.95731	23.84190	23.19189
[209]	20.92622	21.86220	25.39171	24.25460	29.37329	24.05959	26.31548	25.75928
[217]	21.24964	22.06189	25.26530	27.31552	26.60494	29.05052	29.87902	29.18026
[225]	28.74359	25.12107	28.61309	24.17137	23.21153	24.58854	26.74016	26.04356
[233]	23.05670	22.66459	22.45933	23.26225	22.17190	22.04711	21.56955	23.34098
[241]	28.04299	29.81450	31.70470	22.22109	20.79809	20.90811	28.94061	30.15434
[249]	23.51946	21.99396	24.43237	25.46705	23.61689	25.43410	27.08552	24.02738
[257]	23.07023	26.78301	23.95405	24.72919	23.06131	21.60401	24.49198	22.68087
[265]	22.86974	24.67731	26.56307	27.25869	23.25337	22.98054	21.69558	26.96149
[273]	20.32074	24.11678	32.75180	34.26981	34.06457	26.60480	21.58668	22.97997
[281]	22.32653	23.30251	22.85355	23.02015	22.30287	16.74509	18.68352	23.74926
[289]	23.18520	25.92480	29.06550	29.07999	29.71512	23.25821	24.23991	21.58191
[297]	19.17358	18.57138	20.74536	21.93699	22.88888	25.71792	25.87497	23.12613
[305]	18.79846	18.72941	21.35355	20.58782	21.96133	25.07692	27.67563	26.50174
[313]	17.97276	16.53716	16.97229	19.30668	17.84913	19.71203	18.97348	20.17718
[321]	18.86856	18.19310	20.24439	17.38522	17.86961	17.84165	17.55355	18.34449
[329]	16.97243	16.88366	17.15874	17.37589	16.46852	16.85004	15.97152	21.01196



```
[337] 17.81845 14.55153 15.22849 14.53917 11.41759 10.77754 14.79383 16.88453
[345] 13.18359 12.97920 18.07641 14.87311 15.29998 15.01526 14.62115 10.23834
[353] 13.52607 14.07384 11.09789 14.36658
```

## Death valley

```
## initial assumption
```

```
modelMonthlySeasonal100.110DeathValley = Arima(tsCaliforniaDeathValleyTemp,
  order = c(1, 0, 0), seasonal = list(order = c(1, 1, 0), period = 12))
```

```
modelMonthlySeasonal100.011DeathValley = Arima(tsCaliforniaDeathValleyTemp,
  order = c(1, 0, 0), seasonal = list(order = c(0, 1, 1), period = 12))
```

```
modelMonthlySeasonal200.210DeathValley = Arima(tsCaliforniaDeathValleyTemp,
  order = c(2, 0, 0), seasonal = list(order = c(2, 1, 0), period = 12))
```

```
modelMonthlySeasonal200.012DeathValley = Arima(tsCaliforniaDeathValleyTemp,
  order = c(2, 0, 0), seasonal = list(order = c(0, 1, 2), period = 12))
```

```
modelMonthlySeasonal001.110DeathValley = Arima(tsCaliforniaDeathValleyTemp,
  order = c(0, 0, 1), seasonal = list(order = c(1, 1, 0), period = 12))
```

```
modelMonthlySeasonal001.011DeathValley = Arima(tsCaliforniaDeathValleyTemp,
  order = c(0, 0, 1), seasonal = list(order = c(0, 1, 1), period = 12))
```

```
modelMonthlySeasonal002.210DeathValley = Arima(tsCaliforniaDeathValleyTemp,
  order = c(0, 0, 2), seasonal = list(order = c(2, 1, 0), period = 12))
```

```
modelMonthlySeasonal002.012DeathValley = Arima(tsCaliforniaDeathValleyTemp,
  order = c(0, 0, 2), seasonal = list(order = c(0, 1, 2), period = 12))
```

```
modelMonthlySeasonal100.110DeathValley
```

Series: tsCaliforniaDeathValleyTemp

ARIMA(1,0,0)(1,1,0)[12]

Coefficients:

	ar1	sar1
	0.8082	-0.4482
s.e.	0.0321	0.0491

```
sigma^2 = 9.802: log likelihood = -881.56  
AIC=1769.13 AICc=1769.2 BIC=1780.65
```

```
modelMonthlySeasonal100.011DeathValley
```

```
Series: tsCaliforniaDeathValleyTemp  
ARIMA(1,0,0)(0,1,1)[12]
```

```
Coefficients:
```

	ar1	sma1
	0.9679	-0.8985
s.e.	0.0164	0.0443

```
sigma^2 = 7.777: log likelihood = -850.22  
AIC=1706.45 AICc=1706.52 BIC=1717.97
```

```
modelMonthlySeasonal200.210DeathValley
```

```
Series: tsCaliforniaDeathValleyTemp  
ARIMA(2,0,0)(2,1,0)[12]
```

```
Coefficients:
```

	ar1	ar2	sar1	sar2
	0.9640	-0.1497	-0.5483	-0.1975
s.e.	0.0538	0.0546	0.0555	0.0578

```
sigma^2 = 9.163: log likelihood = -869.61  
AIC=1749.21 AICc=1749.39 BIC=1768.42
```

```
modelMonthlySeasonal200.012DeathValley
```

```
Series: tsCaliforniaDeathValleyTemp  
ARIMA(2,0,0)(0,1,2)[12]
```

```
Coefficients:
```

	ar1	ar2	sma1	sma2
	1.0154	-0.0521	-0.8544	-0.0521
s.e.	0.0538	0.0544	0.0608	0.0560

```
sigma^2 = 7.769: log likelihood = -849.32  
AIC=1708.63 AICc=1708.81 BIC=1727.84
```

```
modelMonthlySeasonal001.110DeathValley
```

```
Series: tsCaliforniaDeathValleyTemp  
ARIMA(0,0,1)(1,1,0)[12]
```

```
Coefficients:
```

	ma1	sar1
	0.7398	-0.2828
s.e.	0.0308	0.0525

```
sigma^2 = 13.16: log likelihood = -931.29  
AIC=1868.57 AICc=1868.64 BIC=1880.09
```

```
modelMonthlySeasonal001.011DeathValley
```

```
Series: tsCaliforniaDeathValleyTemp  
ARIMA(0,0,1)(0,1,1)[12]
```

```
Coefficients:
```

	ma1	sma1
	0.7460	-0.2453
s.e.	0.0315	0.0468

```
sigma^2 = 13.31: log likelihood = -933.06  
AIC=1872.12 AICc=1872.19 BIC=1883.64
```

```
modelMonthlySeasonal002.210DeathValley
```

```
Series: tsCaliforniaDeathValleyTemp  
ARIMA(0,0,2)(2,1,0)[12]
```

```
Coefficients:
```

	ma1	ma2	sar1	sar2
	0.9398	0.3858	-0.3944	0.0105
s.e.	0.0493	0.0432	0.0571	0.0564

```
sigma^2 = 10.94: log likelihood = -899.09
AIC=1808.17 AICc=1808.35 BIC=1827.38
```

```
modelMonthlySeasonal002.012DeathValley
```

```
Series: tsCaliforniaDeathValleyTemp
ARIMA(0,0,2)(0,1,2)[12]
```

```
Coefficients:
```

```
      ma1      ma2      sma1      sma2
      0.9326 0.3751 -0.3968 0.1196
s.e. 0.0497 0.0429 0.0581 0.0464
```

```
sigma^2 = 10.99: log likelihood = -899.77
AIC=1809.53 AICc=1809.71 BIC=1828.73
```

## Forecast

```
forecast(modelMonthlySeasonal100.011DeathValley, h = 8)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2012.9727	17.55612	13.98162	21.13061	12.089397	23.02284
2012.9754	17.24408	12.26961	22.21855	9.636279	24.85188
2012.9781	17.65015	11.65380	23.64650	8.479525	26.82078
2012.9809	18.51313	11.69716	25.32909	8.089008	28.93725
2012.9836	19.08324	11.58033	26.58615	7.608526	30.55795
2012.9863	19.63382	11.54013	27.72751	7.255590	32.01205
2012.9891	20.56363	11.95323	29.17402	7.395168	33.73209
2012.9918	19.34573	10.27799	28.41346	5.477826	33.21363

```
toPredictValues$`Death Valley`
```

```
[1] 31.7 23.9 20.0 19.4 20.6 21.7 22.2 21.1 21.7
```

```
# Create a new time series object with missing values replaced by zeros
tsCaliforniaDeathValleyTempFilled <- tsCaliforniaDeathValleyTemp
tsCaliforniaDeathValleyTempFilled[is.na(tsCaliforniaDeathValleyTempFilled)] <- 0
```

```

# Update the ARIMA model with the filled time series (missing values
# replaced by zeros)
updatedArimaModel <- Arima(tsCaliforniaDeathValleyTempFilled, model = modelMonthlySeasonal

# Generate predictions for the entire time series (including missing
# values)
predictionsDeath <- fitted(updatedArimaModel)

predictionsDeath

```

Time Series:

Start = c(2012, 1)

End = c(2012, 356)

Frequency = 366

```

[1] 20.57940 21.07890 20.57940 21.07890 21.67830 21.07890 21.07890 24.97500
[9] 23.27670 22.17780 21.07890 19.98000 21.67780 21.91276 19.34235 19.29272
[17] 19.66073 19.02549 17.75190 20.29710 16.41786 18.77081 22.02248 17.59681
[25] 20.17676 22.99799 22.70613 22.29451 25.79705 22.32645 21.78327 21.17008
[33] 23.34820 26.94421 20.51919 19.77519 27.09028 23.01569 19.74543 21.27679
[41] 25.56964 26.21471 24.89544 27.97729 25.14322 23.96421 19.10951 19.97846
[49] 26.95407 21.09615 20.86545 25.16594 22.83447 23.54662 29.48653 34.08597
[57] 27.39721 29.14605 20.88770 21.42029 23.44473 21.94290 22.76629 23.68910
[65] 25.92133 29.11928 28.44781 29.58285 22.60103 27.33867 26.24473 29.50642
[73] 32.30507 30.24799 29.49926 30.23966 31.69946 30.53573 22.60224 19.59023
[81] 22.38490 25.64981 27.42916 34.93076 33.11055 30.34658 28.17113 26.60402
[89] 30.88495 26.91566 28.26786 33.28171 34.41245 28.82895 27.38373 31.01342
[97] 35.94535 28.59734 23.58160 33.03561 34.30621 35.83783 35.91012 29.56991
[105] 28.12685 23.50038 21.48905 31.05075 32.48294 34.74233 35.69854 42.20036
[113] 41.57295 42.35389 43.20927 42.20428 37.69458 36.71038 33.81299 35.52635
[121] 37.17004 36.72093 39.97341 42.58311 37.99070 36.75480 34.61973 33.82180
[129] 34.89606 36.90426 40.38283 44.29562 42.65459 42.53250 41.71468 46.20519
[137] 44.81586 40.00266 42.71756 45.32715 38.36127 40.63868 44.40707 47.86842
[145] 45.39370 41.00407 39.33004 35.72385 32.56070 34.52381 37.86661 39.86655
[153] 43.47674 48.50080 49.04810 47.12830 46.03215 43.46255 35.60745 38.06966
[161] 42.44473 45.22769 40.97615 37.68632 43.53513 46.98340 45.02509 46.49370
[169] 44.80521 46.03423 47.00668 49.81370 47.07290 47.55562 46.63968 43.66865
[177] 43.29710 44.53699 43.49245 43.69545 45.13372 45.69657 45.65564 45.70559
[185] 45.50598 45.93583 44.73305 38.32792 44.00059 46.05867 48.88860 49.79899
[193] 50.19495 50.64656 51.96134 49.37196 40.16100 44.30633 42.47766 44.22876
[201] 43.15878 44.71694 46.81645 48.81883 49.24551 48.79430 47.38519 47.74684
[209] 48.14764 47.02178 44.32794 46.30712 48.12675 49.26908 46.38906 46.75240
[217] 47.77793 47.74030 46.42010 48.27225 50.17821 50.21823 49.99685 51.78223

```

[225]	51.39687	50.00217	49.55089	50.05117	48.95892	48.23090	47.74066	47.32373
[233]	48.09195	47.66252	46.93639	47.92482	40.51354	42.39486	44.69877	46.73566
[241]	45.23164	45.11023	47.16745	47.79472	48.63691	44.43631	42.24029	42.85343
[249]	46.10688	46.95477	46.48540	44.91232	45.90382	47.08419	47.27946	42.95377
[257]	41.23508	42.77252	43.47598	43.62004	45.00458	43.57614	43.56154	44.49870
[265]	44.93723	44.49885	43.09824	43.77751	45.07238	44.48395	39.31440	41.08747
[273]	42.21681	42.48347	43.11291	44.48890	45.52241	42.84377	42.16535	39.58126
[281]	38.03604	38.70174	37.42985	37.85888	37.37461	35.54309	29.15878	27.31569
[289]	29.13620	33.21339	33.91248	35.08105	38.12435	37.54424	35.38994	35.12673
[297]	33.84123	32.67165	28.36250	29.54453	29.07657	27.00955	29.24685	31.52308
[305]	32.76810	31.44265	31.63724	31.16621	31.60619	31.11109	33.19644	36.43114
[313]	35.80094	32.63552	26.78675	26.12133	27.23849	24.40382	24.57953	26.80508
[321]	29.48547	25.86305	25.31159	25.44446	21.98219	24.04989	26.23382	28.82801
[329]	25.30500	27.69377	26.49143	24.88879	23.23749	24.06981	27.92742	23.45428
[337]	26.53193	23.73810	20.02171	24.02118	20.71926	18.33711	18.30661	17.70819
[345]	19.58159	20.70958	16.17865	16.12332	15.36293	13.19632	17.27411	20.93727
[353]	17.27967	18.28266	18.73881	16.76495				

### 3 d)

The predictions could be improved by utilising a method more specific to predicting values before the end of our current data, as the forecast package is not the most optimized for this, unlike packages that utilize Kalman smoother method another way to improve the prediction could be to utilize dlms and see if they have a better performance than the arima function Lastly obtaining more data would increasing any model performance by having more data for each point in the cycle