

SpaceAndTime

Table of contents

TODO check if I checked all the residuals for all models	2
Question 1 Spatial modelling Kingdom of the Netherlands	2
1 a)	2
1 b)	3
1 c)	5
1 d)	7
1 e)	12
1 f)	30
1 g)	33
1 h)	35
Question 2	41
2 a)	41
Trend analysis	42
2 b)	42
ACF	42
PACF	43
best model selection	45
talk about the model being more easily explainability becaues $MA = 0$	50
Best model residual validation	50
Forecasting	51
2 c)	52
Initial assumptions	52
model fitting	53
Forecast	56
2 d)	58
2 e)	60
Model testing	63
Seasonal check	63
auto arima check	68


Forecasting	69
DLM	71
model fitting	71
Forecast	75
2 f)	78
Question 3	82
Question 3 a)	82
TODO WARNING	82
3 b)	84

TODO check if I checked all the residuals for all models

Question 1 Spatial modelling Kingdom of the Netherlands

1 a)

```
ggplot(data = netherlandsDF) + geom_point(aes(x = longitude, y = latitude,
size = precip, color = precip)) + scale_color_continuous(low = "blue",
high = "red") + labs(title = "Precipitation in Netherlands Stations",
x = "Longitude", y = "Latitude", size = "Precipitation", color = "Precipitation") +
theme_minimal()
```



From what we can see from the data it does seem to be spatially correlated as we can see that the Dutch provinces of north Holland, Friesland and Groningen have higher precipitation and as we go south the precipitation does decrease as we can see from the Dutch provinces of Zeeland, north Brabant and Limburg where precipitation is significantly lower than their northern counterparts.

From this data, latitude seems to be the biggest factor in the variation of the precipitation as the longitude only suggests some slight variations in the data.

```
# Create geodata object
precipitationNetherlands_geoR = as.geodata(netherlandsDF, coords.col = c("longitude",
  "latitude"), data.col = "precip")

summary(precipitationNetherlands_geoR)
```

Number of data points: 220

```
Coordinates summary
      longitude latitude
min       3.500   50.767
max       7.033   53.483
```

```
Distance summary
      min      max
0.001000 3.998498
```

```
Data summary
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
33.9000  80.8500 100.1500 106.5705 137.3500 185.6000
```

As we can see from the numerical summary of the data the median is different from the mean, which indicates it is not a symmetric distribution of data points and is instead positively skewed since the mean is bigger than the median. As such there are more values on the left side of the distribution.

1 b)

```
# set seed for reproducibility
set.seed(26041999)

# Select 3 random rows from the data frame
```

```

randomRowsPrecipitation = netherlandsDF %>%
  sample_n(3)

# Add a new column with labels
randomRowsPrecipitation$label = c("A", "B", "C")

# Print the randomly selected rows
randomRowsPrecipitation

```

```

# A tibble: 3 x 5
  station_name longitude latitude precip label
  <chr>          <dbl>    <dbl>  <dbl> <chr>
1 NIJKERK        5.47      52.2   89.1 A
2 WOLPHAARTSDIJK 3.73      51.5   95.9 B
3 EEXT           6.73      53    147. C

```

```

# Remove the selected rows from the original dataset
netherlandsDF_filtered = netherlandsDF %>%
  anti_join(randomRowsPrecipitation)

```

Joining, by = c("station_name", "longitude", "latitude", "precip")

```

# Print the resulting dataframe
netherlandsDF_filtered

```

```

# A tibble: 217 x 4
  station_name longitude latitude precip
  <chr>          <dbl>    <dbl>  <dbl>
1 WEST TERSCHELLING 5.22      53.4  130.
2 GRONINGEN-1      6.6       53.2  157.
3 HOORN            5.07      52.6  146.
4 HOOFDORP         4.7       52.3  130.
5 WINTERSWIJK      6.7       52.0   77.7
6 KERKWERVE        3.87      51.7   91.8
7 WESTDORPE-1      3.87      51.2   87.7
8 OUDENBOSCH       4.53      51.6   84.2
9 ROERMOND         5.97      51.2   56.4
10 PETTEN          4.65      52.8  158.
# ... with 207 more rows

```


```
## recreate the geoData object with the new filtered dataframe
precipitationNetherland_geoR = as.geodata(netherlandsDF_filtered, coords.col = c("longitud",
"latitude"), data.col = "precip")
```

1 c)

```
# # Calculate empirical variogram
variogramPrecipitationNetherlands = variog(precipitationNetherland_geoR)
```

variog: computing omnidirectional variogram

```
# Plot empirical variogram
plot(variogramPrecipitationNetherlands)
```



```
variogramPrecipitationNetherlands$n
```

```
[1] 1069 2482 3384 3834 3647 3221 2534 1541 787 468 318 133 17
```

From the plotted variogram we can see there a very clear need for a nugget as there is a non-zero value around zero distance, this values seems to be around 75 to 100 at the zero distance from how much is it decreasing.


The semi variance continuous to increase with distance till around the distance of 2 degrees distance wise, after this there is a decrease in variance that is not representative of the data as we are more and more uncertain the further we are from our known points, as such we will choose the distance of two as the cut off for the maximum distance.

we know change the maximum distance change and recut our previous variogram.

```
variogramPrecipitationNetherlands = variog(precipitationNetherland_geor,  
      option = "bin", max.dist = 2)
```

variog: computing omnidirectional variogram

```
plot(variogramPrecipitationNetherlands)
```



As we can see from the newly updated variogram the increase is almost linear with a curve near 0 where we can see the need for the nugget.

1 d)

Now that we have the variogram we will start by fitting a model to estimate the covariance via weighted least squares. Fitting this variogram we get the estimated values of σ^2 , ϕ and τ^2 also known as the nugget

We will first start with the default Matérn = 0,5 which is equivalent to an exponential as the form observed in the previous variogram seems to not fully linear and therefore require the curvature from a function like the exponential function to account for the behaviour at the near 0 distance.

From this we will try different models to search for the model with the best fit.

```
# ?variofit thau = nugget variability sigmasq = if the model can
# capture more or less of the total variability phi = if the
# correlation extends over a bigger or smaller distance loss value =
# goodness of fit (smaller means better fit)

krigingVariogramFittedDefault = variofit(variogramPrecipitationNetherlands,
  nugget = 85)
```

```
variofit: covariance model used is matern
variofit: weights used: npairs
variofit: minimisation function used: optim
```

```
Warning in variofit(variogramPrecipitationNetherlands, nugget = 85): initial
values not provided - running the default search
```

```
variofit: searching for best initial value ... selected values:
      sigmasq  phi  tausq kappa
initial.value "1984.67" "1.54" "85"  "0.5"
status        "est"    "est"  "est" "fix"
loss value: 818327620.928191
```

```
krigingVariogramFittedDefault
```

```
variofit: model parameters estimated by WLS (weighted least squares):
covariance model is: matern with fixed kappa = 0.5 (exponential)
parameter estimates:
      tausq    sigmasq      phi
      0.00 2561207.09 2590.45
```

Practical Range with $\text{cor}=0.05$ for asymptotic range: 7760.294

variofit: minimised weighted sum of squares = 35360382

Now we will increase the kappa of the Matérn to see if the increased flexibility and smoothness leads to a better fit

```
krigingVariogramFittedMatérn1.5 = variofit(variogramPrecipitationNetherlands,  
      kappa = 1.5, nugget = 85)
```

variofit: covariance model used is matern

variofit: weights used: npairs

variofit: minimisation function used: optim

Warning in variofit(variogramPrecipitationNetherlands, kappa = 1.5, nugget = 85): initial values not provided - running the default search

variofit: searching for best initial value ... selected values:

	sigmasq	phi	tausq	kappa
initial.value	"1984.67"	"0.62"	"85"	"1.5"
status	"est"	"est"	"est"	"fix"
loss value:	190482114.063677			

```
krigingVariogramFittedMatérn1.5
```

variofit: model parameters estimated by WLS (weighted least squares):

covariance model is: matern with fixed kappa = 1.5

parameter estimates:

tausq	sigmasq	phi
182.0638	3461.1492	1.1316

Practical Range with $\text{cor}=0.05$ for asymptotic range: 5.368367


variofit: minimised weighted sum of squares = 15435206

We will first visually compare these 2 models to see which one has a better

```
par(mar = c(4, 4, 2, 2))  
plot(variogramPrecipitationNetherlands, pch = 19)  
lines(krigingVariogramFittedDefault)
```



```
lines(krigingVariogramFittedMatrén1.5, lty = 2)
```



```
# lines(krigingVariogramFittedMatrén2.5, lty = 3)
```

Immediately we can see that the extra flexibility of the Matérn 1,5 not only better follows the actual data, it actually accounts correctly for the initial variance from the nugget which the Matérn 0,5 does not as it simply decreases to 0.

We now will test if any additional flexibility changes can improve the model fit

We will first try to again increase the kappa to see if the model again benefits from the extra flexibility

```
krigingVariogramFittedMatrén2.0 = variofit(variogramPrecipitationNetherlands,
      kappa = 2, nugget = 85)
```

```
variofit: covariance model used is matern
variofit: weights used: npairs
variofit: minimisation function used: optim
```

```
Warning in variofit(variogramPrecipitationNetherlands, kappa = 2, nugget = 85):
initial values not provided - running the default search
```

```

variofit: searching for best initial value ... selected values:
           sigmasq  phi   tausq   kappa
initial.value "1984.67" "0.62" "198.47" "2"
status        "est"    "est"  "est"   "fix"
loss value: 227233080.389154

```

```
krigingVariogramFittedMatrén2.0
```

```

variofit: model parameters estimated by WLS (weighted least squares):
covariance model is: matern with fixed kappa = 2
parameter estimates:
      tausq  sigmasq      phi
197.6127 3037.3813   0.8386
Practical Range with cor=0.05 for asymptotic range: 4.502076

```

```
variofit: minimised weighted sum of squares = 18197964
```

Here we will instead see if the model will benefit instead from a cut of flexibility to make it less smooth

```
krigingVariogramFittedMatrén1.0 = variofit(variogramPrecipitationNetherlands,
      kappa = 1, nugget = 85)
```

```

variofit: covariance model used is matern
variofit: weights used: npairs
variofit: minimisation function used: optim

```

```
Warning in variofit(variogramPrecipitationNetherlands, kappa = 1, nugget = 85):
initial values not provided - running the default search
```

```

variofit: searching for best initial value ... selected values:
           sigmasq  phi   tausq   kappa
initial.value "1984.67" "0.92" "198.47" "1"
status        "est"    "est"  "est"   "fix"
loss value: 352001017.756763

```

```
krigingVariogramFittedMatrén1.0
```


```
variofit: model parameters estimated by WLS (weighted least squares):  
covariance model is: matern with fixed kappa = 1  
parameter estimates:
```

```
    tausq  sigmasq    phi  
146.1354 4832.1862  2.0470
```

```
Practical Range with cor=0.05 for asymptotic range: 8.185098
```

```
variofit: minimised weighted sum of squares = 12019899
```

```
par(mar = c(4, 4, 2, 2))  
plot(variogramPrecipitationNetherlands, pch = 19)  
lines(krigingVariogramFittedMatrén1.5)  
lines(krigingVariogramFittedMatrén1.0, lty = 2)  
lines(krigingVariogramFittedMatrén2.0, lty = 3)
```



As we can see from the new graph it does seem that actually a lower flexibility Matérn has a better fit since the extra flexibility near the start and end of the data points made the models deviate too much from the points.

1 e)

To fit a model using the maximum likelihood we will have to try multiple initial values to make sure this is indeed the maximum likelihood and not just a local maximum.

```
# ?likest
```

```
maximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(10, 1))
```

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
  arguments for the maximisation function.  
  For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
  times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

```
maximumLikelihoodNetherlandsInitial1.10 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(1, 10))
```

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
  arguments for the maximisation function.  
  For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
  times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

```
maximumLikelihoodNetherlandsInitial100.10 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(100, 10))
```

```
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
maximumLikelihoodNetherlandsInitial10.100 = likfit(precipitationNetherlands_geoR,
  ini.cov.pars = c(10, 100))
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

WARNING: estimated range is more than 10 times bigger than the biggest distance between two points

- 1) excluding spatial dependence if estimated sill is too low and/or
- 2) taking trends (covariates) into account

```
maximumLikelihoodNetherlandsInitial1.1 = likfit(precipitationNetherlands_geoR,
  ini.cov.pars = c(1, 1))
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
maximumLikelihoodNetherlandsInitial1000.1000 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(1000, 1000))
```

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
       arguments for the maximisation function.  
       For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
       times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

WARNING: estimated range is more than 10 times bigger than the biggest distance between two
1) excluding spatial dependence if estimated sill is too low and/or
2) taking trends (covariates) into account

```
maximumLikelihoodNetherlandsInitial500.500 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(500, 500))
```

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
       arguments for the maximisation function.  
       For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
       times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

WARNING: estimated range is more than 10 times bigger than the biggest distance between two
1) excluding spatial dependence if estimated sill is too low and/or
2) taking trends (covariates) into account

```
maximumLikelihoodNetherlandsInitial10.1
```

```
likfit: estimated model parameters:
```

```

      beta      tausq      sigmasq      phi
" 102.376" " 114.249" "3036.627" "    7.132"
Practical Range with cor=0.05 for asymptotic range: 21.36583

```

```
likfit: maximised log-likelihood = -896.9
```

```
maximumLikelihoodNetherlandsInitial1.10
```

```
likfit: estimated model parameters:
      beta      tausq      sigmasq      phi
" 102.449" " 114.921" "4184.402" "    9.991"
Practical Range with cor=0.05 for asymptotic range: 29.9294

```

```
likfit: maximised log-likelihood = -896.9
```

```
maximumLikelihoodNetherlandsInitial100.10
```

```
likfit: estimated model parameters:
      beta      tausq      sigmasq      phi
" 102.449" " 114.921" "4184.402" "    9.991"
Practical Range with cor=0.05 for asymptotic range: 29.9294

```

```
likfit: maximised log-likelihood = -896.9
```

```
maximumLikelihoodNetherlandsInitial10.100
```

```
likfit: estimated model parameters:
      beta      tausq      sigmasq      phi
" 102.7" " 117.7" "39625.5" " 100.0"
Practical Range with cor=0.05 for asymptotic range: 299.5729

```

```
likfit: maximised log-likelihood = -897.8
```

```
maximumLikelihoodNetherlandsInitial1.1
```

```
likfit: estimated model parameters:
      beta      tausq      sigmasq      phi

```

```
" 102.376" " 114.249" "3036.627" " 7.132"
Practical Range with cor=0.05 for asymptotic range: 21.36583
```

```
likfit: maximised log-likelihood = -896.9
```

```
maximumLikelihoodNetherlandsInitial1000.1000
```

```
likfit: estimated model parameters:
```

```
      beta      tausq      sigmasq      phi
" 103.2" " 148.3" "246755.1" " 1000.0"
```

```
Practical Range with cor=0.05 for asymptotic range: 2995.732
```

```
likfit: maximised log-likelihood = -900.9
```

```
maximumLikelihoodNetherlandsInitial500.500
```

```
likfit: estimated model parameters:
```

```
      beta      tausq      sigmasq      phi
" 103.1" " 145.0" "130046.9" " 500.0"
```

```
Practical Range with cor=0.05 for asymptotic range: 1497.866
```

```
likfit: maximised log-likelihood = -900.1
```

As we can see from the maximised log-likelihoods it does seem that have indeed reached the maximum log-likelihood as none of the values are too fastly different and the likelihood worsedned as we started to increase much more our starting values.

Next we will try the REML that takes into account the fact that some of the parameters of the model are related to the variance of the residuals and not the mean.

```
REMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherland_geoR,
  ini.cov.pars = c(10, 1), lik.method = "REML")
```

```
-----
likfit: likelihood maximisation using the function optim.
```

```
likfit: Use control() to pass additional
```

```
arguments for the maximisation function.
```

```
For further details see documentation for optim.
```

```
likfit: It is highly advisable to run this function several
```


times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!

likfit: end of numerical maximisation.

WARNING: estimated range is more than 10 times bigger than the biggest distance between two
1) excluding spatial dependence if estimated sill is too low and/or
2) taking trends (covariates) into account

```
REMLmaximumLikelihoodNetherlandsInitial1.10 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(1, 10), lik.method = "REML")
```

likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
arguments for the maximisation function.
For further details see documentation for optim.
likfit: It is highly advisable to run this function several
times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!

likfit: end of numerical maximisation.

WARNING: estimated range is more than 10 times bigger than the biggest distance between two
1) excluding spatial dependence if estimated sill is too low and/or
2) taking trends (covariates) into account

```
REMLmaximumLikelihoodNetherlandsInitial100.1 = likfit(precipitationNetherland_geoR,  
  ini.cov.pars = c(100, 1), lik.method = "REML")
```

likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
arguments for the maximisation function.
For further details see documentation for optim.
likfit: It is highly advisable to run this function several
times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!

likfit: end of numerical maximisation.

WARNING: estimated range is more than 10 times bigger than the biggest distance between two
 1) excluding spatial dependence if estimated sill is too low and/or
 2) taking trends (covariates) into account

```
REMLmaximumLikelihoodNetherlandsInitial1.100 = likfit(precipitationNetherlands_geoR,
  ini.cov.pars = c(1, 100), lik.method = "REML")
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

WARNING: estimated range is more than 10 times bigger than the biggest distance between two
 1) excluding spatial dependence if estimated sill is too low and/or
 2) taking trends (covariates) into account

```
REMLmaximumLikelihoodNetherlandsInitial10.1
```

```
likfit: estimated model parameters:
      beta      tausq      sigmasq      phi
" 102.61" " 114.60" "18253.12" " 43.35"
Practical Range with cor=0.05 for asymptotic range: 129.8698

likfit: maximised log-likelihood = -888.9
```

```
REMLmaximumLikelihoodNetherlandsInitial1.10
```

```
likfit: estimated model parameters:
      beta      tausq      sigmasq      phi
" 102.62" " 114.83" "18124.93" " 43.19"
Practical Range with cor=0.05 for asymptotic range: 129.3933

likfit: maximised log-likelihood = -888.9
```

```
REMLmaximumLikelihoodNetherlandsInitial100.1
```

```
likfit: estimated model parameters:
```

```
      beta      tausq      sigmasq      phi  
" 102.61" " 114.60" "18253.12" " 43.35"
```

```
Practical Range with cor=0.05 for asymptotic range: 129.8698
```

```
likfit: maximised log-likelihood = -888.9
```

```
REMLmaximumLikelihoodNetherlandsInitial1.100
```

```
likfit: estimated model parameters:
```

```
      beta      tausq      sigmasq      phi  
" 102.7" " 116.5" "40885.8" " 100.0"
```

```
Practical Range with cor=0.05 for asymptotic range: 299.5729
```

```
likfit: maximised log-likelihood = -888.9
```

As we can see from these new models, the new likelihood method actually did improve our model as we have a lower log-likelihood.

Since the data did not seem to be perfectly stationary as seen in the previous questions, we will now check if adding a linear trend improves our model

```
linearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherland_geoR,  
  trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML")
```

```
-----  
likfit: likelihood maximisation using the function optim.
```

```
likfit: Use control() to pass additional
```

```
  arguments for the maximisation function.
```

```
  For further details see documentation for optim.
```

```
likfit: It is highly advisable to run this function several
```

```
  times with different initial values for the parameters.
```

```
likfit: WARNING: This step can be time demanding!
```

```
-----  
likfit: end of numerical maximisation.
```

```
linearREMLmaximumLikelihoodNetherlandsInitial1.10 = likfit(precipitationNetherland_geoR,  
  trend = "1st", ini.cov.pars = c(1, 10), lik.method = "REML")
```

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
       arguments for the maximisation function.  
       For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
       times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

```
linearREMLmaximumLikelihoodNetherlandsInitial100.10 = likfit(precipitationNetherland_geoR,  
  trend = "1st", ini.cov.pars = c(100, 10), lik.method = "REML")
```

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
       arguments for the maximisation function.  
       For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
       times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

```
linearREMLmaximumLikelihoodNetherlandsInitial10.100 = likfit(precipitationNetherland_geoR,  
  trend = "1st", ini.cov.pars = c(10, 100), lik.method = "REML")
```

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
       arguments for the maximisation function.  
       For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
       times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!
```

likfit: end of numerical maximisation.

WARNING: estimated range is more than 10 times bigger than the biggest distance between two
1) excluding spatial dependence if estimated sill is too low and/or
2) taking trends (covariates) into account

`linearREMLmaximumLikelihoodNetherlandsInitial10.1`

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2194.6747" "	-11.4587" "	45.2352" "	112.7395" "	191.5497" "	0.4033"

Practical Range with cor=0.05 for asymptotic range: 1.208203

likfit: maximised log-likelihood = -872

`linearREMLmaximumLikelihoodNetherlandsInitial1.10`

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2084.141" "	-6.973" "	42.674" "	125.053" "	3238.867" "	9.905"

Practical Range with cor=0.05 for asymptotic range: 29.67433

likfit: maximised log-likelihood = -873

`linearREMLmaximumLikelihoodNetherlandsInitial100.10`

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2084.141" "	-6.973" "	42.674" "	125.053" "	3238.867" "	9.905"

Practical Range with cor=0.05 for asymptotic range: 29.67433

likfit: maximised log-likelihood = -873

`linearREMLmaximumLikelihoodNetherlandsInitial10.100`

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2084.09"	"-6.91"	"42.67"	"125.89"	"32191.57"	"100.00"

Practical Range with cor=0.05 for asymptotic range: 299.5701

likfit: maximised log-likelihood = -873

From these models we can see that the linear trend does indeed improve our model, now we will check if there are any other covariance functions that can improve the model further.

```
Matren0.5linearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherlan
  trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML", cov.model = "matern",
  kappa = 0.5)
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
Matren1.0linearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherlan
  trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML", cov.model = "matern",
  kappa = 1)
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
Matren1.5linearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherlan
  trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML", cov.model = "matern",
  kappa = 1.5)
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
Matren2.0linearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherlan
  trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML", cov.model = "matern",
  kappa = 2)
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
likfit: It is highly advisable to run this function several
      times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
-----
likfit: end of numerical maximisation.
```

```
Matren2.5linearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherlan
  trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML", cov.model = "matern",
  kappa = 2.5)
```

```
-----
likfit: likelihood maximisation using the function optim.
likfit: Use control() to pass additional
      arguments for the maximisation function.
      For further details see documentation for optim.
```

```
likfit: It is highly advisable to run this function several
        times with different initial values for the parameters.
likfit: WARNING: This step can be time demanding!
```

```
-----
likfit: end of numerical maximisation.
```

```
Matren0.5linearREMLmaximumLikelihoodNetherlandsInitial10.1
```

```
likfit: estimated model parameters:
```

beta0	beta1	beta2	tausq	sigmasq	phi
"-2194.6747"	" -11.4587"	" 45.2352"	" 112.7395"	" 191.5497"	" 0.4033"

Practical Range with cor=0.05 for asymptotic range: 1.208203

```
likfit: maximised log-likelihood = -872
```

```
Matren1.0linearREMLmaximumLikelihoodNetherlandsInitial10.1
```

```
likfit: estimated model parameters:
```

beta0	beta1	beta2	tausq	sigmasq	phi
"-2221.5117"	" -12.1627"	" 45.8184"	" 125.4406"	" 160.1807"	" 0.2088"

Practical Range with cor=0.05 for asymptotic range: 0.8347704

```
likfit: maximised log-likelihood = -871.3
```

```
Matren1.5linearREMLmaximumLikelihoodNetherlandsInitial10.1
```

```
likfit: estimated model parameters:
```

beta0	beta1	beta2	tausq	sigmasq	phi
"-2233.7541"	" -12.4219"	" 46.0784"	" 129.6185"	" 149.9124"	" 0.1529"

Practical Range with cor=0.05 for asymptotic range: 0.7254113

```
likfit: maximised log-likelihood = -871.1
```

```
Matren2.0linearREMLmaximumLikelihoodNetherlandsInitial10.1
```

```
likfit: estimated model parameters:
```

beta0	beta1	beta2	tausq	sigmasq	phi
"-2240.7056"	" -12.5595"	" 46.2252"	" 131.4685"	" 144.9394"	" 0.1248"

Practical Range with cor=0.05 for asymptotic range: 0.669996

likfit: maximised log-likelihood = -871.1

```
Matren2.5linearREMLmaximumLikelihoodNetherlandsInitial10.1
```

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2245.1703"	" -12.6453"	" 46.3192"	" 132.4280"	" 142.0693"	" 0.1074"

Practical Range with cor=0.05 for asymptotic range: 0.6358818

likfit: maximised log-likelihood = -871

It does seem that the matrén covariance function did indeed slightly improved the model so we will compare it to a model using a spherical covariance function. The spherical covariance function is appropriate for this scenario has the spatial correlation between data points decreases rapidly as the distance between the points increases and we are limited with the range of correlation has after 2 degrees of distance we loose sensible correlation, hence the cut in the variogram.

```
SphericallinearREMLmaximumLikelihoodNetherlandsInitial10.1 = likfit(precipitationNetherlan  
trend = "1st", ini.cov.pars = c(10, 1), lik.method = "REML", cov.model = "spherical")
```

kappa not used for the spherical correlation function

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
arguments for the maximisation function.  
For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

```
SphericallinearREMLmaximumLikelihoodNetherlandsInitial1.10 = likfit(precipitationNetherlan  
trend = "1st", ini.cov.pars = c(1, 10), lik.method = "REML", cov.model = "spherical")
```

kappa not used for the spherical correlation function

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
        arguments for the maximisation function.  
        For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
        times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

```
SphericallinearREMLmaximumLikelihoodNetherlandsInitial100.10 = likfit(precipitationNetherl  
trend = "1st", ini.cov.pars = c(100, 10), lik.method = "REML", cov.model = "spherical")
```

kappa not used for the spherical correlation function

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
        arguments for the maximisation function.  
        For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
        times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----  
likfit: end of numerical maximisation.
```

```
SphericallinearREMLmaximumLikelihoodNetherlandsInitial10.100 = likfit(precipitationNetherl  
trend = "1st", ini.cov.pars = c(10, 100), lik.method = "REML", cov.model = "spherical")
```

kappa not used for the spherical correlation function

```
-----  
likfit: likelihood maximisation using the function optim.  
likfit: Use control() to pass additional  
        arguments for the maximisation function.  
        For further details see documentation for optim.  
likfit: It is highly advisable to run this function several  
        times with different initial values for the parameters.  
likfit: WARNING: This step can be time demanding!  
-----
```

likfit: end of numerical maximisation.

WARNING: estimated range is more than 10 times bigger than the biggest distance between two

- 1) excluding spatial dependence if estimated sill is too low and/or
- 2) taking trends (covariates) into account

SphericallinearREMLmaximumLikelihoodNetherlandsInitial10.1

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2182.422"	" -11.539"	" 45.002"	" 127.066"	" 200.406"	" 1.003"

Practical Range with cor=0.05 for asymptotic range: 1.00273

likfit: maximised log-likelihood = -872.2

SphericallinearREMLmaximumLikelihoodNetherlandsInitial1.10

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2082.850"	" -6.686"	" 42.622"	" 125.174"	" 2151.559"	" 9.905"

Practical Range with cor=0.05 for asymptotic range: 9.905214

likfit: maximised log-likelihood = -873

SphericallinearREMLmaximumLikelihoodNetherlandsInitial100.10

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2082.850"	" -6.686"	" 42.622"	" 125.174"	" 2151.559"	" 9.905"

Practical Range with cor=0.05 for asymptotic range: 9.905214

likfit: maximised log-likelihood = -873

SphericallinearREMLmaximumLikelihoodNetherlandsInitial10.100

likfit: estimated model parameters:

beta0	beta1	beta2	tausq	sigmasq	phi
"-2083.890"	" -6.902"	" 42.667"	" 125.695"	"21538.008"	" 99.999"

Practical Range with cor=0.05 for asymptotic range: 99.99893

likfit: maximised log-likelihood = -873

As we can see the spherical covariance function does not provide as good of a fit as the Matérn.

We will now validate the model by doing cross-validation on the model.

```
xv.ml = xvalid(precipitationNetherland_geoR, model = Matren2.5linearREMLmaximumLikelihoodM)
```

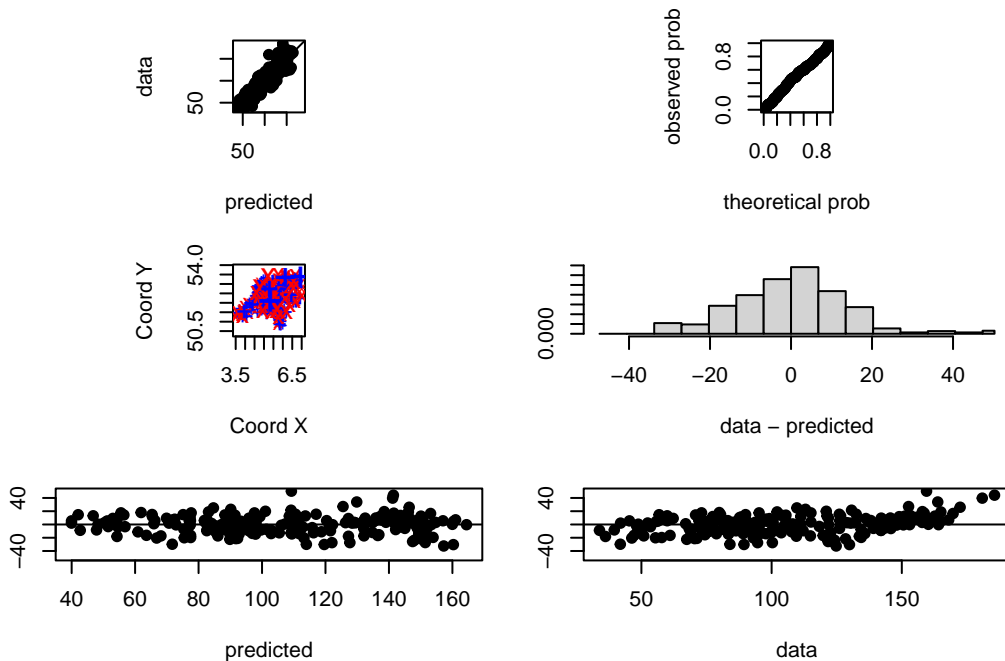
xvalid: number of data locations = 217

xvalid: number of validation locations = 217

xvalid: performing cross-validation at location ... 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13

xvalid: end of cross-validation

```
par(mfrow = c(3, 2), mar = c(4, 2, 2, 2))
plot(xv.ml, error = TRUE, std.error = FALSE, pch = 19)
```



From these plots we can see that the residuals seem mostly normal without any quickly identifiable patterns or bias.


From the first top left graph we can see however that we seem to slightly underestimate more data points.

To account for bias we will also perform cross-validation to the next best performing model using the spherical function instead

```
xv.ml = xvalid(precipitationNetherland_geoR, model = SphericallylinearREMLmaximumLikelihoodN
```

```
xvalid: number of data locations      = 217
xvalid: number of validation locations = 217
xvalid: performing cross-validation at location ... 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13
xvalid: end of cross-validation
```

```
par(mfrow = c(3, 2), mar = c(4, 2, 2, 2))
plot(xv.ml, error = TRUE, std.error = FALSE, pch = 19)
```



As we can see the data at the start seems to be systematically underestimated and at the end it seems to be overestimated. Furthermore the theoretical data plot seems to be less linear.

This confirms that the spherical covariance function is indeed a worse model than the Matérn model.

1 f)

First we will start by making the predictions using the variogram

```
spatialPointsABC = randomRowsPrecipitation[, c("longitude", "latitude")]

# Set the kriging control parameters
krigeControl = krigingVariogramFittedMatrén1.0$krige.control(type.krige = "OK", cov.model = krigingVariogramFittedMatrén1.0$cov.pars)

# Kriging with the fitted variogram model
krigeResults = krigingVariogramFittedMatrén1.0$krige.conv(precipitationNetherlands_geoR, locations = spatialPointsABC,
  krigingVariogramFittedMatrén1.0$krigeControl)
```

krige.conv: model with constant mean

krige.conv: Kriging performed using global neighbourhood

```
# Extract predictions from the kriging results
predictions = krigeResults$predict

# Compare the predicted values with the actual precipitation values

actualPrecipitationValues = randomRowsPrecipitation[, 4]

comparisonvariogram = data.frame(actualPrecipitationValues, predictions)
```

Now for the maximum likelihood function

```
# done above spatialPointsABC = randomRowsPrecipitation[,
# c('longitude', 'latitude')]

# Set the kriging control parameters
krigeControl = krigingVariogramFittedMatrén1.0$krige.control(type.krige = "OK", cov.model = Matren2.5linearREMLmaximumLikelihoodNetherlandsInitial10.1$cov.pars)
cov.pars = Matren2.5linearREMLmaximumLikelihoodNetherlandsInitial10.1$cov.pars)
```

```
# Kriging with the fitted variogram model
krigeResults = krige.conv(precipitationNetherlands_geoR, locations = spatialPointsABC,
  krige = krigeControl)
```

krige.conv: model with constant mean
krige.conv: Kriging performed using global neighbourhood

```
# Extract predictions from the kriging results
predictions = krigeResults$predict

# Compare the predicted values with the actual precipitation values

# done above actualPrecipitationValues = randomRowsPrecipitation[,4]
comparisonMaximumLikelihood = data.frame(actualPrecipitationValues, predictions)
```

Now that we have made the predictions for our 2 models we will check the predicted values compared to the real values for each of the models.

```
comparisonvariogram
```

```
precip predictions
1  89.1    95.41904
2  95.9    98.90558
3 147.2   147.82038
```

```
comparisonMaximumLikelihood
```

```
precip predictions
1  89.1    99.00780
2  95.9    99.36194
3 147.2   140.07734
```

```
# Calculate Mean Absolute Error (MAE)
MAEVariogram = mean(abs(comparisonvariogram$precip - comparisonvariogram$predictions))
MAEMaximumLikelihood = mean(abs(comparisonMaximumLikelihood$precip - comparisonMaximumLike
```

```

# Calculate Mean Squared Error (MSE)
mseVariogram = mean((comparisonvariogram$precip - comparisonvariogram$predictions)^2)
mseMaximumLikelihood = mean((comparisonMaximumLikelihood$precip - comparisonMaximumLikelihood$predictions)^2)

# Calculate Root Mean Squared Error (RMSE)
rmseVariogram = sqrt(mseVariogram)
rmseMaximumLikelihood = sqrt(mseMaximumLikelihood)

# Display the calculated metrics
cat("Mean Absolute Error (MAE) of the Variogram:", MAEVariogram, "\n")

```

Mean Absolute Error (MAE) of the Variogram: 3.315002

```

cat("Mean Squared Error (MSE) of the Variogram:", mseVariogram, "\n")

```

Mean Squared Error (MSE) of the Variogram: 16.44957

```

cat("Root Mean Squared Error (RMSE) of the Variogram:", rmseVariogram, "\n")

```

Root Mean Squared Error (RMSE) of the Variogram: 4.055807

```

cat("\n\n")

```

```

cat("Mean Absolute Error (MAE) of the maximum likelihood:", MAEMaximumLikelihood,
    "\n")

```

Mean Absolute Error (MAE) of the maximum likelihood: 6.830801

```

cat("Mean Squared Error (MSE) of the maximum likelihood:", mseMaximumLikelihood,
    "\n")

```

Mean Squared Error (MSE) of the maximum likelihood: 53.62729


```
cat("Root Mean Squared Error (RMSE) of the maximum likelihood:", rmseMaximumLikelihood,
    "\n")
```

Root Mean Squared Error (RMSE) of the maximum likelihood: 7.323066

As we can see from both the real values and the MAE, MSE and RMSE the variogram has as much better performance predicting those 3 points than our maximum likelihood model

1 g)

```
# Determine the range of the coordinates
xRange = range(precipitationNetherland_geoR$coords[, 1])
yRange = range(precipitationNetherland_geoR$coords[, 2])

# Create a grid with 0.05-degree spacing
gridPoints = expand.grid(x = seq(xRange[1], xRange[2], by = 0.05), y = seq(yRange[1],
    yRange[2], by = 0.05))

# Kriging with the fitted variogram model
krigeResults = krige.conv(precipitationNetherland_geoR, locations = gridPoints,
    krige = krigeControl)
```

krige.conv: model with constant mean

krige.conv: Kriging performed using global neighbourhood

```
# Create a data frame for the grid points with the predicted mean and
# variance
gridData = data.frame(gridPoints, mean = krigeResults$predict, variance = krigeResults$kri

# Mean plot
meanPlot = ggplot(gridData, aes(x = x, y = y, fill = mean)) + geom_tile() +
    scale_fill_gradientn(colors = c("blue", "green", "yellow", "red")) +
    theme_minimal() + ggtitle("Mean Plot") + labs(x = "Longitude", y = "Latitude",
    fill = "Mean")

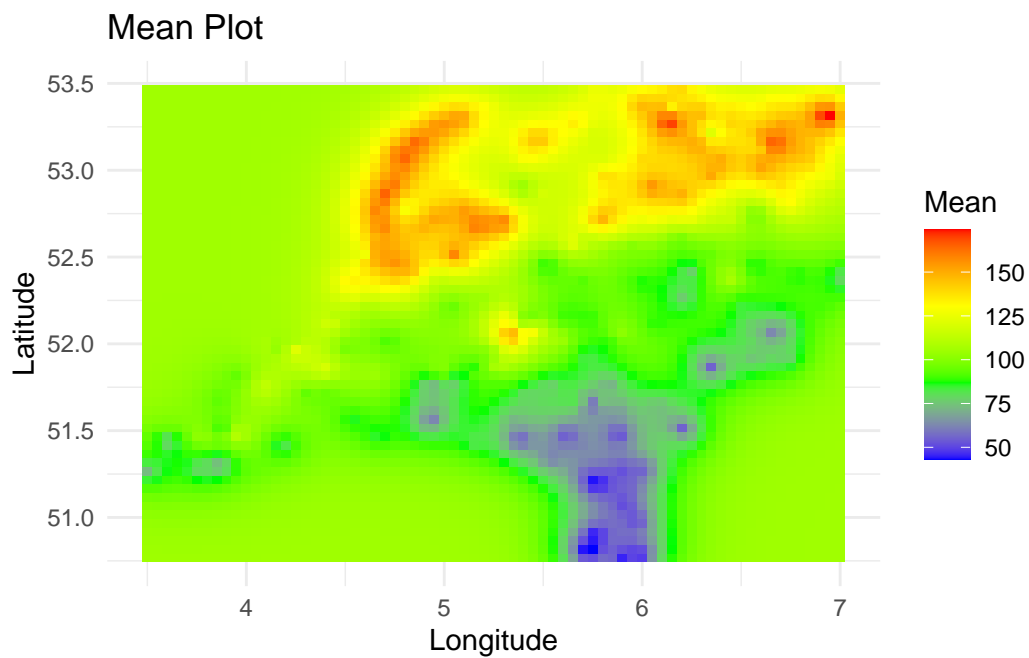
# Variance plot
variancePlot = ggplot(gridData, aes(x = x, y = y, fill = variance)) + geom_tile() +
```

```

scale_fill_gradientn(colors = c("white", "blue", "green", "yellow", "red")) +
theme_minimal() + ggtitle("Variance Plot") + labs(x = "Longitude", y = "Latitude",
fill = "Variance")

# Display the plots
print(meanPlot)


```



```

print(variancePlot)

```



1 h)

For the priors I will be using the estimated values from our latest maximum likelihood model.

```
# Extract the estimated parameters

priorPhiVariogram = krigingVariogramFittedMatrén1.0$cov.pars[1]
priorTauSQVariogram = krigingVariogramFittedMatrén1.0$cov.pars[2]

priorPhiMaximumLikelihood = Matren1.0linearREMLmaximumLikelihoodNetherlandsInitial10.1$cov
priorTauSQMaximumLikelihood = Matren1.0linearREMLmaximumLikelihoodNetherlandsInitial10.1$cov
```

The function does not support continuous priors directly so we will fit them as discrete priors.

```
# Creating discrete priors for phi and tau^2_rel AKA this is for the
# model with a nugget
phiDiscrete <- seq(min(priorPhiVariogram, priorPhiMaximumLikelihood) * 0.5,
  max(priorPhiVariogram, priorPhiMaximumLikelihood) * 1.5, length.out = 50)

tauSqDiscrete <- seq(min(priorTauSQVariogram, priorTauSQMaximumLikelihood) *
```



```
krigeBayesModelWithoutNugget <- krige.bayes(geodata = precipitationNetherlands_geoR,
      loc = ex.grid, prior = prior.control(phi.prior = phiProbability, phi.discrete = phiDis
```

krige.bayes: model with constant mean

krige.bayes: computing the discrete posterior of phi/tausq.rel

krige.bayes: computing the posterior probabilities.

Number of parameter sets: 50

1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26

krige.bayes: sampling from posterior distribution

krige.bayes: sample from the (joint) posterior of phi and tausq.rel

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	
phi	80.09036	226.3799	372.6695	518.9591	665.2486	811.5382	957.8278	
tausq.rel	0.00000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	
frequency	16.00000	20.0000	14.0000	31.0000	23.0000	30.0000	23.0000	
	[,8]	[,9]	[,10]	[,11]	[,12]	[,13]	[,14]	
phi	1104.117	1250.407	1396.696	1542.986	1689.276	1835.565	1981.855	
tausq.rel	0.000	0.000	0.000	0.000	0.000	0.000	0.000	
frequency	24.000	21.000	28.000	21.000	27.000	28.000	31.000	
	[,15]	[,16]	[,17]	[,18]	[,19]	[,20]	[,21]	
phi	2128.144	2274.434	2420.723	2567.013	2713.303	2859.592	3005.882	
tausq.rel	0.000	0.000	0.000	0.000	0.000	0.000	0.000	
frequency	46.000	32.000	37.000	28.000	26.000	26.000	28.000	
	[,22]	[,23]	[,24]	[,25]	[,26]	[,27]	[,28]	[,29]
phi	3152.171	3298.461	3444.75	3591.04	3737.33	3883.619	4029.909	4176.198
tausq.rel	0.000	0.000	0.00	0.00	0.00	0.000	0.000	0.000
frequency	30.000	37.000	20.00	18.00	23.00	14.000	24.000	20.000
	[,30]	[,31]	[,32]	[,33]	[,34]	[,35]	[,36]	
phi	4322.488	4468.777	4615.067	4761.357	4907.646	5053.936	5200.225	
tausq.rel	0.000	0.000	0.000	0.000	0.000	0.000	0.000	
frequency	21.000	19.000	16.000	19.000	19.000	20.000	20.000	
	[,37]	[,38]	[,39]	[,40]	[,41]	[,42]	[,43]	
phi	5346.515	5492.804	5639.094	5785.384	5931.673	6077.963	6224.252	
tausq.rel	0.000	0.000	0.000	0.000	0.000	0.000	0.000	
frequency	17.000	14.000	8.000	17.000	5.000	12.000	9.000	
	[,44]	[,45]	[,46]	[,47]	[,48]	[,49]	[,50]	
phi	6370.542	6516.831	6663.121	6809.411	6955.7	7101.99	7248.279	
tausq.rel	0.000	0.000	0.000	0.000	0.0	0.00	0.000	
frequency	9.000	8.000	7.000	5.000	2.0	3.00	4.000	

krige.bayes: starting prediction at the provided locations

krige.bayes: phi/tausq.rel samples for the predictive are same as for the posterior

```
krige.bayes: computing moments of the predictive distribution
krige.bayes: sampling from the predictive
      Number of parameter sets: 50
1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26
krige.bayes: preparing summaries of the predictive distribution
```

```
summary(krigeBayesModelWithoutNugget$posterior$sample)
```

beta	sigmasq	phi	tausq.rel
Min. : -14806.5	Min. : 375543	Min. : 80.09	Min. : 0
1st Qu.: -2113.0	1st Qu.: 8331936	1st Qu.: 1542.99	1st Qu.: 0
Median : 290.8	Median : 14519487	Median : 2713.30	Median : 0
Mean : 299.1	Mean : 15700517	Mean : 2965.95	Mean : 0
3rd Qu.: 2547.6	3rd Qu.: 22223235	3rd Qu.: 4322.49	3rd Qu.: 0
Max. : 16575.4	Max. : 44053886	Max. : 7248.28	Max. : 0


```
summary(krigeBayesModelWithNugget$posterior$sample)
```

beta	sigmasq	phi	tausq.rel
Min. : -49.52	Min. : 2216	Min. : 80.09	Min. : 0.1044
1st Qu.: 66.57	1st Qu.: 3000	1st Qu.: 80.09	1st Qu.: 0.1044
Median : 106.16	Median : 3207	Median : 80.09	Median : 0.1044
Mean : 105.66	Mean : 3238	Mean : 80.09	Mean : 0.1044
3rd Qu.: 145.57	3rd Qu.: 3433	3rd Qu.: 80.09	3rd Qu.: 0.1044
Max. : 282.96	Max. : 4455	Max. : 80.09	Max. : 0.1044

Now we will compare the posterior of both of the models to see the impact of the nugget


```
# Extract posterior samples for phi and tau^2_rel
posterior_samples_model1 <- krigeBayesModelWithNugget$posterior$phi$phi.marginal$phi
posterior_samples_model2 <- krigeBayesModelWithoutNugget$posterior$phi$phi.marginal$phi

# Plot the posterior distributions
hist(posterior_samples_model1, freq = FALSE, main = "Posterior Distributions for phi (Model 1)",
      xlab = "phi", col = "blue", xlim = range(c(posterior_samples_model1,
          posterior_samples_model2)))
```



```
hist(posterior_samples_model2, freq = FALSE, main = "Posterior Distributions for phi (Model 1)",  
      xlab = "phi", col = "blue", xlim = range(c(posterior_samples_model1,  
          posterior_samples_model2)))
```

Posterior Distributions for phi (Model 2)



```
# Compare summary statistics
summary_model1 <- summary(posterior_samples_model1)
summary_model2 <- summary(posterior_samples_model2)

cat("model 1 :\n")
```

model 1 :

```
summary_model1
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
80.09	1872.14	3664.18	3664.18	5456.23	7248.28

```
cat("\n\n model 2 :\n")
```

model 2 :


```
summary_model2
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
80.09	1872.14	3664.18	3664.18	5456.23	7248.28

```
# Reset the plot layout  
par(mfrow = c(1, 1))
```

As we can see with a low number of binds we can't see any significant difference in the summaries or histogram between the models with and without a nugget

Question 2


2 a)

We first start by making the appropriate changes in the data to average the data to quarterly means

```
AMOCDF$Date = as.Date(AMOCDF$Date, format = "%d/%m/%Y")  
  
## I will now make a column with the quarter and year that I will use  
## to create the averages per quarter  
AMOCDF$YearQuarter = paste(AMOCDF$Year, AMOCDF$Quarter, sep = "-")  
  
YearQuarterAverage = AMOCDF %>%  
  group_by(YearQuarter) %>%  
  summarise(AverageStrength = mean(Strength))
```

Now we will convert the average data to a time series object to be able to plot it

```
tsAtlantic = ts(YearQuarterAverage, start = c(2010, 1), frequency = 4)  
  
tsAtlantic = tsAtlantic[, "AverageStrength"]  
  
plot.ts(tsAtlantic)
```



Trend analysis

From this graph we can see a yearly oscillation of Sverdrups. We can also identify that the peaks in Sverdrups are usually in the last quarter before the start of a new year and the valleys are on the second quarter of the year.

The data does seem stationary enough that if we were to differentiate we would start losing some of the structure.


2 b)

ACF

First we will start by checking the ACF(Autocorrelation Function) and PACF(Partial Autocorrelation Function) to check for if we have stationary data or not to help us decide between an ARMA or an ARIMA model.

```
acf(tsAtlantic)
```


Series tsAtlantic



We can see that for ACF OF Average strength slowly decreases as lag increases to infinity with lag = 3 still being a significant values, meaning it is not a simple MA model as AR is clearly not quickly cut-off.

PACF

```
pacf(tsAtlantic)
```



The PACF seems to be cut-off at lag 0,5 indicating an AR model might be a best fit for our data to be a but with some almost significant values after the cut it might be also appropriate to some non-zero q values to confirm our initial assumption

As such we will now proceed to fit multiple model firstly with the initial assumption that, then I will both use models with non-zero q and the model given by the auto.arima function to double check that the assumptions made by the previous analyses is correct.

```
# it is always a good practice to try multiple values of p,d and q to
# see if we can do better we then obviously compare via the AIC of the
# models and their log likelihoods it is never enough to check those we
# also need to check the residuals

## order is p, d ,q

## initial models under our assumptions

model100 = Arima(tsAtlantic, order = c(1, 0, 0))
model200 = Arima(tsAtlantic, order = c(2, 0, 0))
model300 = Arima(tsAtlantic, order = c(3, 0, 0))

## now I will add postive q values
```

```

model101 = Arima(tsAtlantic, order = c(1, 0, 1))
model102 = Arima(tsAtlantic, order = c(1, 0, 2))
model103 = Arima(tsAtlantic, order = c(1, 0, 3))

model201 = Arima(tsAtlantic, order = c(2, 0, 1))
model202 = Arima(tsAtlantic, order = c(2, 0, 2))
model203 = Arima(tsAtlantic, order = c(2, 0, 3))

model301 = Arima(tsAtlantic, order = c(3, 0, 1))
model302 = Arima(tsAtlantic, order = c(3, 0, 2))
model303 = Arima(tsAtlantic, order = c(3, 0, 3))

## lastly we will use auto.arima without seasonality to confirm our
## initial assumptions
modelAuto = auto.arima(tsAtlantic, max.d = 0, max.p = 5, max.q = 5, seasonal = FALSE)

```

best model selection

```
model100
```

```

Series: tsAtlantic
ARIMA(1,0,0) with non-zero mean

Coefficients:
      ar1      mean
    0.0665  16.3878
s.e.  0.1788   0.3726

sigma^2 = 5.572:  log likelihood = -99.2
AIC=204.41   AICc=205.01   BIC=209.76

```

```
model200
```

```

Series: tsAtlantic
ARIMA(2,0,0) with non-zero mean

Coefficients:

```

```

          ar1      ar2      mean
0.0990 -0.5565 16.4298
s.e. 0.1576 0.1488 0.2113

sigma^2 = 4.321: log likelihood = -93.45
AIC=194.9 AICc=195.92 BIC=202.04

```

```
model300
```

```

Series: tsAtlantic
ARIMA(3,0,0) with non-zero mean

```

```

Coefficients:
          ar1      ar2      ar3      mean
0.1626 -0.5690 0.1464 16.4227
s.e. 0.1729 0.1479 0.1708 0.2409

sigma^2 = 4.35: log likelihood = -93.09
AIC=196.17 AICc=197.75 BIC=205.1

```

As we can see from these initial models ARIMA(2,0,0) is the model that has the best fit has we can see from its lower AIC score of 194,9.

Now we will check against the other models to check the validity of our assumptions.

```
model101
```

```

Series: tsAtlantic
ARIMA(1,0,1) with non-zero mean

```

```

Coefficients:
          ar1      ma1      mean
-0.4204 0.7718 16.3721
s.e. 0.2390 0.1466 0.4067

sigma^2 = 5.045: log likelihood = -96.64
AIC=201.29 AICc=202.31 BIC=208.43

```

```
model102
```

Series: tsAtlantic
ARIMA(1,0,2) with non-zero mean

Coefficients:

	ar1	ma1	ma2	mean
	0.0230	0.1275	-0.4485	16.4289
s.e.	0.3051	0.2420	0.1348	0.2224

sigma² = 4.651: log likelihood = -94.41
AIC=198.81 AICc=200.39 BIC=207.73

model103

Series: tsAtlantic
ARIMA(1,0,3) with non-zero mean

Coefficients:

	ar1	ma1	ma2	ma3	mean
	-0.5284	0.7214	-0.3646	-0.3072	16.4299
s.e.	0.9228	0.8649	0.2077	0.3545	0.2195

sigma² = 4.733: log likelihood = -94.25
AIC=200.5 AICc=202.77 BIC=211.21

model201

Series: tsAtlantic
ARIMA(2,0,1) with non-zero mean

Coefficients:

	ar1	ar2	ma1	mean
	-0.0669	-0.5475	0.2187	16.4255
s.e.	0.2740	0.1555	0.2883	0.2300

sigma² = 4.366: log likelihood = -93.15
AIC=196.31 AICc=197.88 BIC=205.23

model202

Series: tsAtlantic
ARIMA(2,0,2) with non-zero mean

Coefficients:

	ar1	ar2	ma1	ma2	mean
	0.0787	-0.9982	-0.0255	0.9999	16.4015
s.e.	0.0285	0.0066	0.0899	0.1158	0.2684

sigma² = 3.378: log likelihood = -89.46
AIC=190.91 AICc=193.18 BIC=201.62

model203

Series: tsAtlantic
ARIMA(2,0,3) with non-zero mean

Coefficients:

	ar1	ar2	ma1	ma2	ma3	mean
	0.0325	-0.9621	0.0499	0.8487	0.4147	16.4028
s.e.	0.0645	0.0442	0.1987	0.2041	0.2511	0.3044

sigma² = 3.315: log likelihood = -89.07
AIC=192.13 AICc=195.25 BIC=204.62

model301

Series: tsAtlantic
ARIMA(3,0,1) with non-zero mean

Coefficients:

	ar1	ar2	ar3	ma1	mean
	0.4092	-0.5931	0.2864	-0.2467	16.4191
s.e.	0.6330	0.1651	0.3580	0.6291	0.2537

sigma² = 4.449: log likelihood = -93.03
AIC=198.06 AICc=200.34 BIC=208.77

model302


```
Series: tsAtlantic
ARIMA(3,0,2) with non-zero mean
```

```
Coefficients:
```

	ar1	ar2	ar3	ma1	ma2	mean
	0.2684	-0.9851	0.2222	-0.3030	1.0000	16.4144
s.e.	0.1999	0.0305	0.1995	0.1453	0.1921	0.2922

```
sigma^2 = 3.392: log likelihood = -89.53
AIC=193.06 AICc=196.17 BIC=205.54
```

```
model303
```

```
Series: tsAtlantic
ARIMA(3,0,3) with non-zero mean
```

```
Coefficients:
```

	ar1	ar2	ar3	ma1	ma2	ma3	mean
	-0.3983	-0.9518	-0.4263	0.4381	0.7816	0.7352	16.4197
s.e.	0.3385	0.0412	0.3442	0.2946	0.1690	0.2422	0.2691

```
sigma^2 = 3.352: log likelihood = -88.54
AIC=193.08 AICc=197.19 BIC=207.35
```

In this initial analysis we have found models that do have a lower AIC lower log likelihood than our previous best model, however these model ma's standard error are to close the the ma values indicating that while we are getting a better fit we might be overfitting to our data.

As such this does confirm our initial assumption for the choice of a zero q value.

Now lastly we will check if the auto.arima function does confirm our initial assumptions.

```
modelAuto
```

```
Series: tsAtlantic
ARIMA(2,0,0) with non-zero mean
```

```
Coefficients:
```

	ar1	ar2	mean
	0.0990	-0.5565	16.4298
s.e.	0.1576	0.1488	0.2113

```
sigma^2 = 4.321: log likelihood = -93.45
AIC=194.9 AICc=195.92 BIC=202.04
```

The function does confirm our assumption that ARIMA(2,0,0) is indeed the best model.


We will now check the residuals to verify if any of our previously selected model validates well or if it is simply the best of bad models.

talk about the model being more easily explainability because MA = 0

Best model residual validation

```
# Set smaller margins
par(mar = c(4, 4, 2, 2))

tsdiag(model200)
```



```
# Reset margins
par(mar = c(5, 4, 4, 2) + 0.1)
```

Initially from the standardised residuals plot we can identify some sort of sinusoidal pattern, this implies that there is a seasonal trend that is not being accounted for in our model and

as such this trends needs to be accounted in future models to better explain and increase the prediction power of a new model.

Forecasting

Now using the forecast function we will forecast the next 4 quarters of 2021

```
forecast(model200, 4)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2021 Q1	16.50240	13.83841	19.16639	12.42818	20.57662
2021 Q2	14.81104	12.13403	17.48804	10.71691	18.90517
2021 Q3	16.22919	13.18168	19.27669	11.56843	20.88994
2021 Q4	17.31076	14.24941	20.37212	12.62882	21.99271

But this data is better visualized in a graph to better understand if the predictions are sensible compared to our real data.

```
predictedArimaDF = data.frame(forecast(model200, 4))

predictedArimaDF$YearQuarter = c("2021-Q1", "2021-Q2", "2021-Q3", "2021-Q4")

# Combine real_data and pred_data into a single data frame
combinedDataframeAMOC = rbind(data.frame(Date = YearQuarterAverage$YearQuarter,
    Temperature = YearQuarterAverage$AverageStrength, Type = "Real"), data.frame(Date = pr
    Temperature = predictedArimaDF$Point.Forecast, Type = "Predicted"))

predictedArimaDF$Temperature = predictedArimaDF$Point.Forecast

predictedArimaDF$Type = "Predicted"


# Create the ggplot
plotARIMA = ggplot(combinedDataframeAMOC, aes(x = Date, y = Temperature,
    color = Type, group = 1)) + geom_line() + scale_color_manual(values = c("blue",
    "red"))

# Add the 95% confidence interval
plotARIMA = plotARIMA + geom_ribbon(data = predictedArimaDF, aes(x = YearQuarter,
    ymin = Lo.95, ymax = Hi.95), fill = "red", alpha = 0.2)
```

```
# Adjust the x-axis labels
plotARIMA = plotARIMA + scale_x_discrete(breaks = combinedDataframeAMOC$Date[c(TRUE,
  rep(FALSE, 3))], labels = combinedDataframeAMOC$Date[c(TRUE, rep(FALSE,
  3))])

plotARIMA = plotARIMA + theme(axis.text.x = element_text(angle = 90, vjust = 0.5,
  size = 8))

# Display the plot
print(plotARIMA)
```



As we can see from the graph the ARIMA (2,0,0) seems to give us a sensible forecast for the 2021 quarter values, however as we can see the interval of the prediction accuracy our model is not too certain on the values most likely due to our model not accounting for the seasonal cycle of our data.

2 c)

Initial assumptions

From the previous exploratory analysis of the data we have established that the data did not

need to be differentiated since it was constant, this translates to polynomial DLM component of order 2 that will use linear model to account for this type of changes in the data.

Furthermore, from the residual analysis we have inferred that there is an underlying seasonal trend present on the data, this seasonal trend will be represented by a seasonal component of frequency 4 to represent the 4 quarters per year.

model fitting

```
## linear model, order = 2, quadratic order = 3 , etc

## what we want is a linear model with a seasonal component so we add
## the 2 components together in a model

## things to try, another term like quadratic, or a arma component
## stacked on top of this

## Initial model with a linear polynomial and a seasonal component

buildFun = function(x) {
  dlmModPoly(order = 2, dV = exp(x[1]), dW = c(0, exp(x[2]))) + dlmModSeas(frequency = 4,
    dV = 0, dW = c(exp(x[3]), rep(0, 2)))
}

linearDLM = dlmMLE(tsAtlantic, parm = c(0, 0, 0), build = buildFun)

linearDLM$par
```

```
[1] 1.151339 -18.078101 -2.189479
```

```
fittedLinearDLM = buildFun(linearDLM$par)
```

```
V(fittedLinearDLM)
```

```
      [,1]
[1,] 3.162425
```

```
W(fittedLinearDLM)
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	0	0.000000e+00	0.000000	0	0
[2,]	0	1.408576e-08	0.000000	0	0
[3,]	0	0.000000e+00	0.111975	0	0
[4,]	0	0.000000e+00	0.000000	0	0
[5,]	0	0.000000e+00	0.000000	0	0

```
## second model with a quadratic polynomial and a seasonal component
```

```
buildFunQuad = function(x) {
  dlmModPoly(order = 3, dV = exp(x[1]), dW = c(0, exp(x[2]), exp(x[3]))) +
  dlmModSeas(frequency = 4, dV = 0, dW = c(exp(x[4]), rep(0, 2)))
}
```

```
quadraticDLM = dlmMLE(tsAtlantic, parm = c(0, 0, 0, 0), build = buildFunQuad)
```

```
quadraticDLM$par
```

```
[1] 1.161355 -17.807081 -28.603103 -2.352292
```

```
fittedQuadraticDLM = buildFunQuad(quadraticDLM$par)
```

```
V(fittedQuadraticDLM)
```

	[,1]
[1,]	3.194257

```
W(fittedQuadraticDLM)
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	0	0.000000e+00	0.000000e+00	0.00000000	0	0
[2,]	0	1.847069e-08	0.000000e+00	0.00000000	0	0
[3,]	0	0.000000e+00	3.782948e-13	0.00000000	0	0
[4,]	0	0.000000e+00	0.000000e+00	0.09515082	0	0
[5,]	0	0.000000e+00	0.000000e+00	0.00000000	0	0
[6,]	0	0.000000e+00	0.000000e+00	0.00000000	0	0

Now we will compare both models through their log likelihood using the `dlmLL` function and see if the extra flexibility from the extra polynomial function is providing a better fit

```
dmlLL(tsAtlantic, fittedLinearDLM)
```

```
[1] 94.98804
```

```
dmlLL(tsAtlantic, fittedQuadraticDLM)
```

```
[1] 108.043
```


As we can see the dlm model using only a linear polynomial has a lower log likelihood than the model with an extra quadratic term, meaning this extra flexibility does not contribute to a better model fit and as such we will use the linear fitted model to do our forecasting.

```
amocPredict = dlmFilter(tsAtlantic, mod = fittedLinearDLM)
summary(amocPredict)
```

	Length	Class	Mode
y	44	ts	numeric
mod	10	dlm	list
m	225	mts	numeric
U.C	45	-none-	list
D.C	225	-none-	numeric
a	220	mts	numeric
U.R	44	-none-	list
D.R	220	-none-	numeric
f	44	ts	numeric

```
x = cbind(tsAtlantic, dropFirst(amocPredict$a[, c(1, 3)]))
x = window(x, start = c(2010, 1))
colnames(x) = c("Gas", "Trend", "Seasonal")
plot(x, type = "o", main = "Atlantic AMOC at 26,5N 2010-2020")
```

Atlantic AMOC at 26,5N 2010–2020



Forecast

```
amocForecast = dlmForecast(amocPredict, nAhead = 4)
summary(amocForecast)
```

```
Length Class Mode
a 20      mts   numeric
R  4      -none- list
f  4      ts    numeric
Q  4      -none- list
```

```
dim(amocForecast$a)
```

```
[1] 4 5
```

```
dim(amocForecast$f)
```

```
[1] 4 1
```




```

sqrtR = sapply(amocForecast$R, function(x) sqrt(x[1, 1]))
pl = amocForecast$a[, 1] + qnorm(0.025, sd = sqrtR)
pu = amocForecast$a[, 1] + qnorm(0.975, sd = sqrtR)
x = ts.union(window(tsAtlantic, start = c(2010, 1)), amocForecast$a[, 1],
              amocForecast$f, pl, pu)
par(mar = c(4, 4, 2, 2))
plot(x, plot.type = "single", type = "o", pch = c(1, 20, 3, NA, NA), col = c("darkgrey",
  "brown", "brown", "blue", "blue"), ylab = "Log gas consumption")

legend("bottomright", legend = c("Observed", "Forecast", "95% interval"),
      bty = "n", pch = c(1, 20, NA), lty = 1, col = c("darkgrey", "brown",
  "blue"))

```




```

# Set smaller margins
par(mar = c(4, 4, 2, 2))

tsdiag(amocPredict)

```




```
# Reset margins
par(mar = c(5, 4, 4, 2) + 0.1)
```

2 d)

Again comparing the forecast values and their respective prediction intervals as we can see from the graphs below the dlm model has smaller prediction intervals, most likely due to being able to explain the underlying seasonal trend reducing therefore the uncertainty in comparison the ARIMA model.

```
print(plotARIMA)
```




```

sqaretRoot = sapply(amocForecast$R, function(x) sqrt(x[1, 1]))
predictionLow = amocForecast$a[, 1] + qnorm(0.025, sd = sqaretRoot) ## Low
predictionUpper = amocForecast$a[, 1] + qnorm(0.975, sd = sqaretRoot) ## Upper
x = ts.union(window(tsAtlantic, start = c(2010, 1)), amocForecast$a[, 1],
              amocForecast$f, predictionLow, predictionUpper)
par(mar = c(4, 4, 2, 2))
plot(x, plot.type = "single", type = "o", pch = c(1, 20, 3, NA, NA), col = c("darkgrey",
  "brown", "brown", "blue", "blue"), ylab = "Log gas consumption")

legend("bottomright", legend = c("Observed", "Forecast", "95% interval"),
      bty = "n", pch = c(1, 20, NA), lty = 1, col = c("darkgrey", "brown",
  "blue"))

```



2 e)

```
# AMOCDFMonthly =AMOCDF %>% mutate(YearMonth = paste0(year(Date), '-',
# month(Date, label = TRUE, abbr = FALSE)))
```

```
## I will now make a column with the month and year that I will use to
## create the monthly averages
```

```
AMOCDF$YearMonth = paste(AMOCDF$Year, AMOCDF$Month, sep = "-")
```


```
YearMonthlyAverage = AMOCDF %>%
  group_by(YearMonth) %>%
  summarise(AverageStrength = mean(Strength))
```

Now we will create a new montly time series object and make it univariate

```
tsAtlanticMonthly = ts(YearMonthlyAverage, start = c(2010, 1), frequency = 12)

tsAtlanticMonthly = tsAtlanticMonthly[, "AverageStrength"]


plot.ts(tsAtlanticMonthly)
```



Seeing this graph we can observe that the data continues being stationary for the ARIMA model but a seasonal trend not only is more apparently but it also appear to need to be differentiated as it seems to have a decreasing linear trend


```
acf(tsAtlanticMonthly)
```

Series tsAtlanticMonthly



```
pacf(tsAtlanticMonthly)
```

Series tsAtlanticMonthly



The acf has a very clear cut-off as only 3 the values are significant which is very similar to what we had observed previously.

The main difference is in the pacf, where we can now say for sure that there is a very clear cut-off after the first value.

Model testing

These pattern suggests that an ARMA/ARIMA model might be the most appropriate so first we will check them out with the seasonal component of order 1, the so quick cut-off of both the ACF and PACF also might suggest that p and q will be smaller values.

Seasonal check

```
## initial assumption

modelMonthlySeasonal100.110 = Arima(tsAtlanticMonthly, order = c(1, 0, 0),
  seasonal = list(order = c(1, 1, 0), period = 12))

modelMonthlySeasonal100.011 = Arima(tsAtlanticMonthly, order = c(1, 0, 0),
  seasonal = list(order = c(0, 1, 1), period = 12))

modelMonthlySeasonal200.210 = Arima(tsAtlanticMonthly, order = c(2, 0, 0),
  seasonal = list(order = c(2, 1, 0), period = 12))
modelMonthlySeasonal200.012 = Arima(tsAtlanticMonthly, order = c(2, 0, 0),
  seasonal = list(order = c(0, 1, 2), period = 12))

modelMonthlySeasonal001.110 = Arima(tsAtlanticMonthly, order = c(0, 0, 1),
  seasonal = list(order = c(1, 1, 0), period = 12))

modelMonthlySeasonal001.011 = Arima(tsAtlanticMonthly, order = c(0, 0, 1),
  seasonal = list(order = c(0, 1, 1), period = 12))

modelMonthlySeasonal002.210 = Arima(tsAtlanticMonthly, order = c(0, 0, 2),
  seasonal = list(order = c(2, 1, 0), period = 12))
modelMonthlySeasonal002.012 = Arima(tsAtlanticMonthly, order = c(0, 0, 2),
  seasonal = list(order = c(0, 1, 2), period = 12))

modelMonthlySeasonal100.110
```

Series: tsAtlanticMonthly
ARIMA(1,0,0)(1,1,0)[12]

Coefficients:

	ar1	sar1
	0.1779	-0.4618
s.e.	0.0909	0.0839

sigma^2 = 12.06: log likelihood = -320.1
AIC=646.2 AICc=646.4 BIC=654.56

modelMonthlySeasonal100.011

Series: tsAtlanticMonthly
ARIMA(1,0,0)(0,1,1)[12]

Coefficients:

	ar1	sma1
	0.1844	-0.8500
s.e.	0.0918	0.1123

sigma^2 = 8.74: log likelihood = -306.88
AIC=619.76 AICc=619.97 BIC=628.12

modelMonthlySeasonal200.210

Series: tsAtlanticMonthly
ARIMA(2,0,0)(2,1,0)[12]

Coefficients:

	ar1	ar2	sar1	sar2
	0.115	0.0374	-0.7271	-0.4814
s.e.	0.094	0.0930	0.0932	0.0933

sigma^2 = 9.741: log likelihood = -309.65
AIC=629.29 AICc=629.82 BIC=643.23

modelMonthlySeasonal200.012


```
Series: tsAtlanticMonthly  
ARIMA(2,0,0)(0,1,2)[12]
```

```
Coefficients:
```

	ar1	ar2	sma1	sma2
	0.1566	0.0546	-0.9817	0.1895
s.e.	0.0947	0.0950	0.1384	0.1503

```
sigma^2 = 8.778: log likelihood = -305.88  
AIC=621.76 AICc=622.29 BIC=635.7
```

```
modelMonthlySeasonal001.110
```

```
Series: tsAtlanticMonthly  
ARIMA(0,0,1)(1,1,0)[12]
```

```
Coefficients:
```

	ma1	sar1
	0.1680	-0.4634
s.e.	0.0861	0.0840

```
sigma^2 = 12.07: log likelihood = -320.19  
AIC=646.39 AICc=646.59 BIC=654.75
```

```
modelMonthlySeasonal001.011
```

```
Series: tsAtlanticMonthly  
ARIMA(0,0,1)(0,1,1)[12]
```

```
Coefficients:
```

	ma1	sma1
	0.1606	-0.8446
s.e.	0.0847	0.1093

```
sigma^2 = 8.804: log likelihood = -307.14  
AIC=620.28 AICc=620.49 BIC=628.65
```

```
modelMonthlySeasonal002.210
```

```
Series: tsAtlanticMonthly
ARIMA(0,0,2)(2,1,0)[12]
```

```
Coefficients:
```

	ma1	ma2	sar1	sar2
	0.1145	0.0452	-0.7275	-0.4806
s.e.	0.0943	0.0882	0.0933	0.0936

```
sigma^2 = 9.745: log likelihood = -309.66
AIC=629.33 AICc=629.85 BIC=643.26
```

```
modelMonthlySeasonal002.012
```

```
Series: tsAtlanticMonthly
ARIMA(0,0,2)(0,1,2)[12]
```

```
Coefficients:
```

	ma1	ma2	sma1	sma2
	0.1583	0.0745	-0.9786	0.1868
s.e.	0.0953	0.0893	0.1377	0.1495

```
sigma^2 = 8.784: log likelihood = -305.89
AIC=621.78 AICc=622.3 BIC=635.72
```

So as suspected from both the time series plot and the last exercise analysis the added seasonality does increase our model goodness of fit while also penalising the increased in complexity with so far.

Now lets compare them to bigger p and q values to see if our initial assumptions do hold up

```
modelMonthlySeasonal301 = Arima(tsAtlanticMonthly, order = c(3, 0, 1), seasonal = list(ord
1, 1), period = 12))
modelMonthlySeasonal302 = Arima(tsAtlanticMonthly, order = c(3, 0, 2), seasonal = list(ord
1, 0), period = 12))
modelMonthlySeasonal303 = Arima(tsAtlanticMonthly, order = c(3, 1, 3), seasonal = list(ord
1, 0), period = 12))

modelMonthlySeasonal103 = Arima(tsAtlanticMonthly, order = c(1, 0, 3), seasonal = list(ord
1, 1), period = 12))
modelMonthlySeasonal203 = Arima(tsAtlanticMonthly, order = c(2, 1, 3), seasonal = list(ord
1, 0), period = 12))
```

modelMonthlySeasonal301

Series: tsAtlanticMonthly

ARIMA(3,0,1)(1,1,1)[12]

Coefficients:

	ar1	ar2	ar3	ma1	sar1	sma1
	-0.5728	0.1878	-0.0048	0.7442	-0.1158	-0.8073
s.e.	0.7595	0.1720	0.1369	0.7564	0.1224	0.1140

$\sigma^2 = 8.939$: log likelihood = -306.02

AIC=626.04 AICc=627.04 BIC=645.55

modelMonthlySeasonal302

Series: tsAtlanticMonthly

ARIMA(3,0,2)(1,1,0)[12]

Coefficients:

	ar1	ar2	ar3	ma1	ma2	sar1
	-1.4367	-0.5902	0.1291	1.6983	1.0000	-0.4382
s.e.	0.0944	0.1532	0.0945	0.0436	0.0489	0.0871

$\sigma^2 = 11.34$: log likelihood = -316.59

AIC=647.19 AICc=648.19 BIC=666.7

modelMonthlySeasonal303

Series: tsAtlanticMonthly

ARIMA(3,1,3)(1,1,0)[12]

Coefficients:

	ar1	ar2	ar3	ma1	ma2	ma3	sar1
	-1.4279	-0.5751	0.1373	0.6984	-0.6984	-1.0000	-0.4313
s.e.	0.0952	0.1547	0.0954	0.0533	0.0552	0.0576	0.0877

$\sigma^2 = 11.54$: log likelihood = -316.86

AIC=649.73 AICc=651.04 BIC=671.96

```
modelMonthlySeasonal103
```

```
Series: tsAtlanticMonthly  
ARIMA(1,0,3)(1,1,1)[12]
```

```
Coefficients:
```

	ar1	ma1	ma2	ma3	sar1	sma1
	-0.6951	0.8708	0.1931	-0.0013	-0.1134	-0.8029
s.e.	0.4854	0.4812	0.1503	0.1321	0.1194	0.1136

```
sigma^2 = 8.954: log likelihood = -306.02  
AIC=626.03 AICc=627.03 BIC=645.55
```

```
modelMonthlySeasonal203
```

```
Series: tsAtlanticMonthly  
ARIMA(2,1,3)(1,1,0)[12]
```

```
Coefficients:
```

	ar1	ar2	ma1	ma2	ma3	sar1
	-1.3087	-0.6017	0.5001	-0.6702	-0.8299	-0.4545
s.e.	0.2576	0.1557	0.2045	0.1332	0.1180	0.0979

```
sigma^2 = 12.02: log likelihood = -318.1  
AIC=650.19 AICc=651.2 BIC=669.65
```

As we can see here the initial assumption that a smaller p and q value would better fit the model.

Now we will use `auto.arima` to verify if our assumptions were indeed correct

auto arima check

```
`?`(auto.arima)
```

```
starting httpd help server ... done
```

```
monthlyModelAuto = auto.arima(tsAtlanticMonthly, max.d = 0, max.p = 5, max.q = 5,
                               D = 1)
monthlyModelAuto
```

Series: tsAtlanticMonthly
ARIMA(1,0,0)(0,1,1)[12]

Coefficients:

```
          ar1      sma1
          0.1844  -0.8500
s.e.      0.0918   0.1123
```

sigma^2 = 8.74: log likelihood = -306.88
AIC=619.76 AICc=619.97 BIC=628.12

From what we can see the auto arima has indeed confirmed our initial assumption by picking a model that we already had seen as the best performer ARIMA(1,0,0)(0,1,1)[12]

Forecasting

Now using the forecast function we will forecast the next 4 quarters of 2021

```
forecast(modelMonthlySeasonal100.011, 12)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Jan 2021	15.29474	11.49064	19.09884	9.476872	21.11261
Feb 2021	17.68442	13.81636	21.55247	11.768741	23.60009
Mar 2021	19.72032	15.85011	23.59053	13.801347	25.63929
Apr 2021	17.53922	13.66894	21.40951	11.620140	23.45831
May 2021	16.17383	12.30355	20.04411	10.254741	22.09292
Jun 2021	14.65775	10.78746	18.52803	8.738657	20.57683
Jul 2021	14.12800	10.25772	17.99829	8.208914	20.04709
Aug 2021	15.37476	11.50447	19.24504	9.455670	21.29385
Sep 2021	15.71081	11.84052	19.58109	9.791718	21.62989
Oct 2021	17.33660	13.46632	21.20689	11.417514	23.25569
Nov 2021	17.56114	13.69086	21.43141	11.642055	23.48022
Dec 2021	16.87677	13.00663	20.74691	10.957906	22.79563

But this data is better visualized in a graph to better understand if the predictions are sensible compared to our real data.

```

predictedArimaSeasonalDF = data.frame(forecast(modelMonthlySeasonal100.011,
12))

predictedArimaSeasonalDF$YearMonth = c("2021-1", "2021-2", "2021-3", "2021-4",
"2021-5", "2021-6", "2021-7", "2021-8", "2021-9", "2021-10", "2021-11",
"2021-12")

predictedArimaSeasonalDF$Type = "Predicted"

predictedArimaSeasonalDF$Temperature = predictedArimaSeasonalDF$Point.Forecast

YearMonthlyAverage$Type = "Real"
YearMonthlyAverage$Temperature = YearMonthlyAverage$AverageStrength

# Combine real_data and pred_data into a single data frame
combinedDataframeAMOC = rbind(data.frame(Date = YearMonthlyAverage$YearMonth,
Temperature = YearMonthlyAverage$Temperature, Type = "Real"), data.frame(Date = predictedArimaSeasonalDF$YearMonth,
Temperature = predictedArimaSeasonalDF$Temperature, Type = "Predicted"))

# Create the ggplot
plotARIMA2 = ggplot(combinedDataframeAMOC, aes(x = Date, y = Temperature,
color = Type, group = 1)) + geom_line() + scale_color_manual(values = c("blue",
"red"))


# Add the 80% and 95% confidence intervals
plotARIMA2 = plotARIMA2 + geom_ribbon(data = predictedArimaSeasonalDF, aes(x = YearMonth,
ymin = Lo.95, ymax = Hi.95), fill = "red", alpha = 0.2) + geom_ribbon(data = predictedArimaSeasonalDF,
aes(x = YearMonth, ymin = Lo.80, ymax = Hi.80), fill = "blue", alpha = 0.2)

# Adjust the x-axis labels
plotARIMA2 = plotARIMA2 + scale_x_discrete(breaks = c("2021-1", "2021-4",
"2021-8", "2021-12"), labels = c("Q1", "Q2", "Q3", "Q4"))

plotARIMA2 = plotARIMA2 + theme(axis.text.x = element_text(angle = 90, vjust = 0.5,
size = 8))

# Display the plot
print(plotARIMA2)

```



DLM

model fitting

```
## linear model, order = 2, quadratic order = 3 , etc

## what we want is a linear model with a seasonal component so we add
## the 2 components together in a model

## things to try, another term like quadratic, or a arma component
## stacked on top of this

## Initial model with a linear polynomial and a seasonal component

buildFun = function(x) {
  dlmModPoly(order = 2, dV = exp(x[1]), dW = c(0, exp(x[2]))) + dlmModSeas(frequency = 12,
    dV = 0, dW = c(exp(x[3]), rep(0, 10)))
}

linearDLM = dlmMLE(tsAtlanticMonthly, parm = c(0, 0, 0), build = buildFun)
```

```
linearDLM$par
```

```
[1] 2.044026 -11.586366 -4.421181
```

```
fittedLinearDLM = buildFun(linearDLM$par)
```

```
V(fittedLinearDLM)
```

```
[,1]
```

```
[1,] 7.721632
```

```
W(fittedLinearDLM)
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]
[1,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[2,]	0	9.291914e-06	0.00000000	0	0	0	0	0	0	0	0
[3,]	0	0.000000e+00	0.01202003	0	0	0	0	0	0	0	0
[4,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[5,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[6,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[7,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[8,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[9,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[10,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[11,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[12,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[13,]	0	0.000000e+00	0.00000000	0	0	0	0	0	0	0	0
[,12] [,13]											
[1,]	0	0									
[2,]	0	0									
[3,]	0	0									
[4,]	0	0									
[5,]	0	0									
[6,]	0	0									
[7,]	0	0									
[8,]	0	0									
[9,]	0	0									
[10,]	0	0									
[11,]	0	0									


```
[12,]    0    0
[13,]    0    0
```

```
## second model with a quadratic polynomial and a seasonal component
```

```
buildFunQuad = function(x) {
  dlmModPoly(order = 3, dV = exp(x[1]), dW = c(0, exp(x[2]), exp(x[3]))) +
  dlmModSeas(frequency = 12, dV = 0, dW = c(exp(x[4]), rep(0, 10)))
}

quadraticDLM = dlmMLE(tsAtlanticMonthly, parm = c(0, 0, 0, 0), build = buildFunQuad)

quadraticDLM$par
```

```
[1] 2.047135 -21.060474 -56.104784 -12.300531
```

```
fittedQuadraticDLM = buildFunQuad(quadraticDLM$par)

V(fittedQuadraticDLM)
```

```
      [,1]
[1,] 7.745678
```

```
W(fittedQuadraticDLM)
```

```
      [,1]      [,2]      [,3]      [,4] [,5] [,6] [,7] [,8] [,9]
[1,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[2,] 0 7.137603e-10 0.000000e+00 0.000000e+00 0 0 0 0 0
[3,] 0 0.000000e+00 4.305286e-25 0.000000e+00 0 0 0 0 0
[4,] 0 0.000000e+00 0.000000e+00 4.549326e-06 0 0 0 0 0
[5,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[6,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[7,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[8,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[9,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[10,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[11,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
[12,] 0 0.000000e+00 0.000000e+00 0.000000e+00 0 0 0 0 0
```

```

[13,]    0 0.000000e+00 0.000000e+00 0.000000e+00    0    0    0    0    0
[14,]    0 0.000000e+00 0.000000e+00 0.000000e+00    0    0    0    0    0
      [,10] [,11] [,12] [,13] [,14]
[1,]      0     0     0     0     0
[2,]      0     0     0     0     0
[3,]      0     0     0     0     0
[4,]      0     0     0     0     0
[5,]      0     0     0     0     0
[6,]      0     0     0     0     0
[7,]      0     0     0     0     0
[8,]      0     0     0     0     0
[9,]      0     0     0     0     0
[10,]     0     0     0     0     0
[11,]     0     0     0     0     0
[12,]     0     0     0     0     0
[13,]     0     0     0     0     0
[14,]     0     0     0     0     0

```

Now we will compare both models through their log likelihood using the `dmlL` function and see if the extra flexibility from the extra polynomial function is providing a better fit

```
dmlL(tsAtlanticMonthly, fittedLinearDLM)
```

```
[1] 309.5446
```

```
dmlL(tsAtlanticMonthly, fittedQuadraticDLM)
```

```
[1] 324.4748
```

As we can see the dlm model using only a linear polynomial has a lower log likelihood than the model with an extra quadratic term, meaning this extra flexibility does not contribute to a better model fit and as such we will use the linear fitted model to do our forecasting.

```
amocPredict = dlmFilter(tsAtlanticMonthly, mod = fittedLinearDLM)
summary(amocPredict)
```

```

      Length Class  Mode
y       132   ts     numeric
mod      10   dlm     list

```

```

m 1729 mts numeric
U.C 133 -none- list
D.C 1729 -none- numeric
a 1716 mts numeric
U.R 132 -none- list
D.R 1716 -none- numeric
f 132 ts numeric


```

```

x = cbind(tsAtlanticMonthly, dropFirst(amocPredict$a[, c(1, 3)]))
x = window(x, start = c(2010, 1))
colnames(x) = c("Gas", "Trend", "Seasonal")
plot(x, type = "o", main = "Atlantic AMOC at 26,5N 2010-2020")

```

Atlantic AMOC at 26,5N 2010–2020



Forecast

```

amocForecastMonthly = dlmForecast(amocPredict, nAhead = 12)
summary(amocForecastMonthly)

```

```

Length Class Mode
a 156 mts numeric

```

```
R 12 -none- list
f 12 ts      numeric
Q 12 -none- list
```

```
dim(amocForecastMonthly$a)
```

```
[1] 12 13
```


```
dim(amocForecastMonthly$f)
```

```
[1] 12 1
```


```
sqrtR = sapply(amocForecastMonthly$R, function(x) sqrt(x[1, 1]))
pl = amocForecastMonthly$a[, 1] + qnorm(0.025, sd = sqrtR)
pu = amocForecastMonthly$a[, 1] + qnorm(0.975, sd = sqrtR)

x = ts.union(window(tsAtlanticMonthly, start = c(2010, 1)), amocForecastMonthly$a[,
  1], amocForecastMonthly$f, pl, pu)
par(mar = c(4, 4, 2, 2))
plot(x, plot.type = "single", type = "o", pch = c(1, 20, 3, NA, NA), col = c("darkgrey",
  "brown", "brown", "blue", "blue"), ylab = "Log gas consumption")

legend("bottomright", legend = c("Observed", "Forecast", "95% interval"),
  bty = "n", pch = c(1, 20, NA), lty = 1, col = c("darkgrey", "brown",
  "blue"))
```



```
# Set smaller margins  
par(mar = c(4, 4, 2, 2))  
  
tsdiag(amocPredict)
```



```
# Reset margins
par(mar = c(5, 4, 4, 2) + 0.1)
```

Lastly checking the residuals, they seem to be mostly normally distributed with a good mixture of over and under estimations, especially in the middle with some slight seasonality on both ends being present

2 f)

Now starting with the ARIMA models

```
predictedArimaDF = data.frame(forecast(model200, 4))

predictedArimaDF$YearQuarter = c("2021-Q1", "2021-Q2", "2021-Q3", "2021-Q4")

# Combine real_data and pred_data into a single data frame
combinedDataframeAMOC = rbind(data.frame(Date = YearQuarterAverage$YearQuarter,
  Temperature = YearQuarterAverage$AverageStrength, Type = "Real"), data.frame(Date = pr
  Temperature = predictedArimaDF$Point.Forecast, Type = "Predicted"))

predictedArimaDF$Temperature = predictedArimaDF$Point.Forecast
```

```

predictedArimaDF$Type = "Predicted"

# Create the ggplot
plotARIMA = ggplot(combinedDataframeAMOC, aes(x = Date, y = Temperature,
  color = Type, group = 1)) + geom_line() + scale_color_manual(values = c("blue",
  "red"))


# Add the 95% confidence interval
plotARIMA = plotARIMA + geom_ribbon(data = predictedArimaDF, aes(x = YearQuarter,
  ymin = Lo.95, ymax = Hi.95), fill = "red", alpha = 0.2)

# Adjust the x-axis labels
plotARIMA = plotARIMA + scale_x_discrete(breaks = combinedDataframeAMOC$Date[c(TRUE,
  rep(FALSE, 3))], labels = combinedDataframeAMOC$Date[c(TRUE, rep(FALSE,
  3))])


plotARIMA = plotARIMA + theme(axis.text.x = element_text(angle = 90, vjust = 0.5,
  size = 8))

print(plotARIMA)

```



```
print(plotARIMA2)
```



The most obvious difference is the level of detail on the seasonality, with the monthly averages capturing an almost opposite effect than the quarterly averages in both the predicted and also in some of the real data, the forecast also has the opposite trend, with an actual expected decrease in Sverdrups around between the 2nd quarter and the mid 3rd quarter which again seems to follow the opposite trend on the quarterly data.

```
sqrtr = sapply(amocForecastMonthly$R, function(x) sqrt(x[1, 1]))
pl = amocForecastMonthly$a[, 1] + qnorm(0.025, sd = sqrtr)
pu = amocForecastMonthly$a[, 1] + qnorm(0.975, sd = sqrtr)

x = ts.union(window(tsAtlanticMonthly, start = c(2010, 1)), amocForecastMonthly$a[,
  1], amocForecastMonthly$f, pl, pu)
par(mar = c(4, 4, 2, 2))
plot(x, plot.type = "single", type = "o", pch = c(1, 20, 3, NA, NA), col = c("darkgrey",
  "brown", "brown", "blue", "blue"), ylab = "Log gas consumption")

legend("bottomright", legend = c("Observed", "Forecast", "95% interval"),
  bty = "n", pch = c(1, 20, NA), lty = 1, col = c("darkgrey", "brown",
  "blue"))
```






Figure 1: 1st dlm



Observing the 2 graphs of the dlm models we can see that although the quarterly predictions do not completely predict the spike during the first half of the year compared to the monthly data the second half of the year seems to be relative similarly forecasted with the exception of December where the monthly data show again a decrease but the quarterly data is not capable of capturing.

Despite these changes the predicted overall trend is quite similar with the quarterly trend very slightly increasing the monthly data seeming to remain constant.

Question 3

Question 3 a)

I will start with the time series analysis of the temperature in California
other approach see max temp in the entire state with 8 cities

TODO WARNING

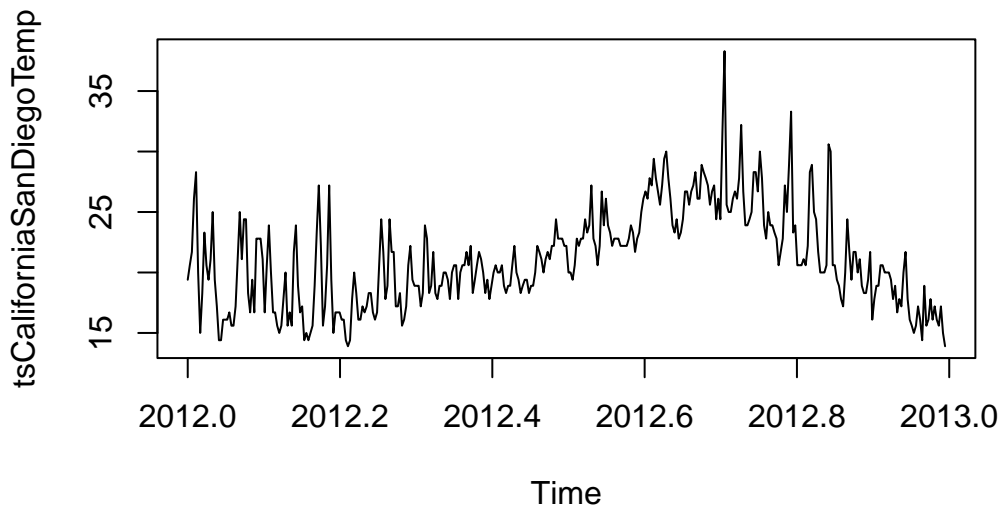
For a dataset of daily data with only 1 year of a cycle data available a daily frequency won't be a very good fit because we only have one observation per cycle, we need a hidden entry to capture the 12 months instead

```
californiaTempDF$Date = as.Date(as.character(californiaTempDF$Date), format = "%Y%m%d",
  origin = "1970-01-01")

## its better to just get a time series object for each city and plot
## each of those in the same plot

tsCaliforniaSanDiegoTemp = ts(californiaTempDF$`San Diego`, start = c(2012,
  1), frequency = 366)

plot.ts(tsCaliforniaSanDiegoTemp)
```

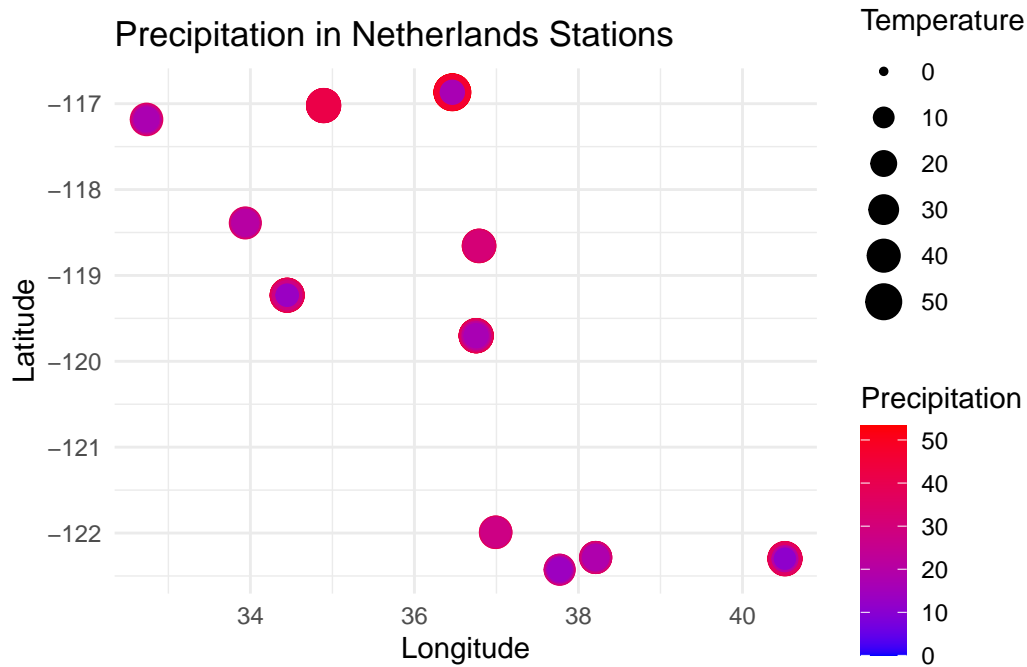


```
californiaLongTempDF = pivot_longer(californiaTempDF, cols = -Date, names_to = "Location",
  values_to = "Temperature")

spatialTemperatureCaliforniaDF = merge(californiaLongTempDF, californiaSpatialDataDF)

ggplot(data = spatialTemperatureCaliforniaDF) + geom_point(aes(x = Lat, y = Long,
  color = Temperature, size = Temperature)) + scale_color_continuous(low = "blue",
  high = "red") + labs(title = "Precipitation in Netherlands Stations",
```

```
x = "Longitude", y = "Latitude", color = "Precipitation") + theme_minimal()
```



3 b)

```
geoDataCalifornia = as.geodata(spatialTemperatureCaliforniaDF, coords.col = 4:5,  
  data.col = "Temperature", covar.col = "Elev")
```

as.geodata: 4004 replicated data locations found.

Consider using jitterDupCoords() for jittering replicated locations.

WARNING: there are data at coincident or very closed locations, some of the geoR's functions

Use function dup.coords() to locate duplicated coordinates.

Consider using jitterDupCoords() for jittering replicated locations

```
variogramCalifornia = variog(geoDataCalifornia)
```

variog: computing omnidirectional variogram

variog: co-located data found, adding one bin at the origin

```
plot(variogramCalifornia)
```

